



# Oceananigans.jl



## an ocean-flavoured fluid dynamics library



Navid Constantinou (@navidcy)  & the CliMA Ocean Dev Team




JuliaCon2024



near-global ( $75^{\circ}\text{S}$ – $75^{\circ}\text{N}$ ) ocean simulation at  $1/12^{\circ}$  horizontal resolution, 48 vertical levels  
@ 68 Nvidia A100 achieving 10 simulated years per day

# requirements for a climate/ocean model

Computational  
efficiency



Flexibility and  
ease of use

- ◆ Necessary for global calibration
- ◆ Possibility of high-resolution


- ◆ Simulate physics from meters to global-scale
- ◆ Support rapid prototyping of parameterizations

*“A fast model can be a good model,  
but a good model must be a fast model!  
Computational efficiency is crucial...”*

# Oceananigans: Fast and Efficient




global ocean simulation forced with atmospheric JRA55 reanalysis



**fast compute**

written from scratch for GPUs




**memory leanness**

minimize temporary array creations



**scalability**

overlap communication & computation




**16 km** horizontal resolution: **10 SYPD** on **8 GPUs**

10 Simulated years per day (SYPD):  
threshold for climate projections

**8 km** horizontal resolution: **10 SYPD** on **64 GPUs**

**2 km** horizontal resolution: **>1 SYPD** on **512 GPUs**

# Oceananigans: Easy to use and Accessible



```
1 using Oceananigans
2 using GLMakie
3
4 grid = RectilinearGrid(CPU(),
5                         size = (64, 64),
6                         x = (-5, 5),
7                         y = (-5, 5),
8                         topology = (Bounded, Bounded, Flat))
9
10 model = NonhydrostaticModel(grid=grid, tracers=:c, advection=WENO())
11
12 gaussian(x, y) = exp(-(x^2 + y^2)) ← initial conditions
13 set!(model, c=gaussian)
14
15 c = model.tracers.c
16
17 ∇c² = ∂x(c)^2 + ∂y(c)^2 ← diagnostics
18 ∇c² = Field(∇c²)
19 compute!(∇c²)
```

used in *more than 20* scientific papers

55+ contributors to the codebase


*"...I have never experienced getting a useful calculation done as easily as I was able to do with Oceananigans. It not only has a sophisticated interface, but it is remarkably fast...".*

*Linux magazine*


User interface:

- Designed so code “reads like a paper”
- Should not require comments

# Oceananigans: Flexible




# Oceananigans: Flexible



# Injecting code in a simulation: forcing with a neural net

```
using Oceananigans,  
    Oceananigans.Units  
using Lux  
  
grid = LatitudeLongitudeGrid(GPU(); kw...)  
  
u_sgs = Field(grid)  
  
model = HydrostaticFreeSurfaceModel(; grid, forcing = (; u = u_sgs), kw...)  
  
simulation = Simulation(model; Δt = 10minutes, stop_time = 10days)  
  
NN = Chain(args...) |> gpu # A neural network that computes u -> u_sgs  
  
function neural_network_inference(simulation)  
    u_sgs = simulation.model.forcing.u  
    u = simulation.model.velocities.u  
    u_sgs .= NN(u)  
end  
  
simulation.callbacks[:apply_nn] = Callback(neural_network_inference,  
                                         IterationInterval(1))  
  
run!(simulation)
```


Simple and effective way to add a NN in Oceananigans thanks to:



- Inject the function `neural_network_inference` in the time-stepping loop
- A callback has access to all the variables of the simulation
- Each iteration `u_sgs` is used as forcing and then recalculated



<http://clima.caltech.edu>



*thanks*

[github.com/ClIMA/Oceananigans.jl](https://github.com/ClIMA/Oceananigans.jl)

