



# Waves of the future



## how modern software transforms ocean research & teaching

Navid Constantinou

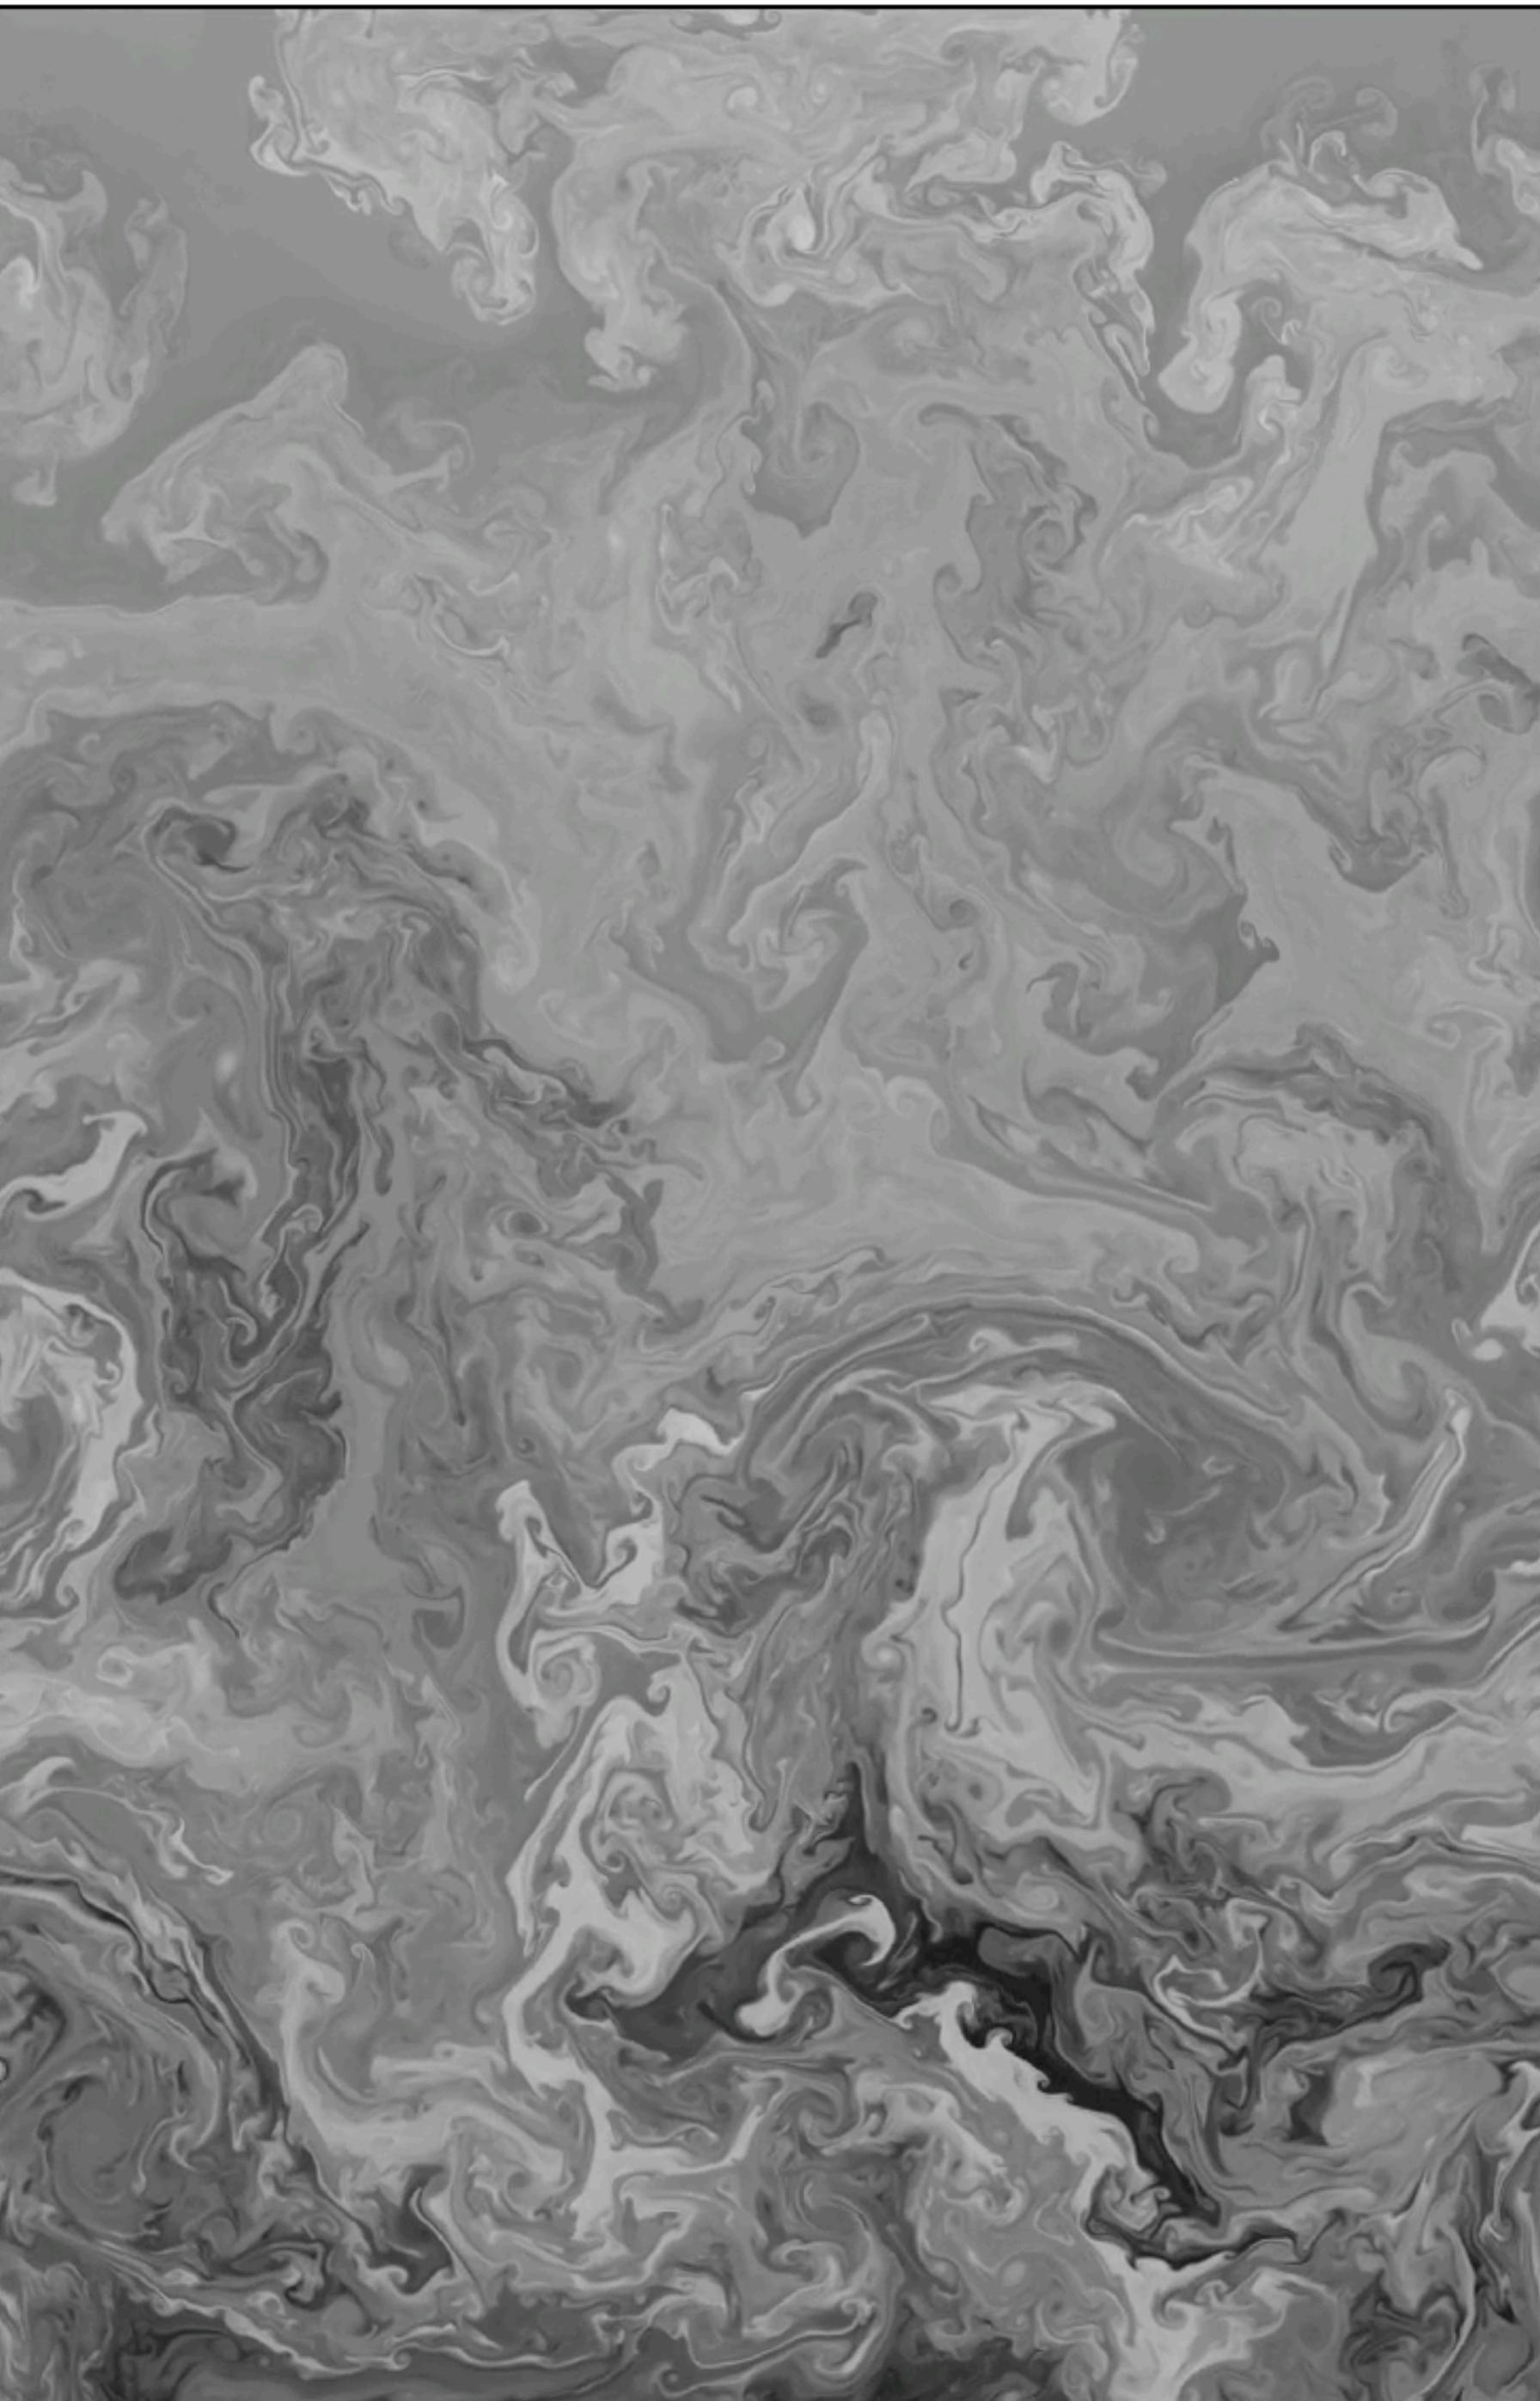


### Sean Haney 2025 Memorial Symposium

Feb 2025



stratified turbulence simulated with Oceananigans  
animation made using Greg Wagner's TurbulentImages.jl package



why climate is tad difficult problem? ☀️☁️earth🌀

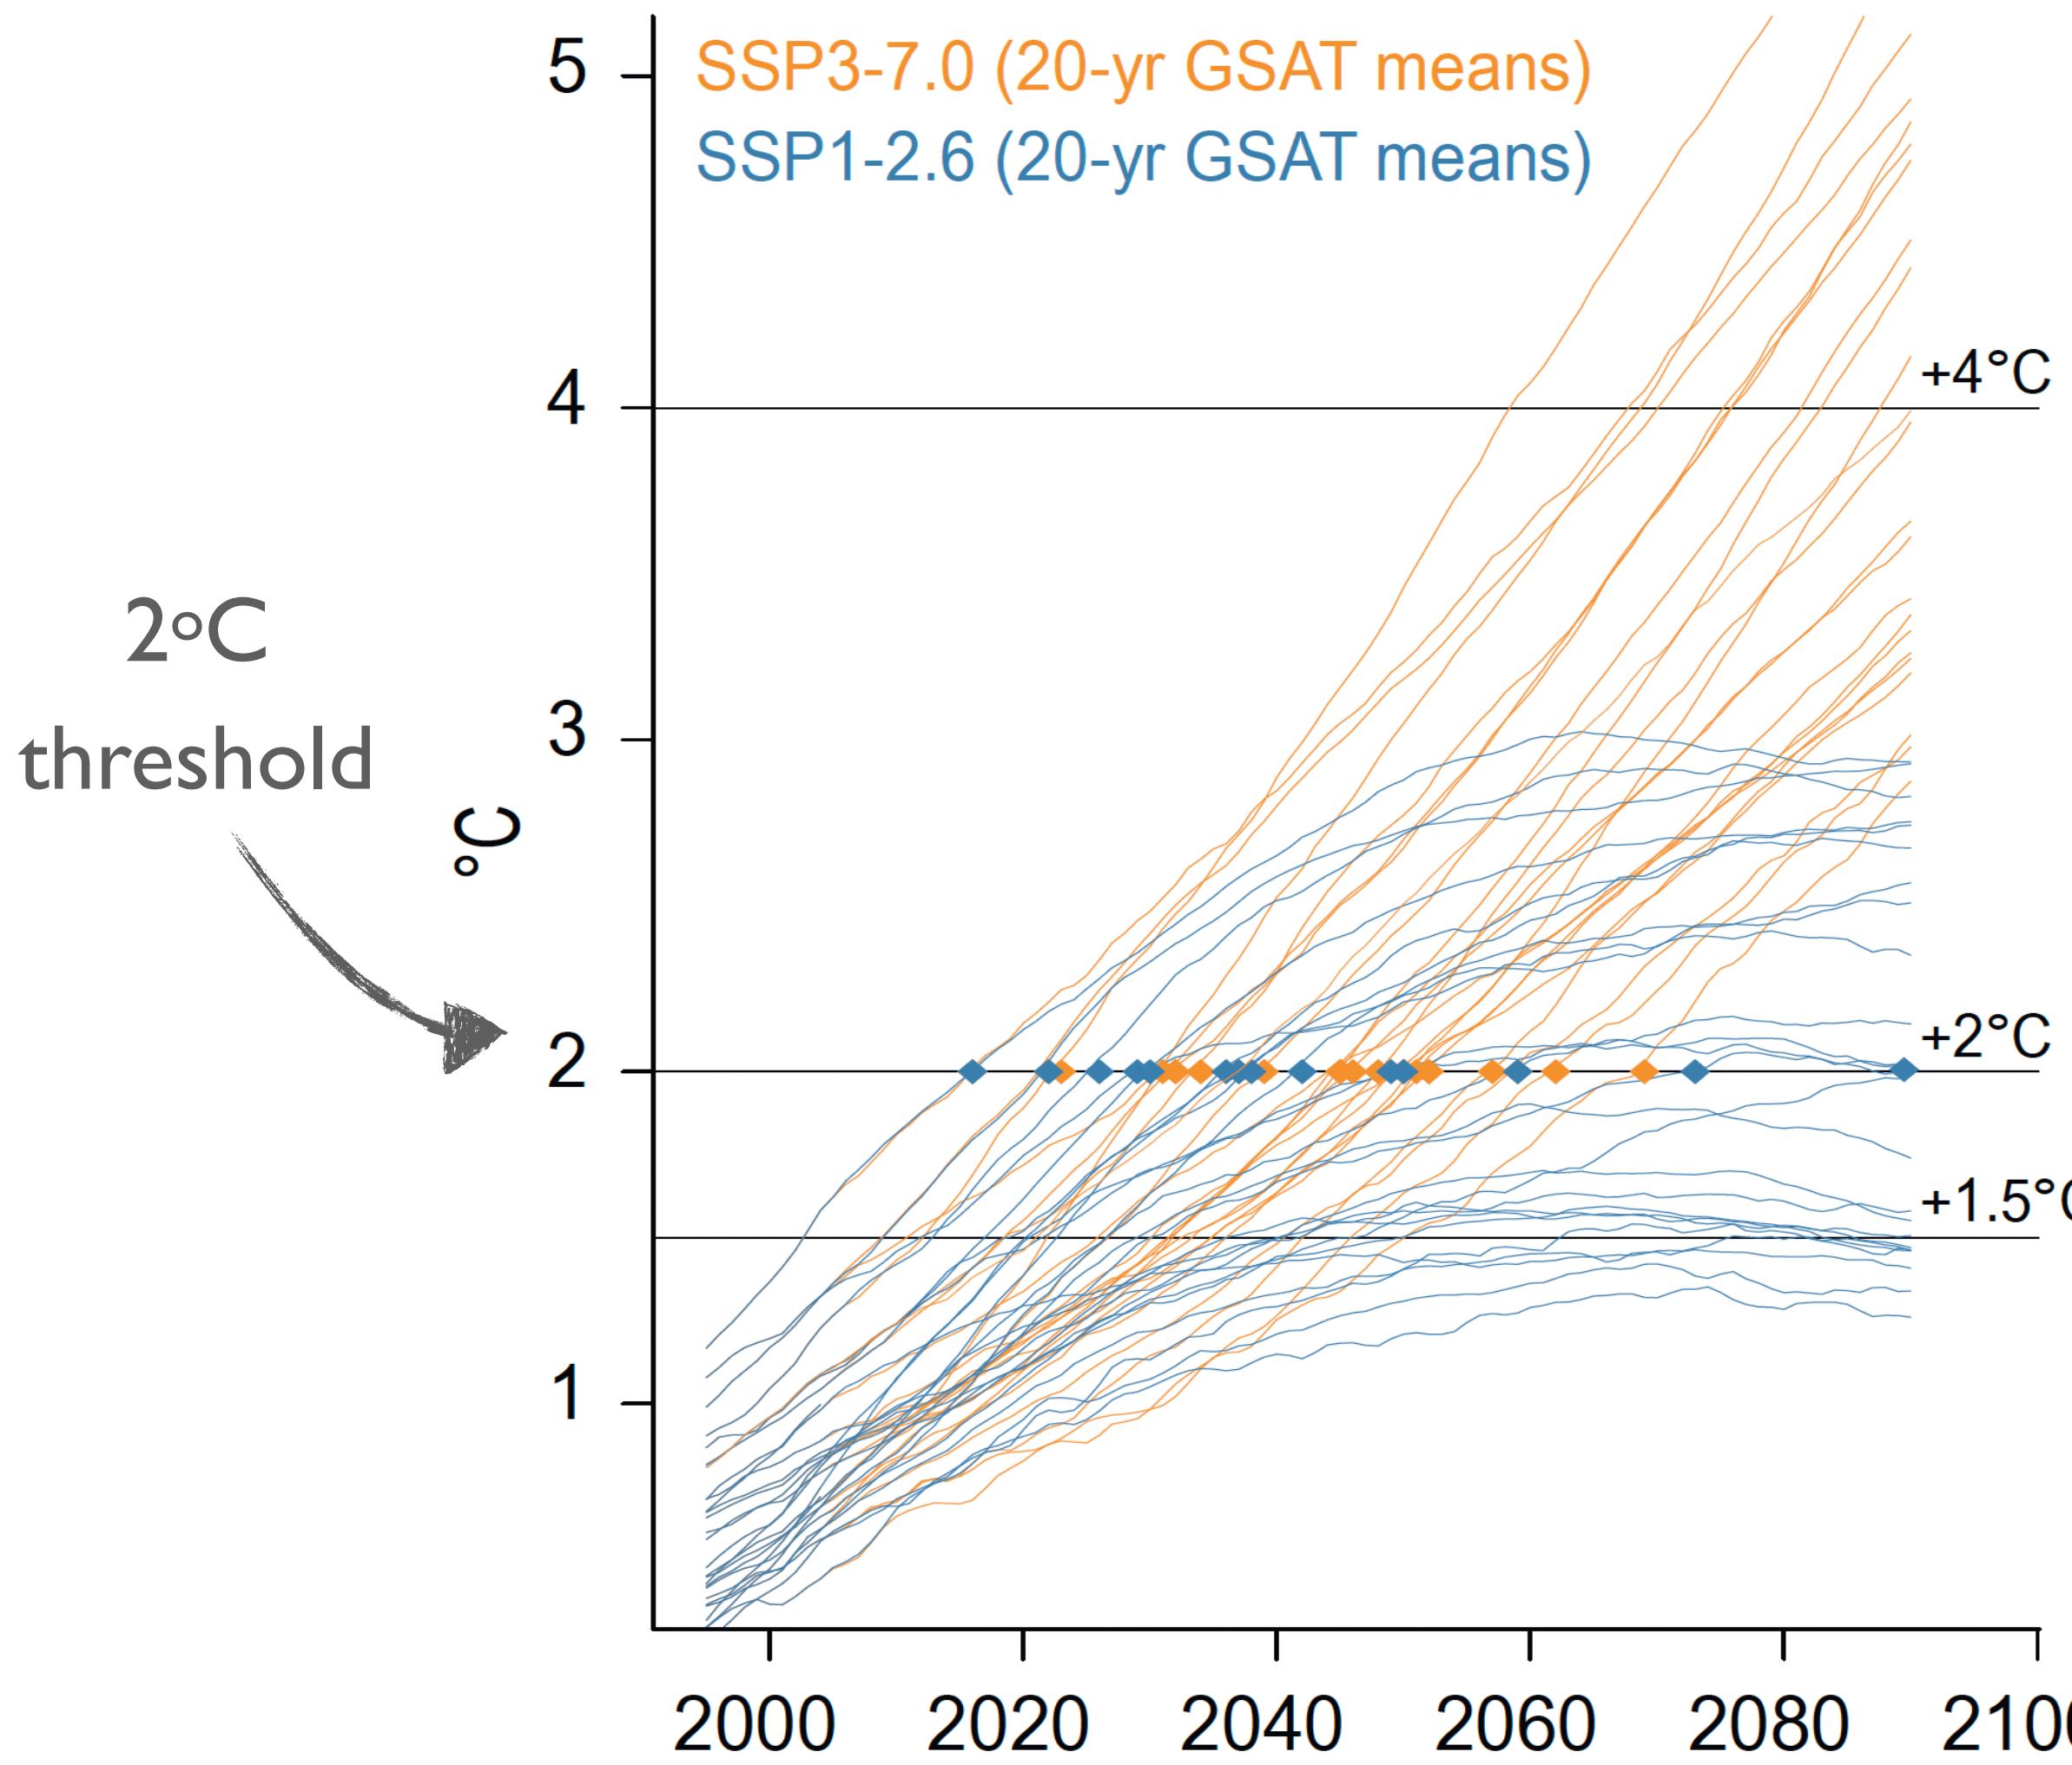
climate models 🚧💾

modern software 🛸  
& its impact on our work flow in  
research 🧫 and teaching 🎓

# why climate is tad difficult problem? ☀️ 🌎 🌪️

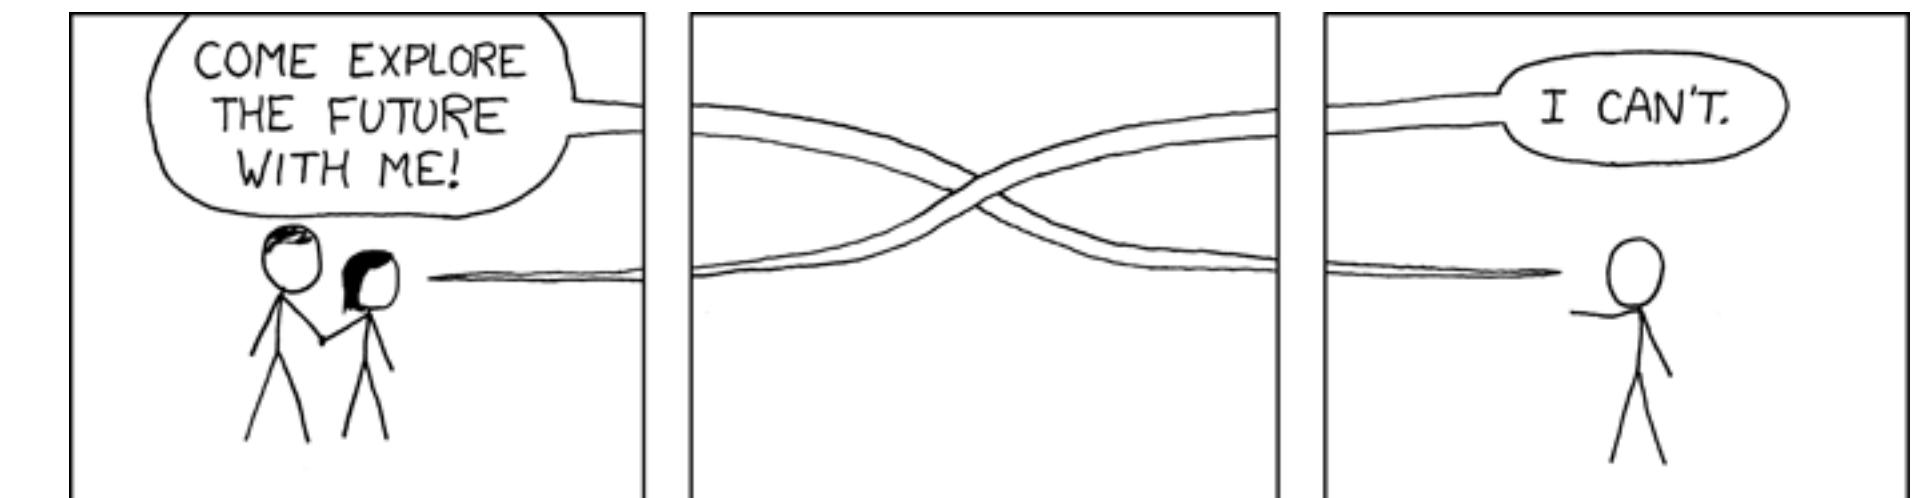
climate 🌎 is what you expect,  
weather is what you get ☀️RAIN

# **predictions of future warming are *uncertain***

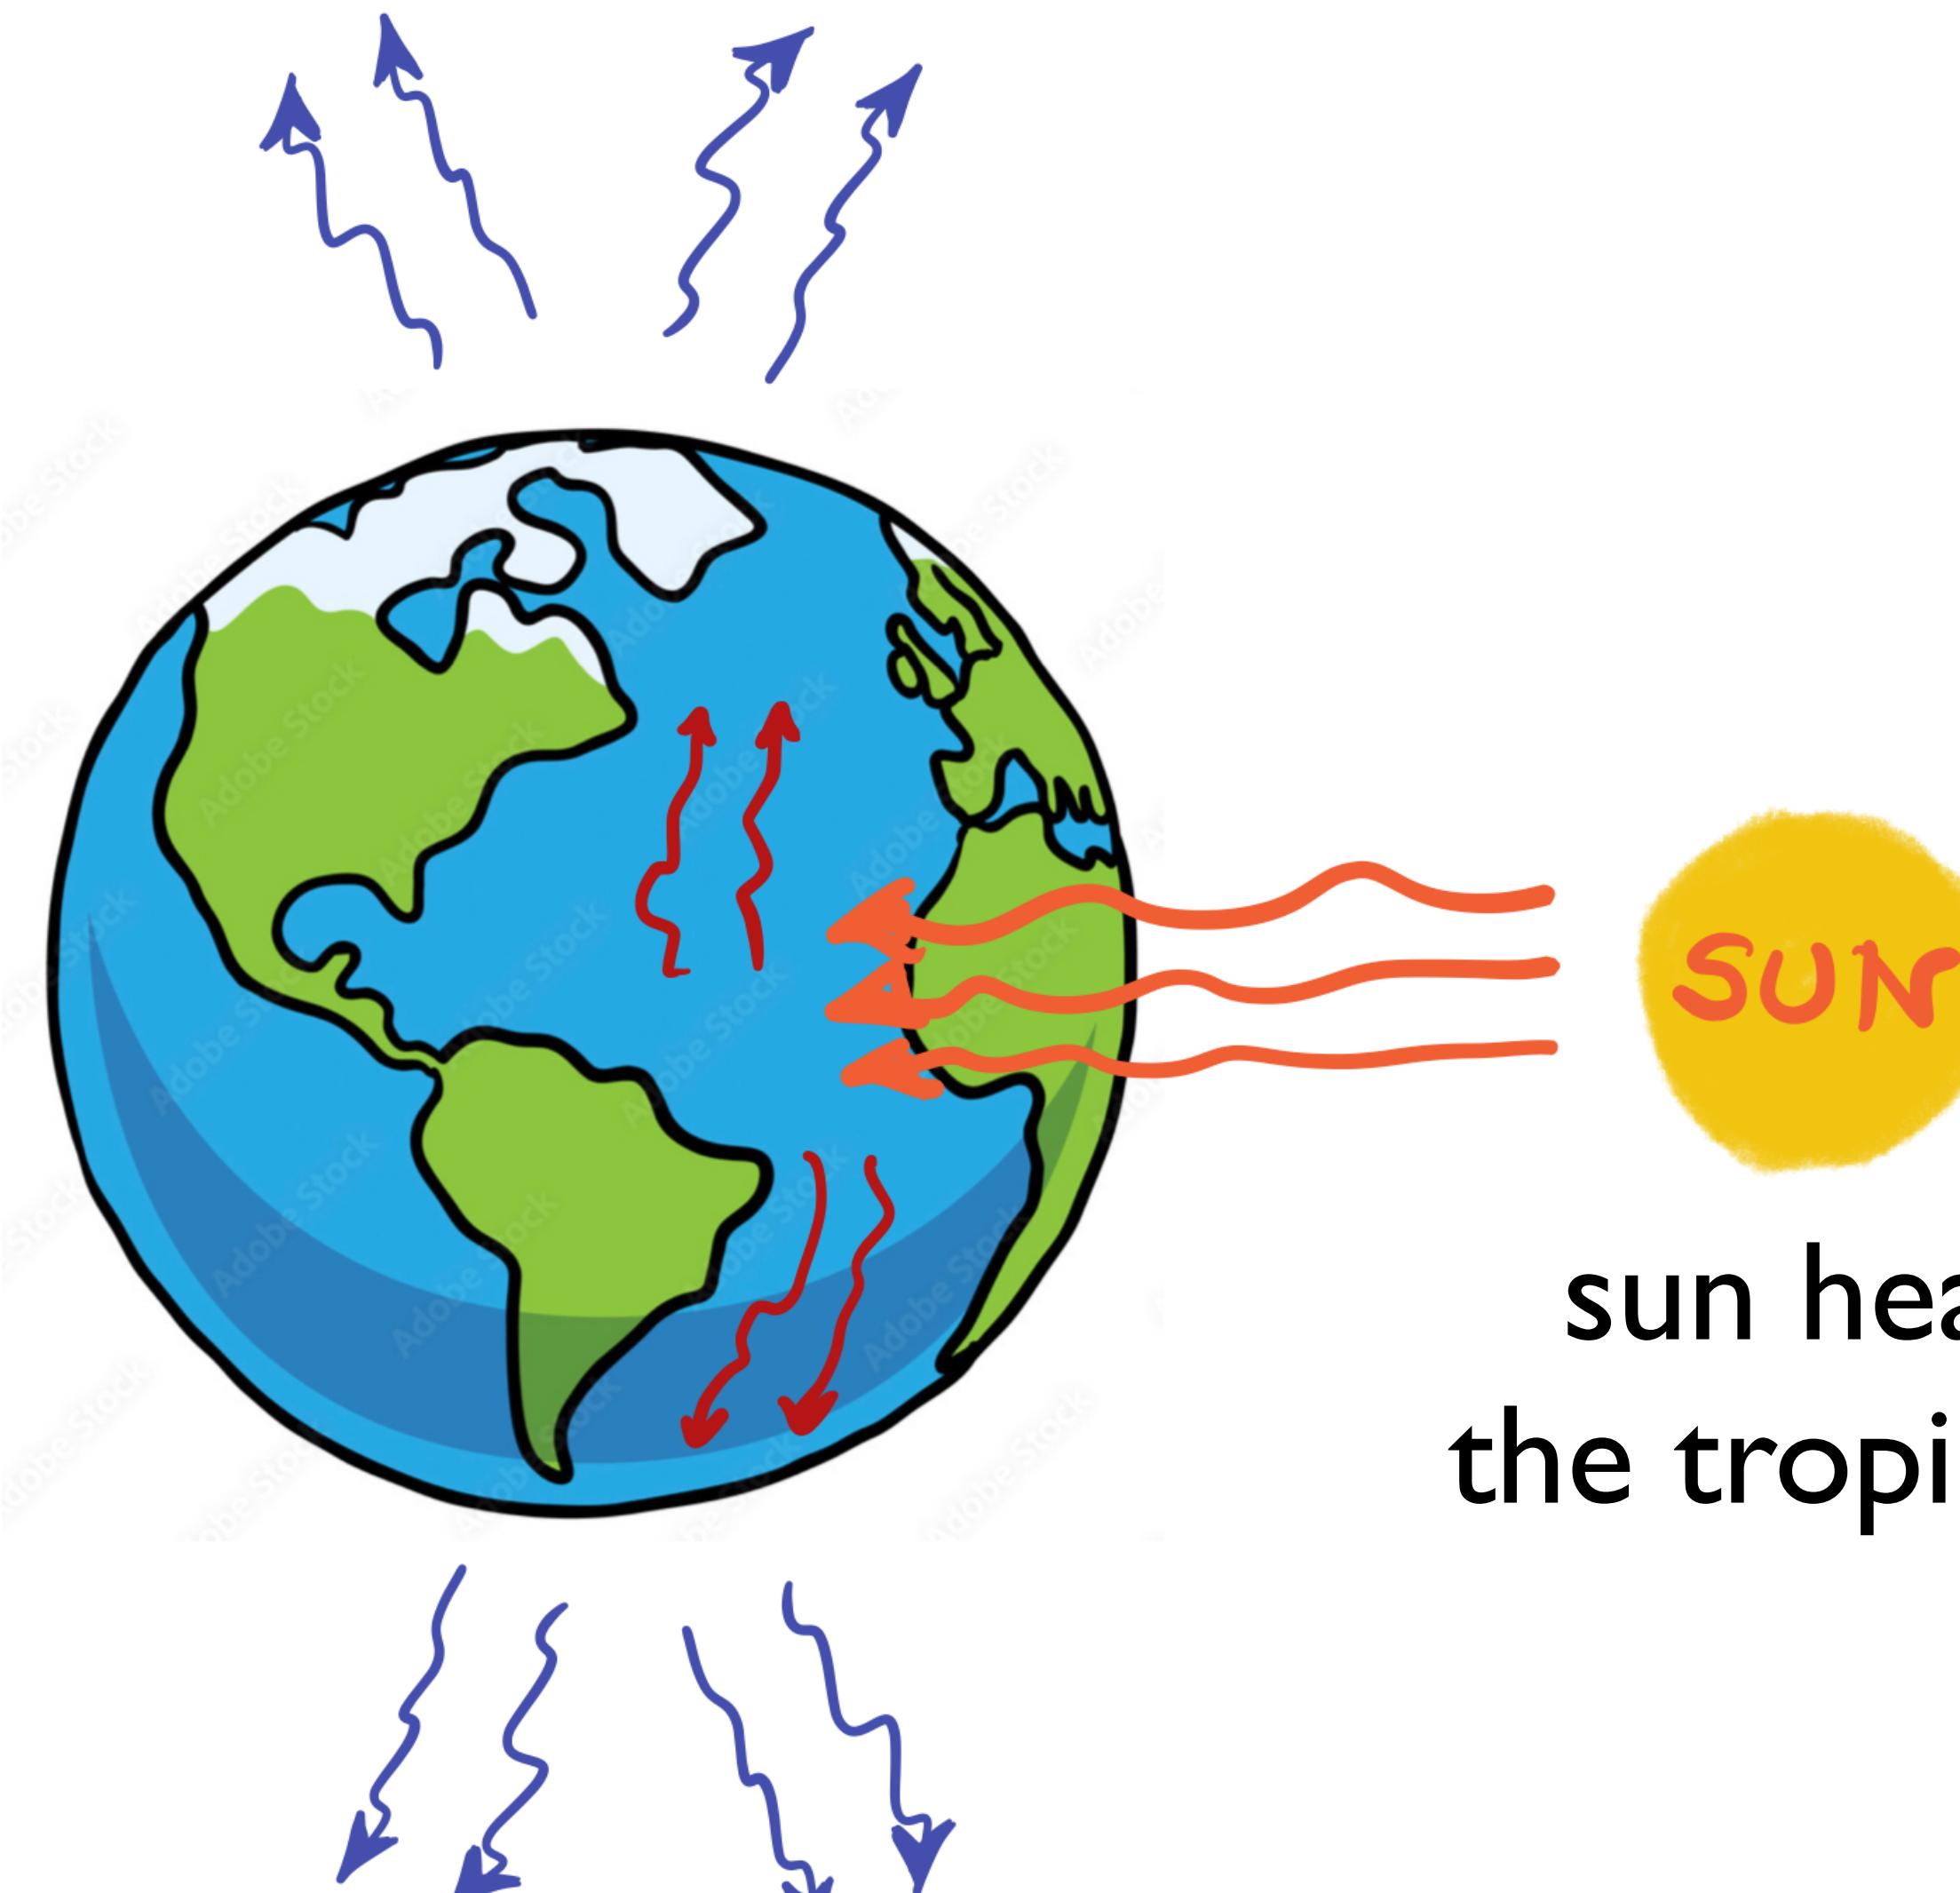


High-emissions (SSP3-7.0)  
and  
low-emissions (SSP1-2.6)  
scenarios in CMIP6

IPCC AR 6 Working Group I

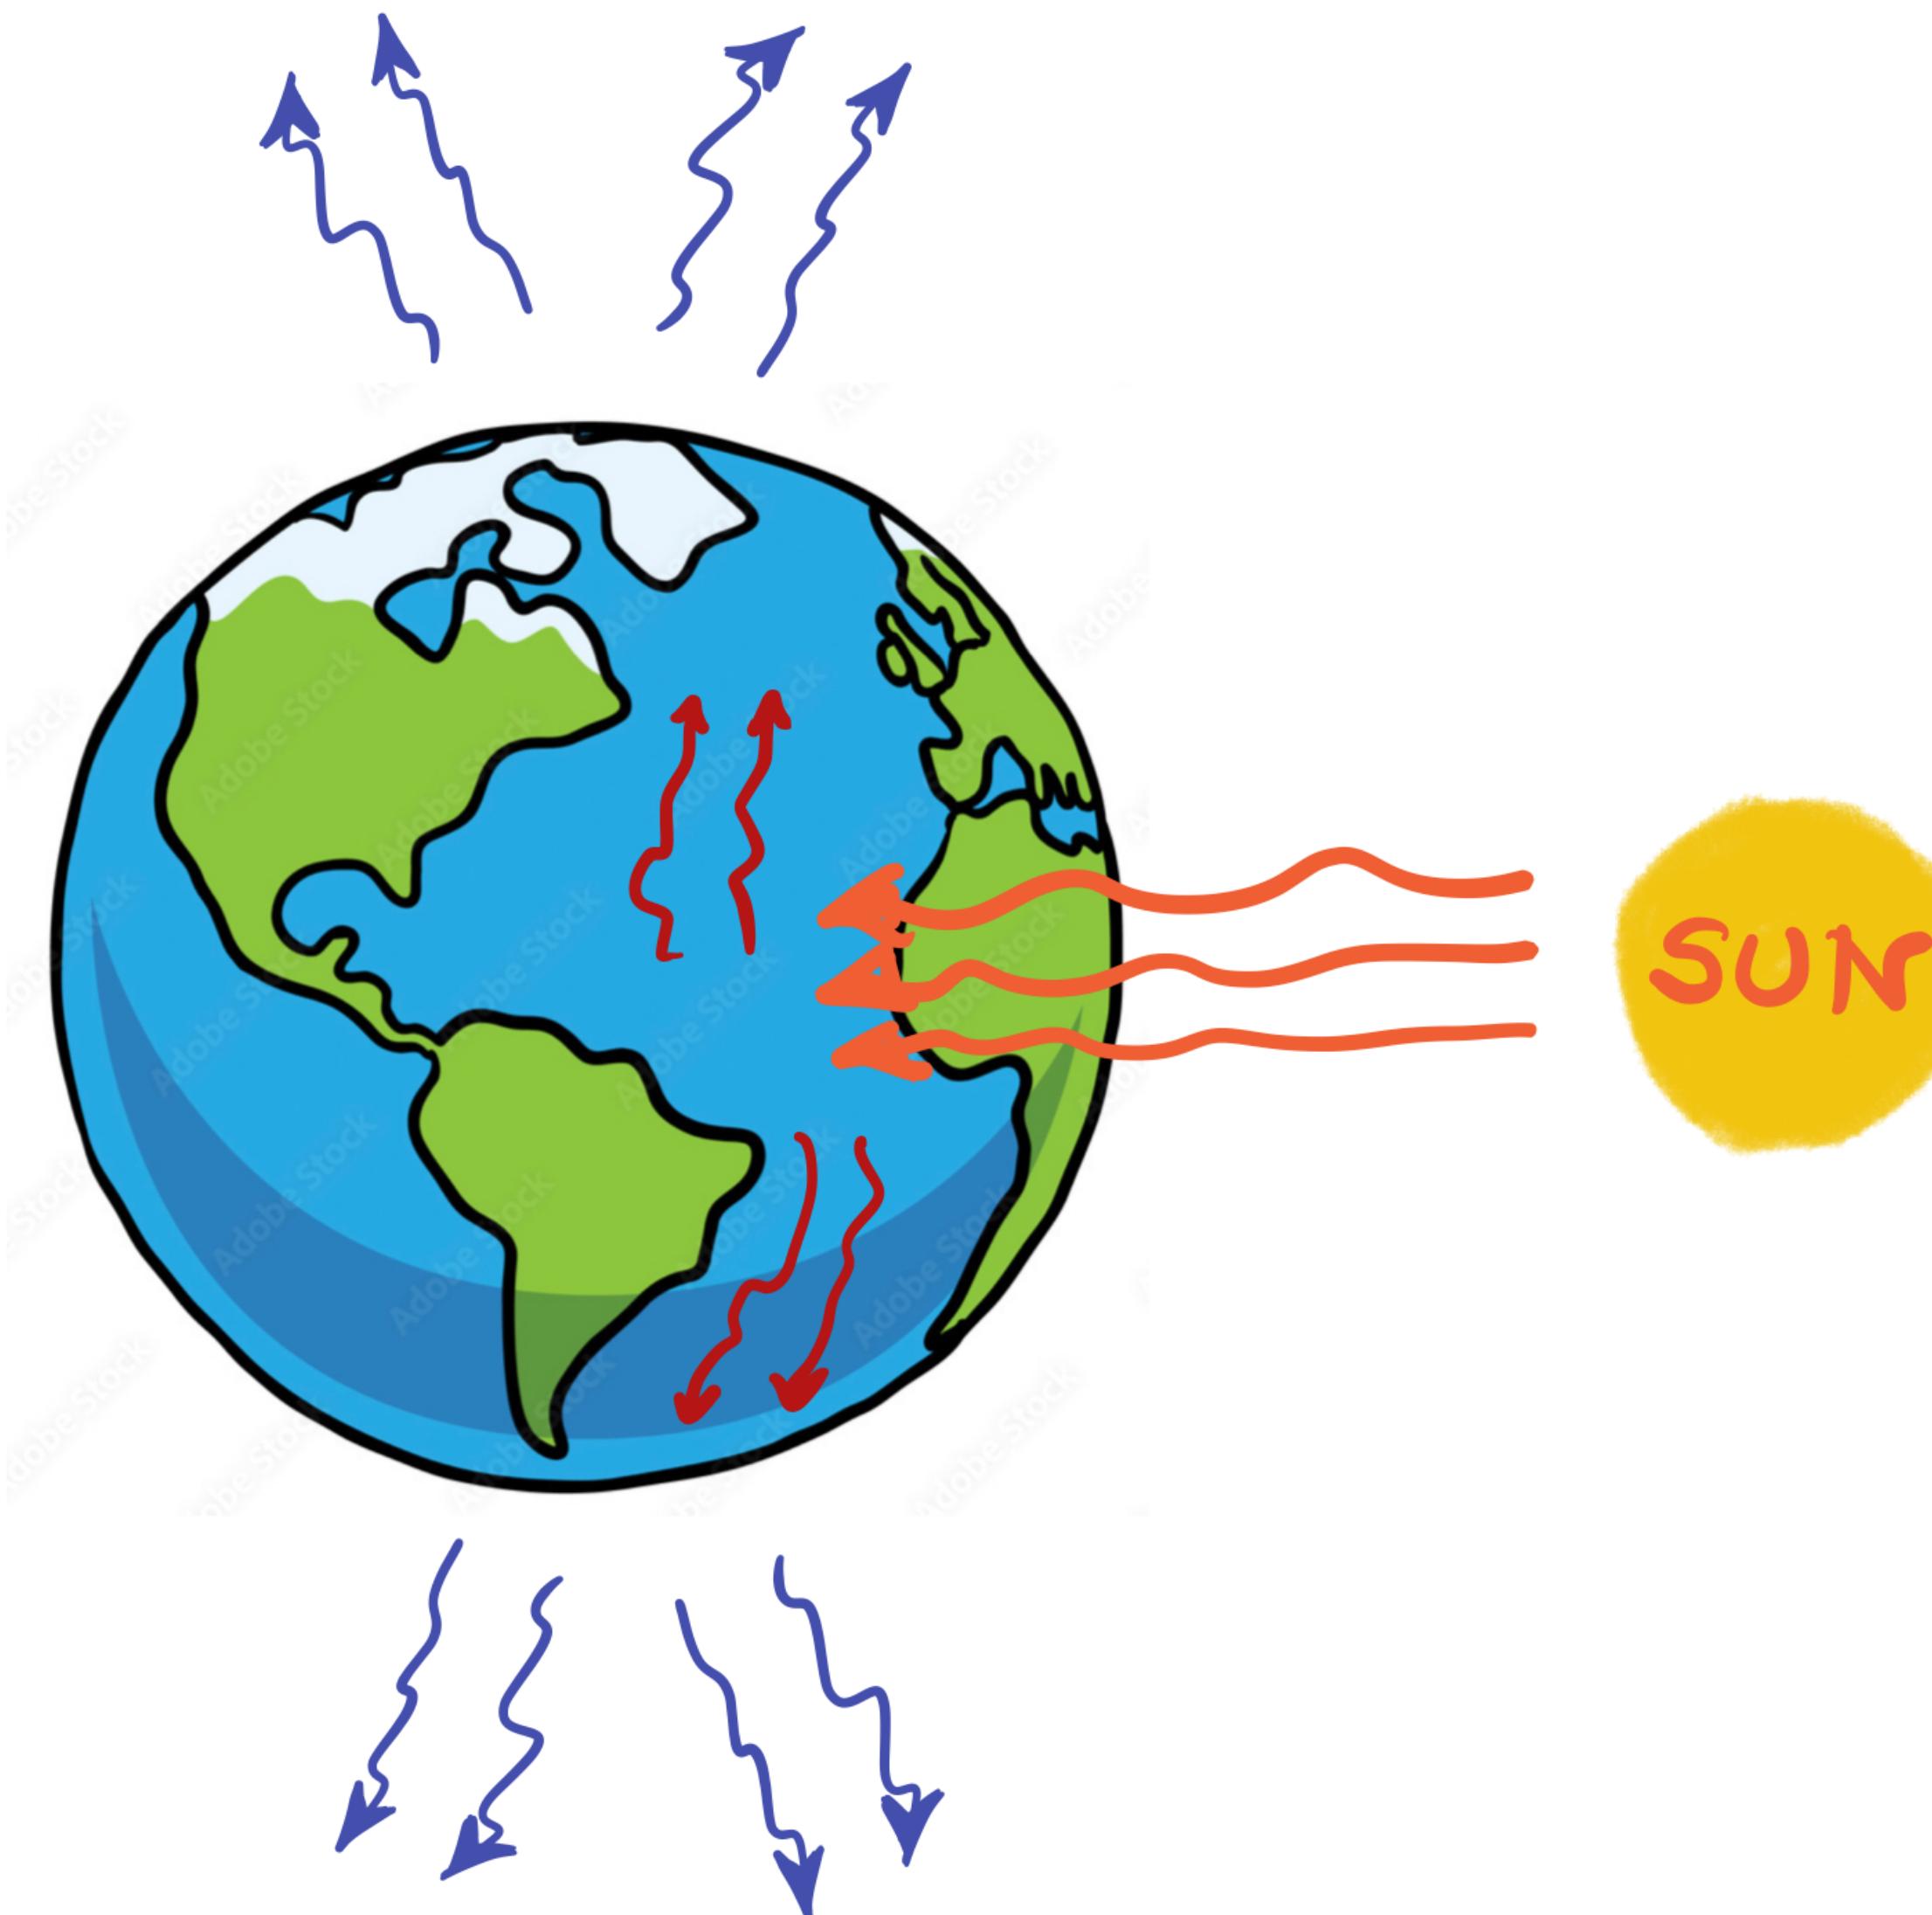


# what sets the ocean (and atmosphere) in motion



...ocean & atmosphere   
carry some of the heat to the poles

# 'steady' climate / climate in equilibrium

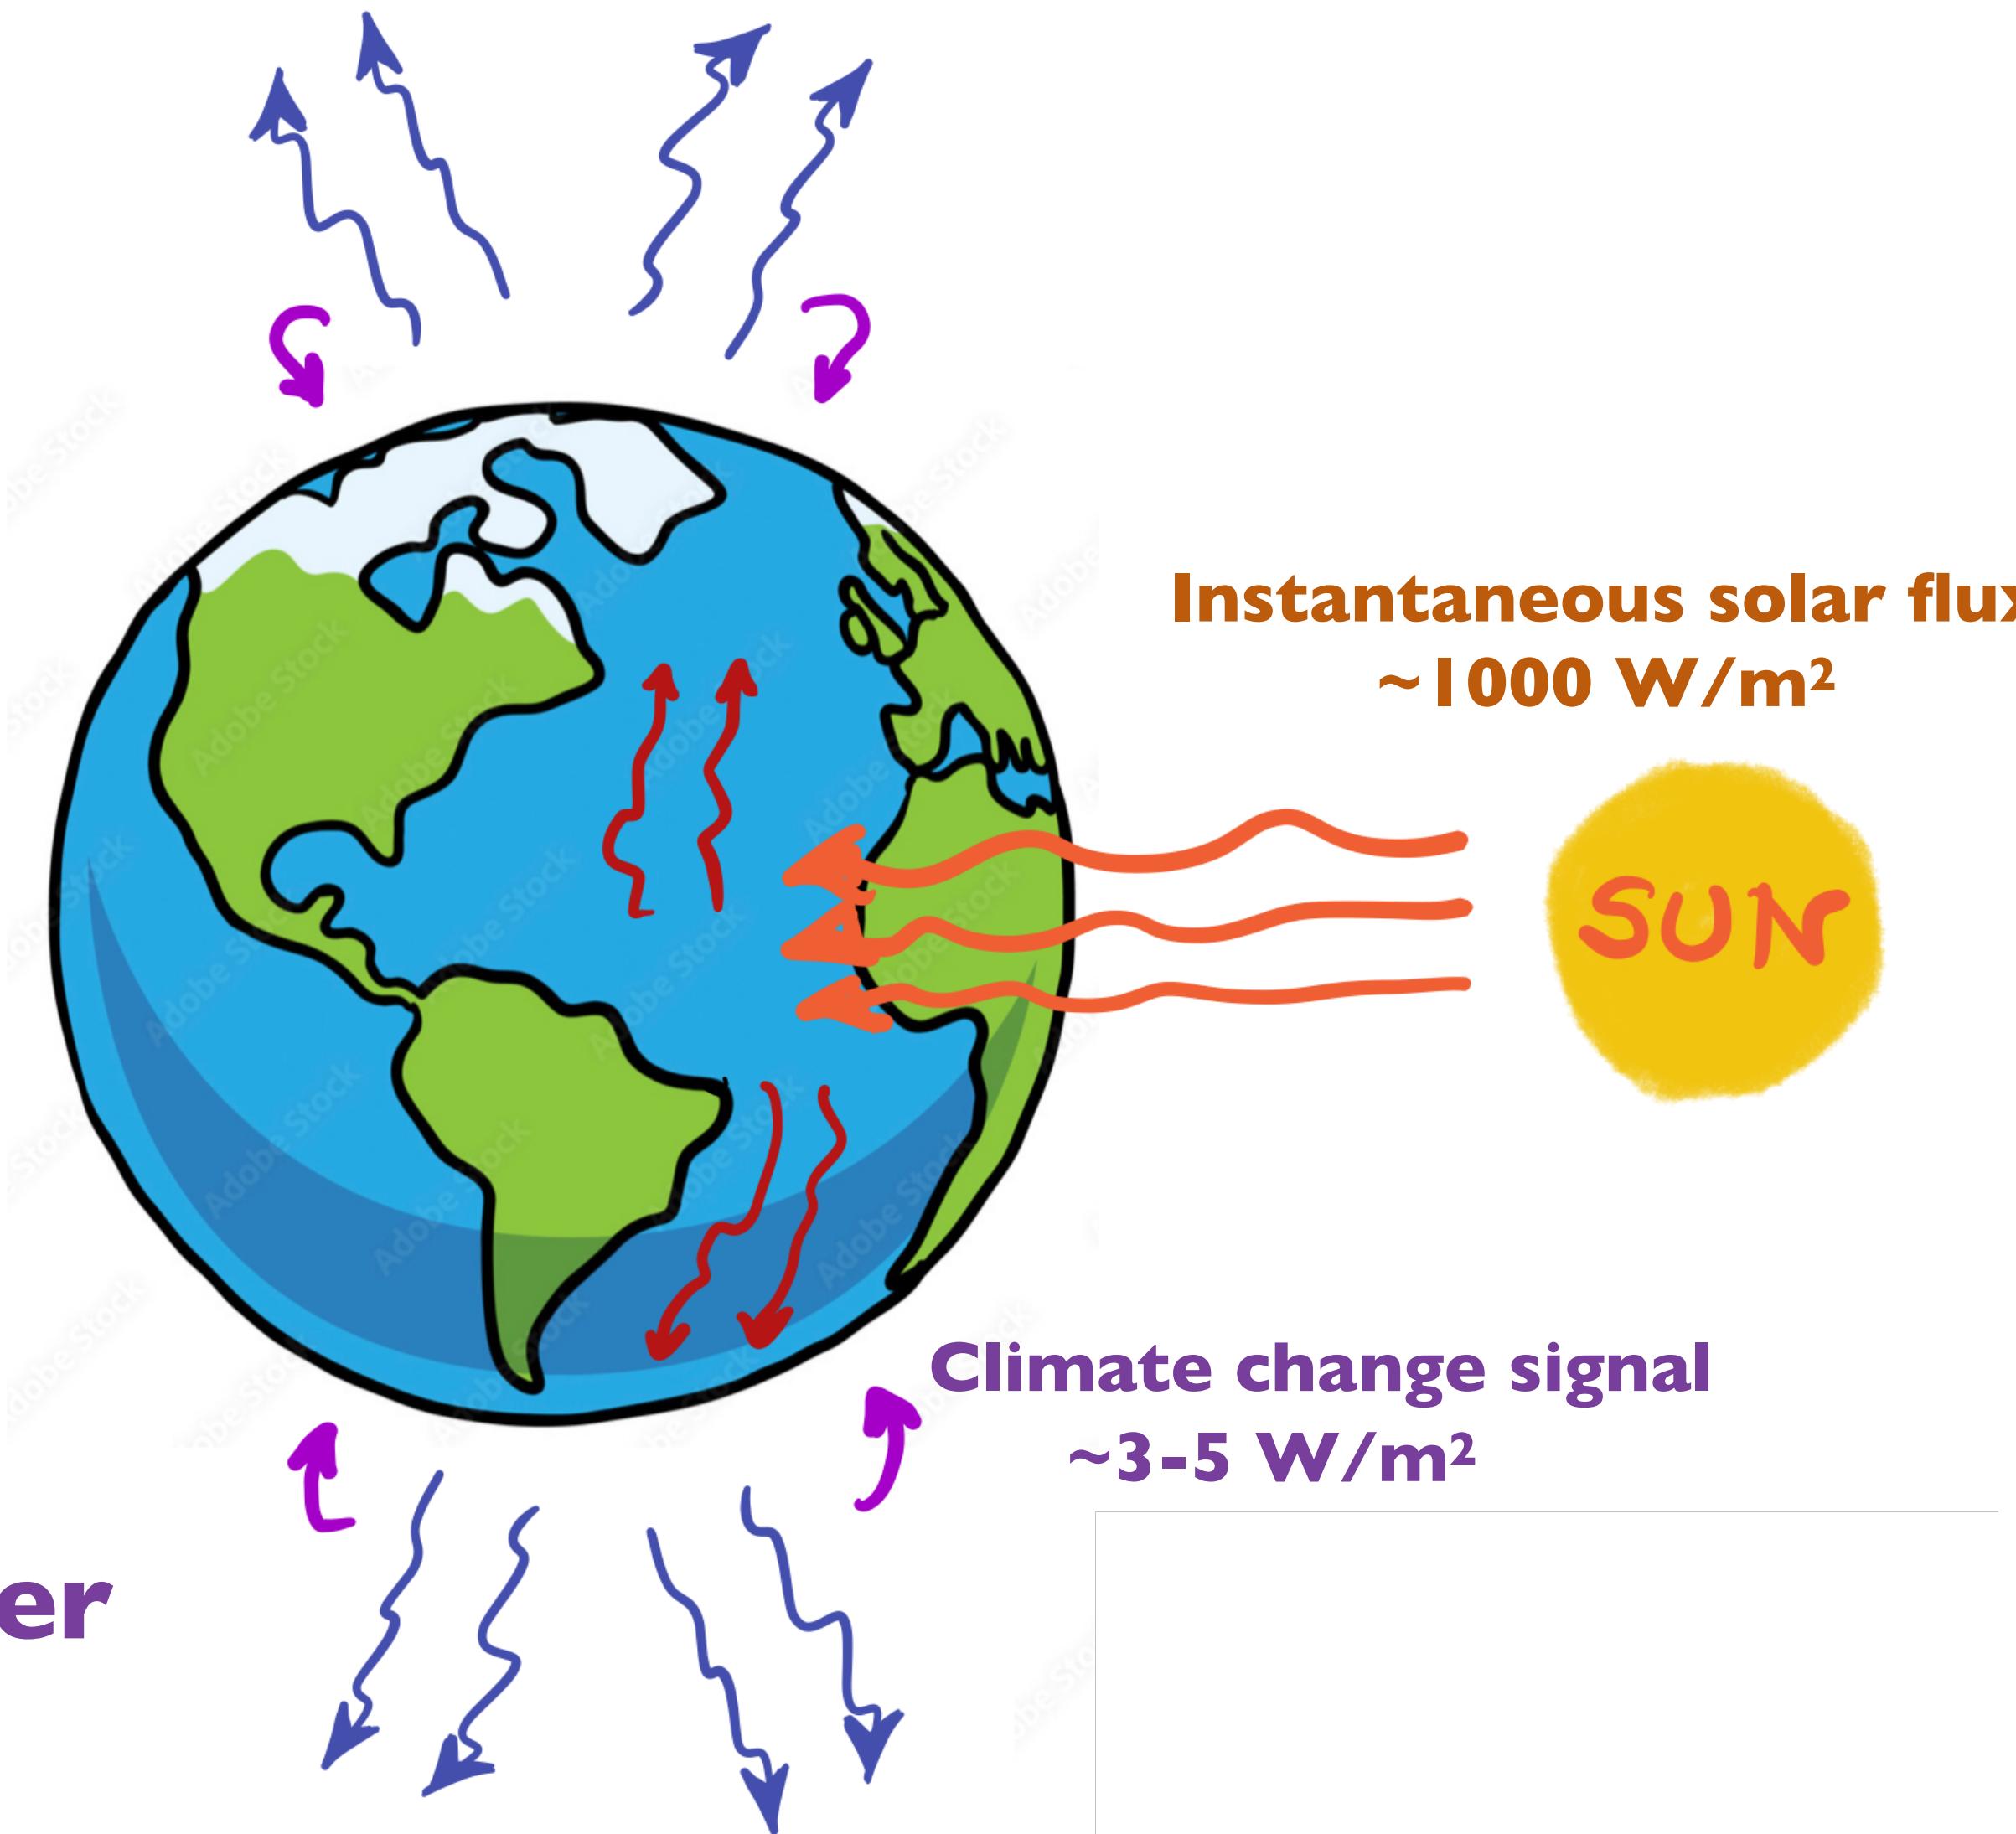


**energy in = energy out**

**energy in** is exactly equal to **energy out**

# a changing climate

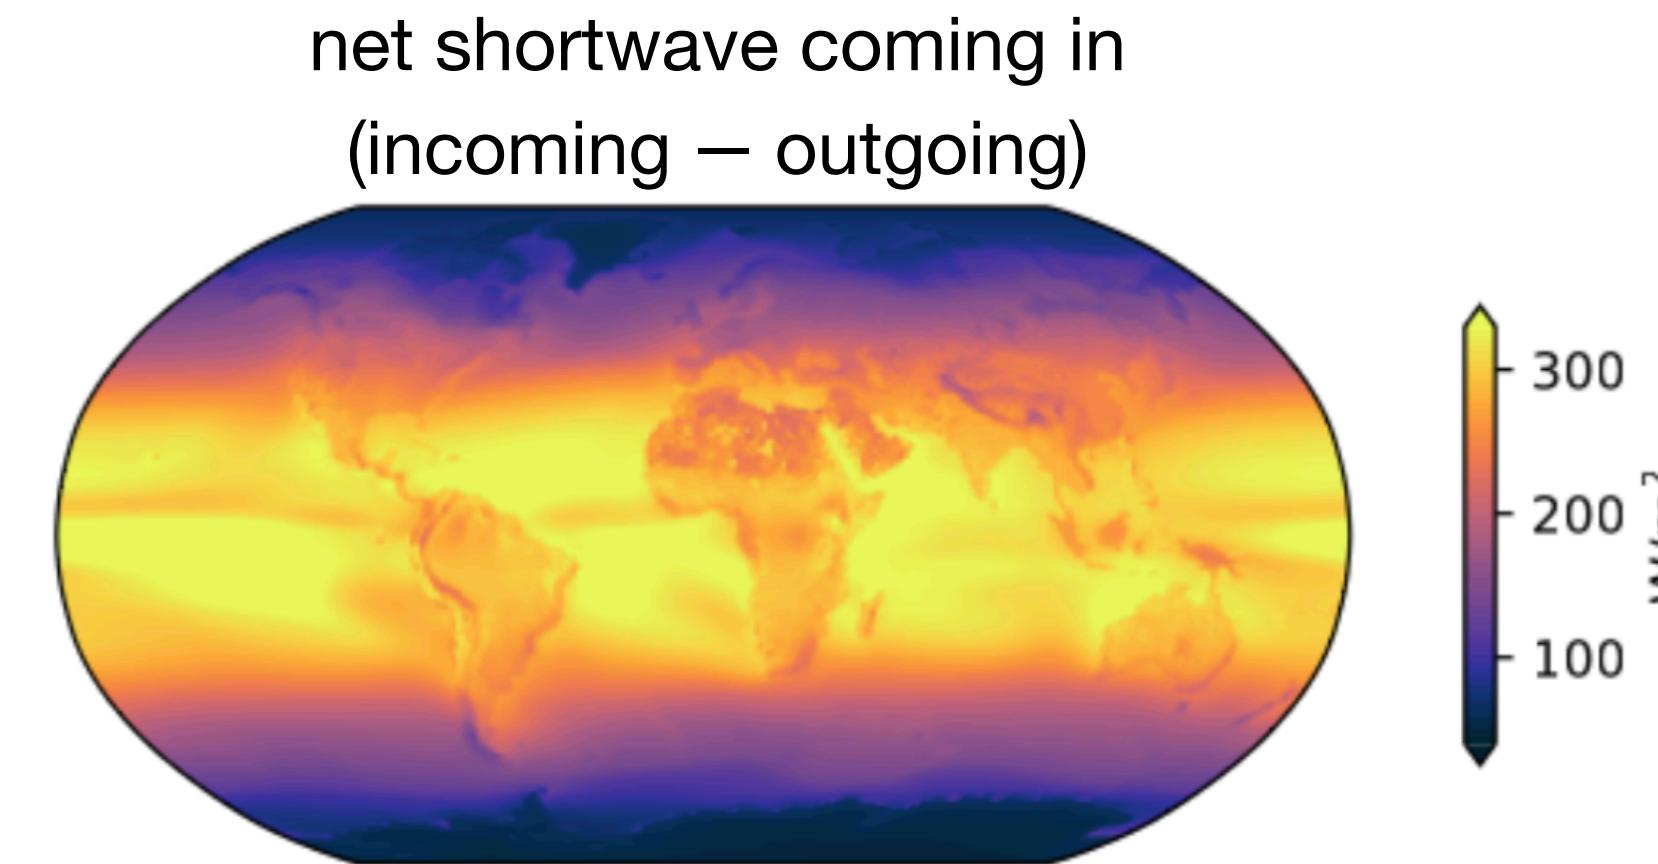
**energy in = energy out**  
+ a little  
bit leftover



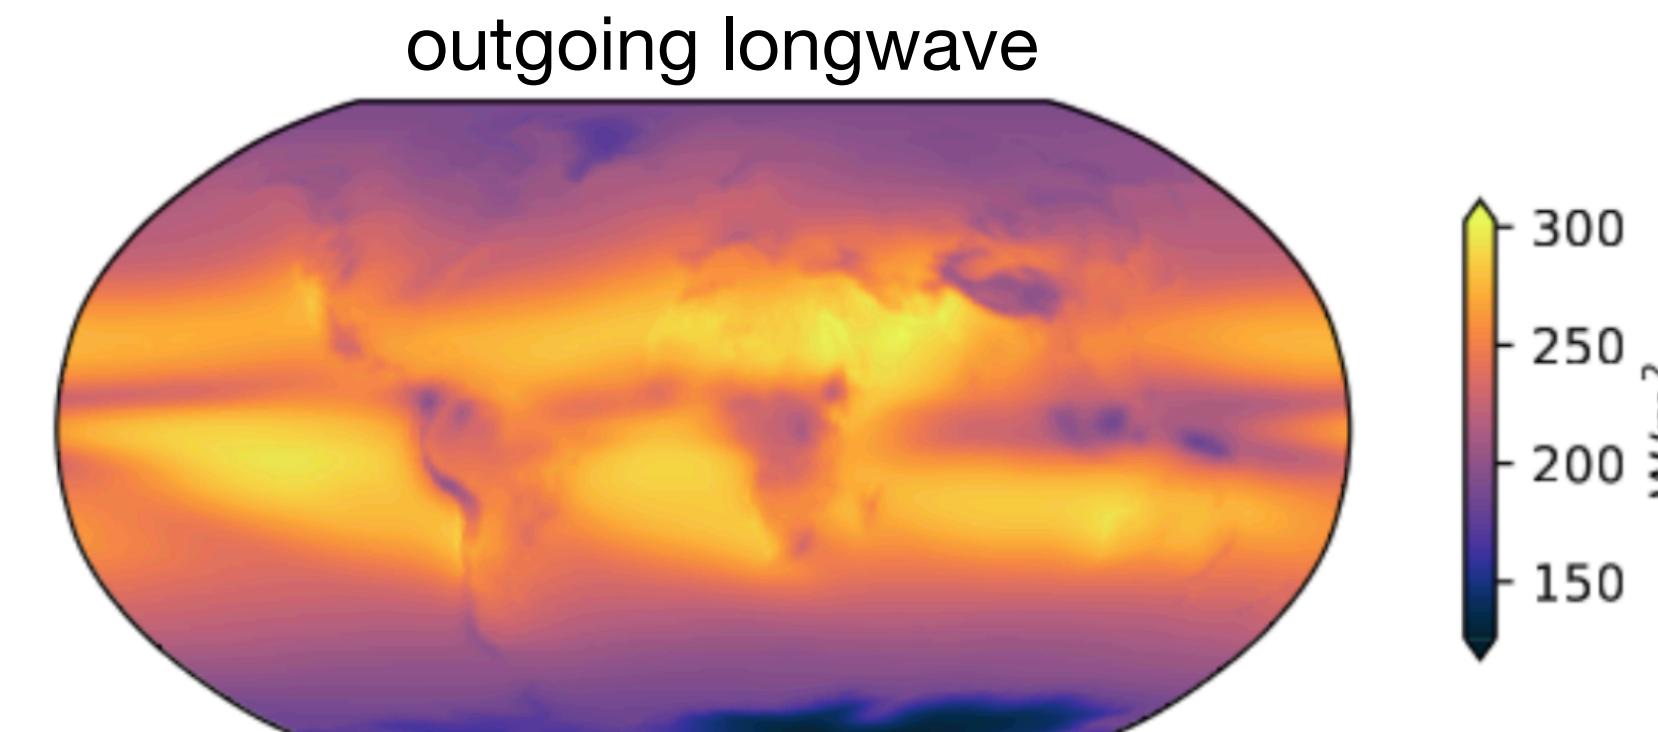
**energy in** is not quite equal to **energy out**

# if you like being a bit more quantitative..

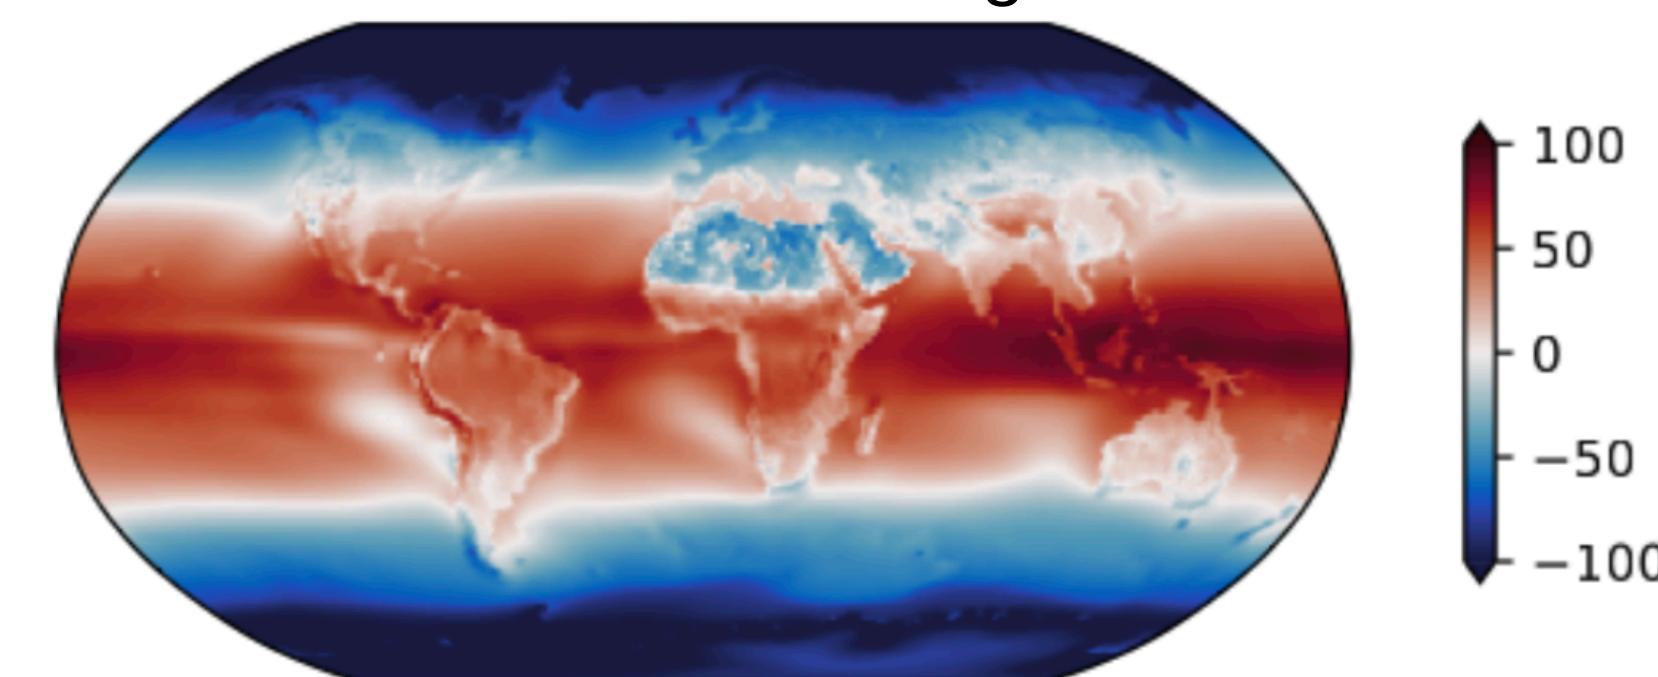
**energy in**



**energy out**

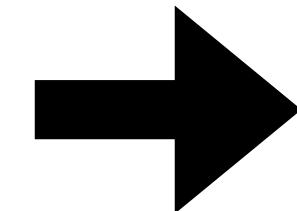


**their difference**



[CESM pre-industrial control  
climate simulation for CMIP6]

# why is climate prediction so hard?

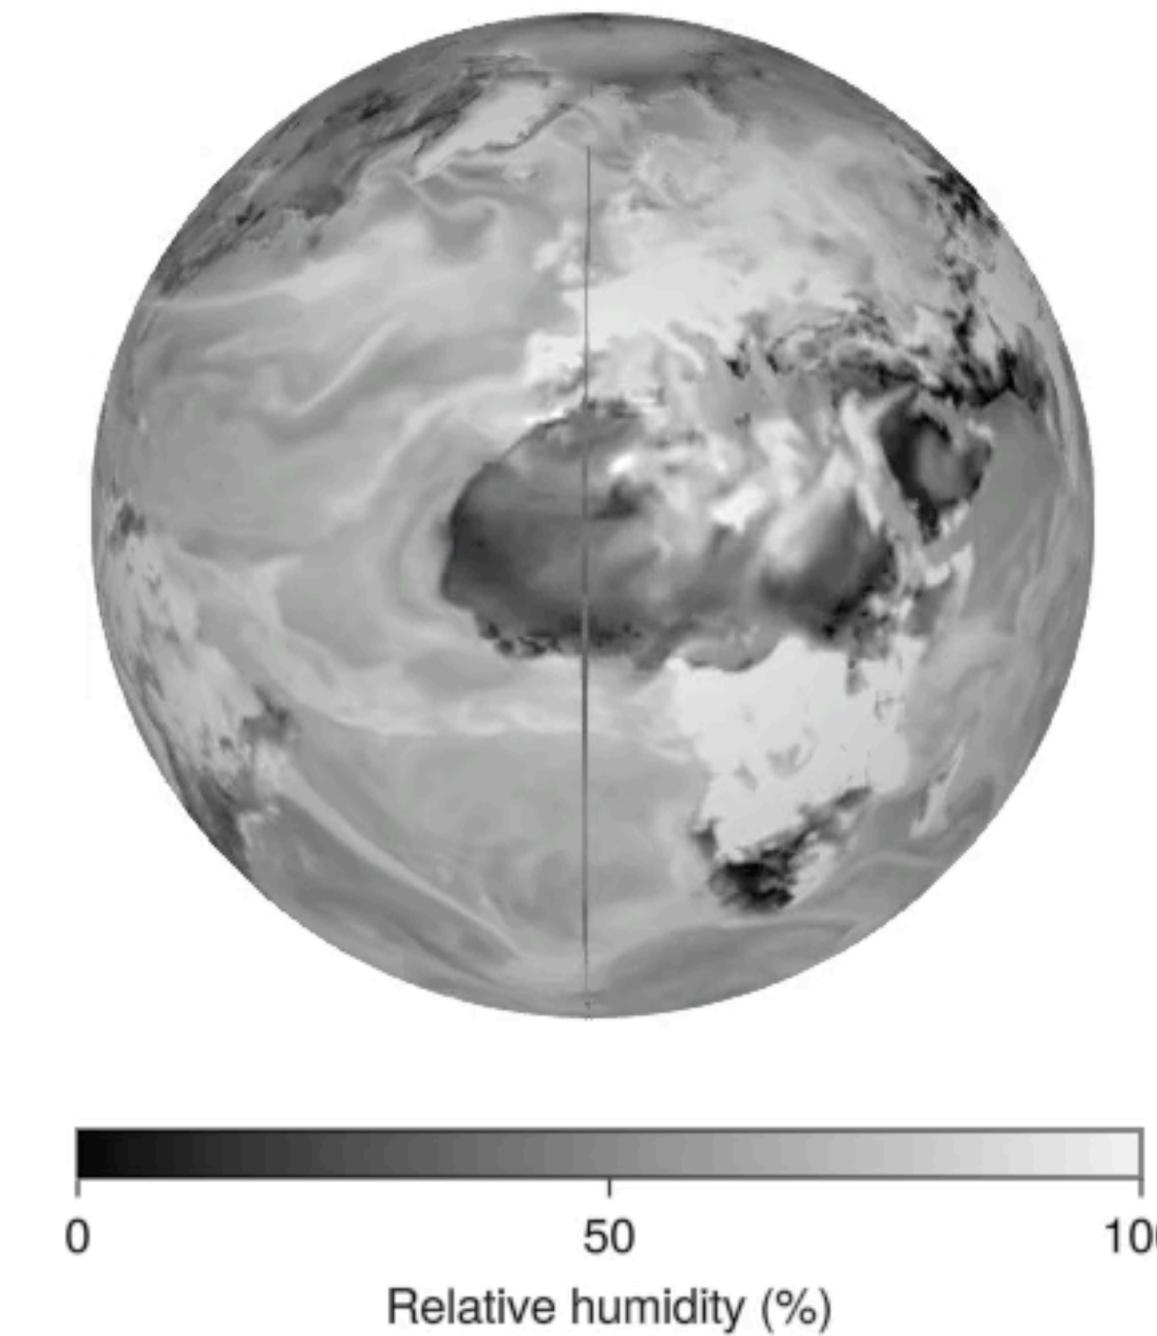
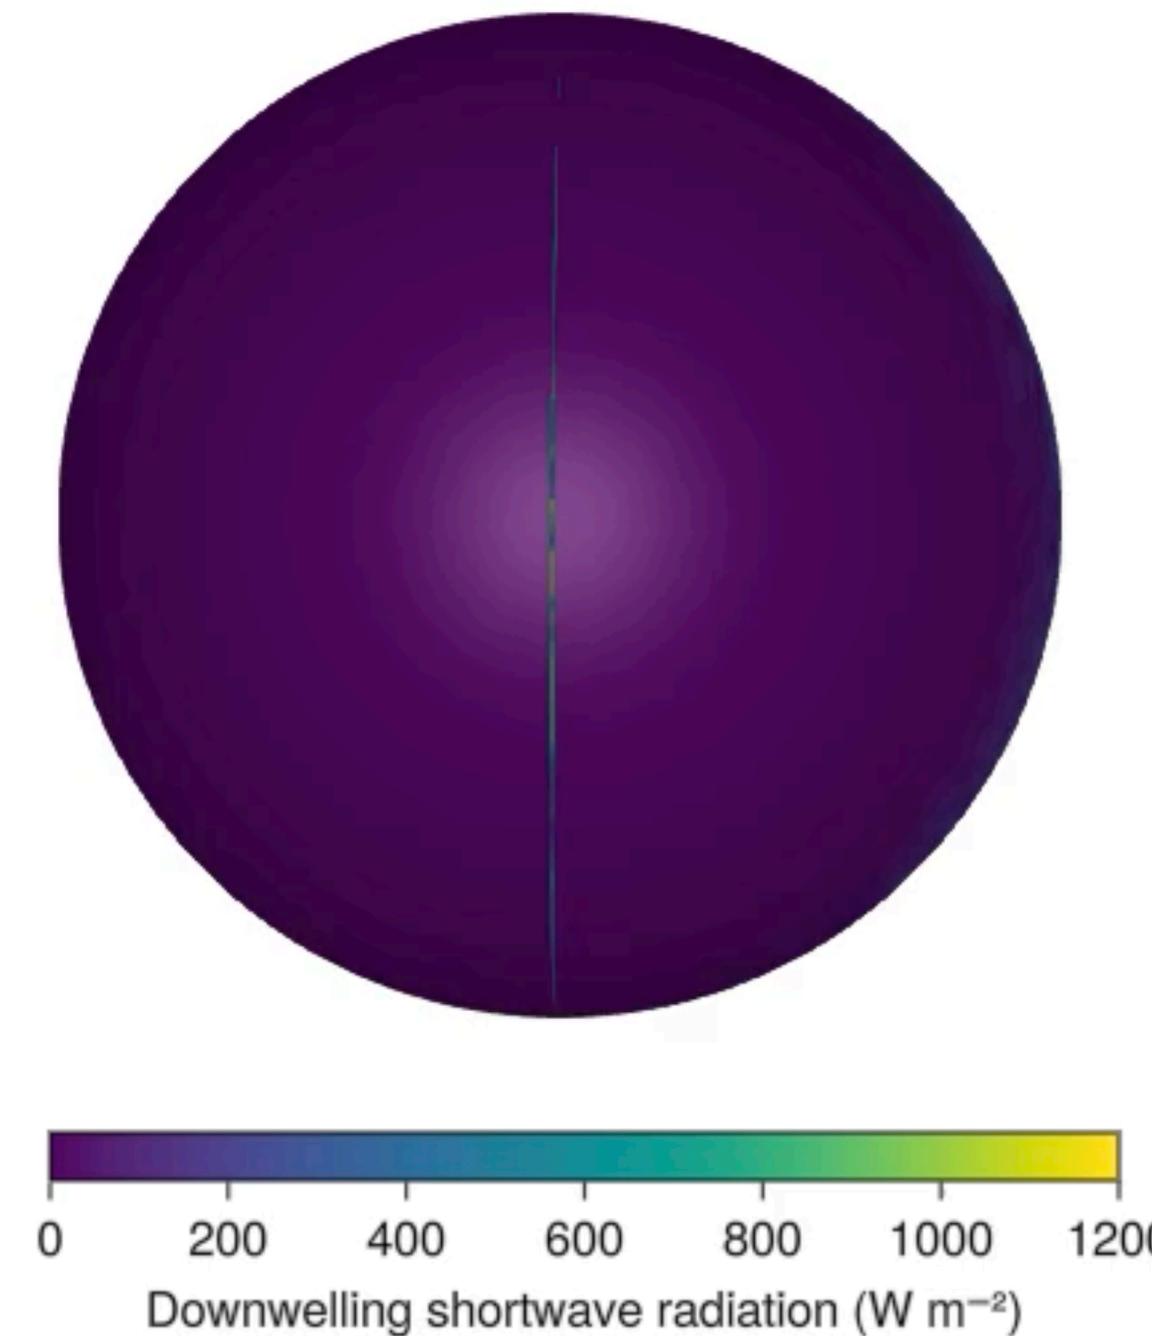


The Earth system is **strongly** forced  
but very close to equilibrium

JRA55 forcing on year-day 0.0

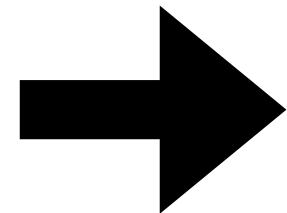
Climate change signal: ~3-5 W/m<sup>2</sup>  
Instantaneous solar flux: ~1300 W/m<sup>2</sup>

**“energy in”**



Earth system models redistribute these fluxes across all components  
(atmosphere, ocean, land, ice, trees, etc)

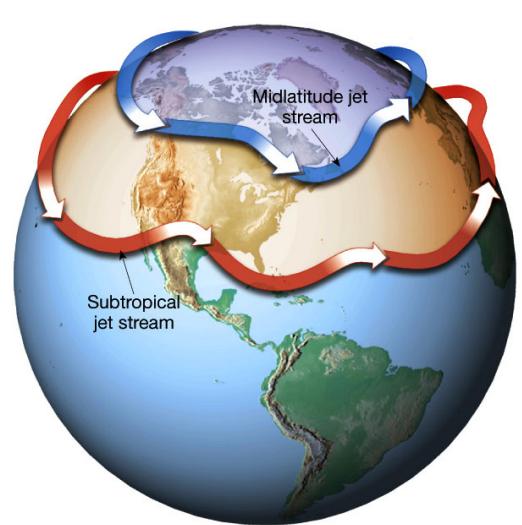
# why is climate prediction so hard?



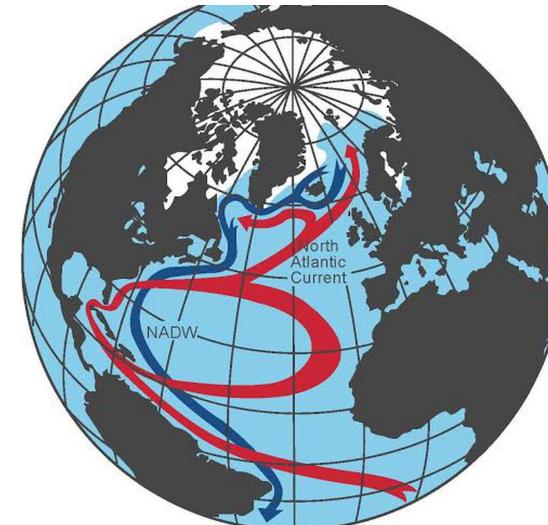
Interconnections among many processes

Some things we can't resolve..!

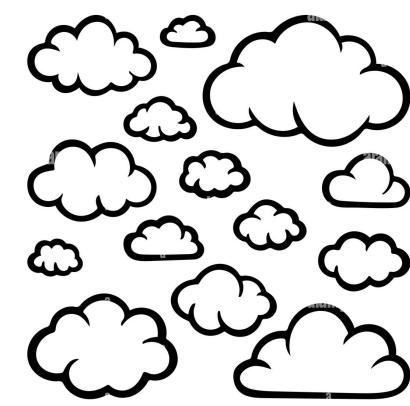
Other things we don't even know the equations!



atmosphere



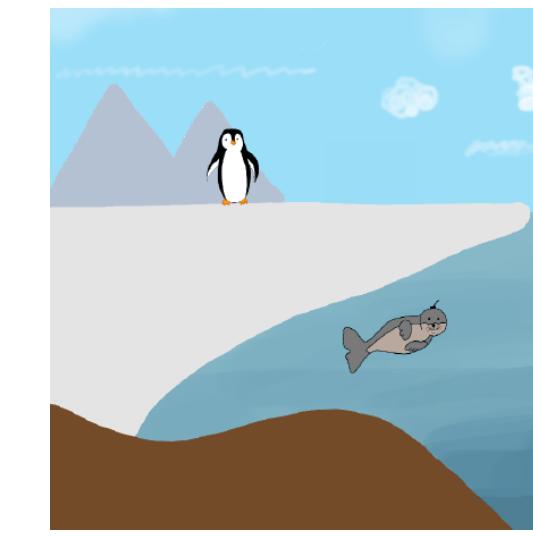
ocean



clouds



trees



ice-ocean  
interactions

[made by Claire Yung]



air-sea exchanges

...

many more  
processes

Approximate models for many processes (parameterizations)  
dominate the uncertainty of climate projections



# ocean: processes at many scales interact

turbulence

~1-100 m  
~minutes

eddies

~200 km  
~months

currents

~1000 km  
~years-decades

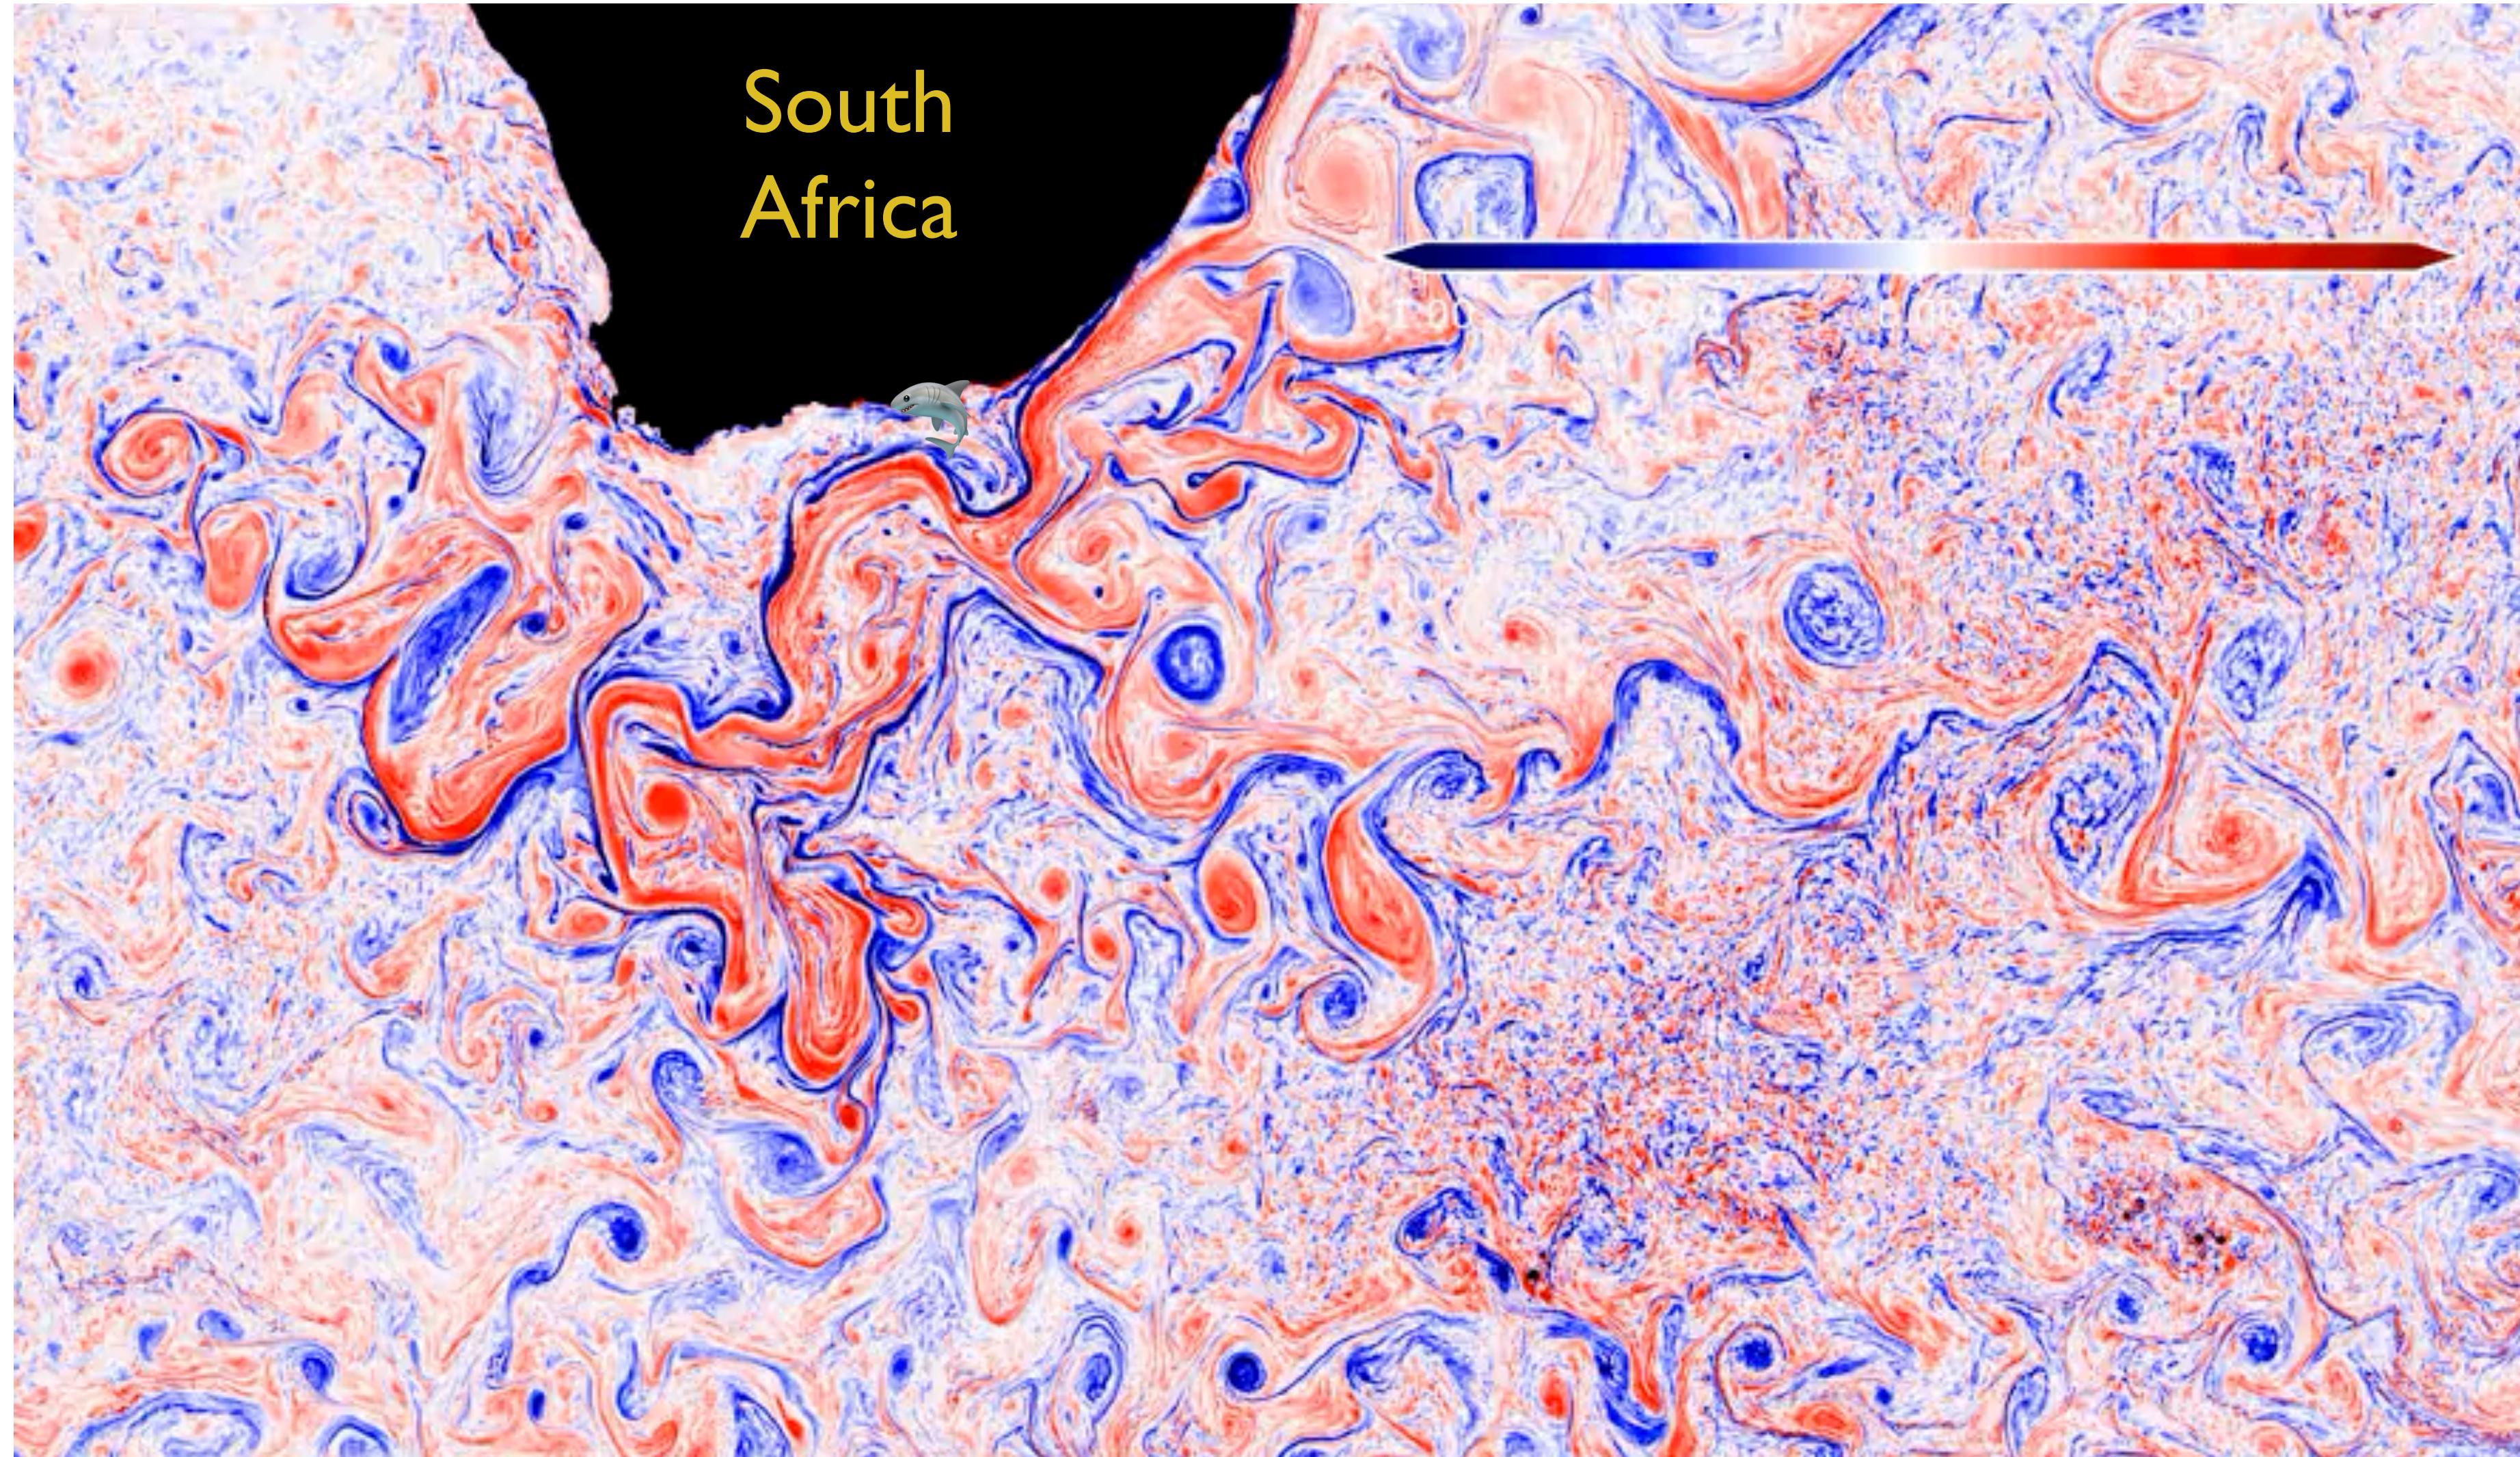
overturning circulation

~10.000 km  
~decades-centuries

LLC4320 sea surface speed animation  
by Henze and Menemenlis (NASA/JPL)  
1/48<sup>th</sup> degree or ~1-2 km & 90 vertical levels  
(one of the biggest ocean simulation ever run...)



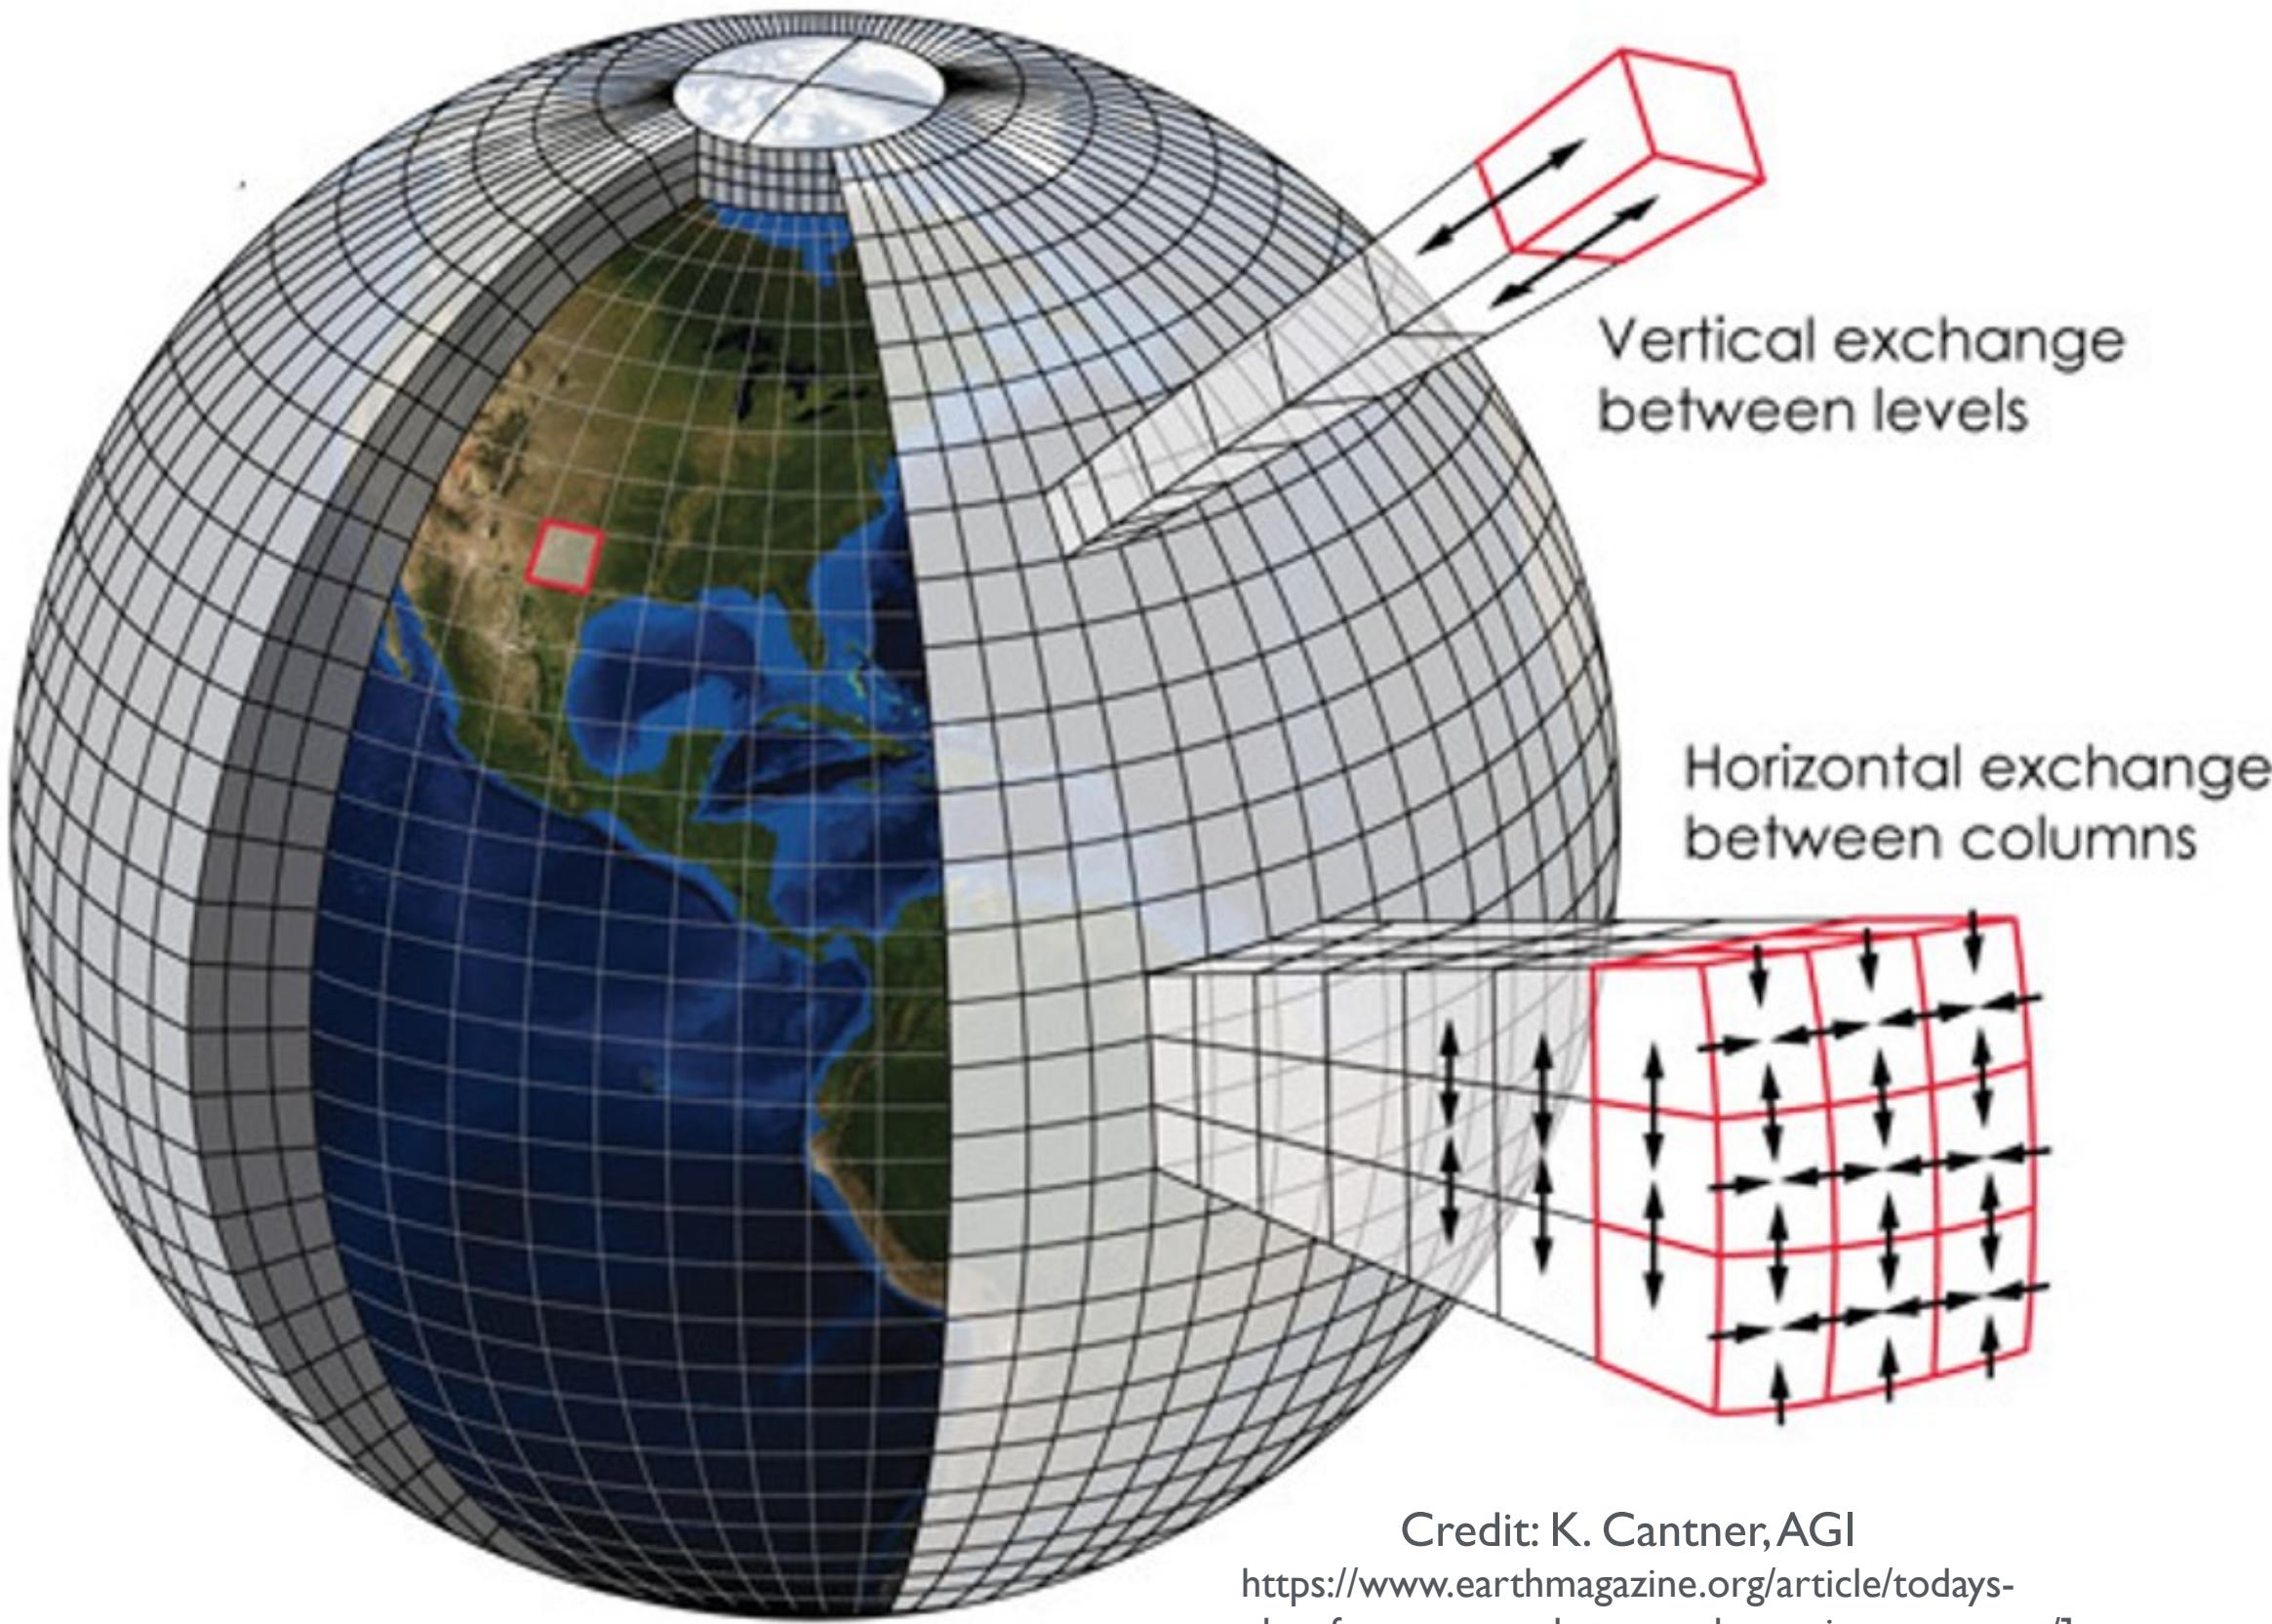
**turbulence comes at many scales!**  
the more we zoom in, the more see...



climate models



# how do we simulate climate?



Newton's 2nd law  
(mass  $\times$  acceleration = force)  
for every grid box!

(Navier–Stokes equations + thermodynamics)



IBM Blue Gene/P supercomputer  
with 164,000 processors!

(imagine 164,000 laptops side by side)

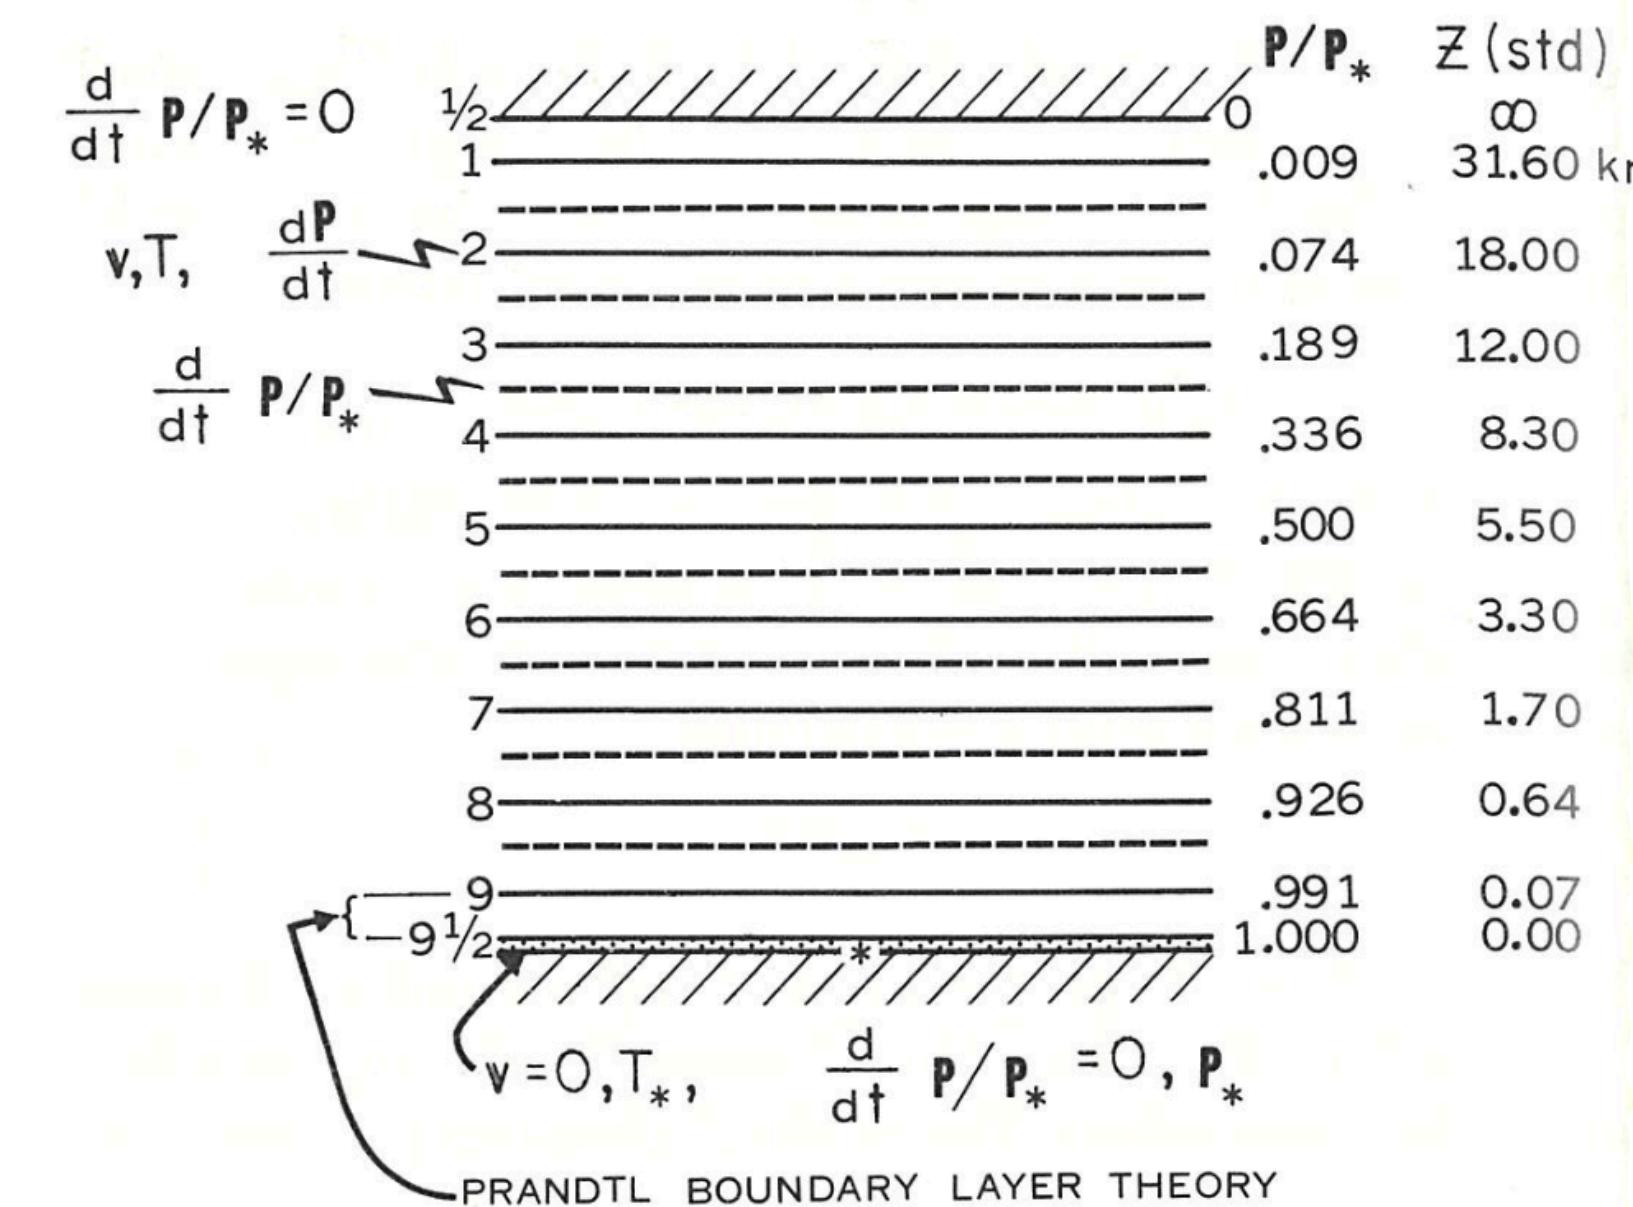
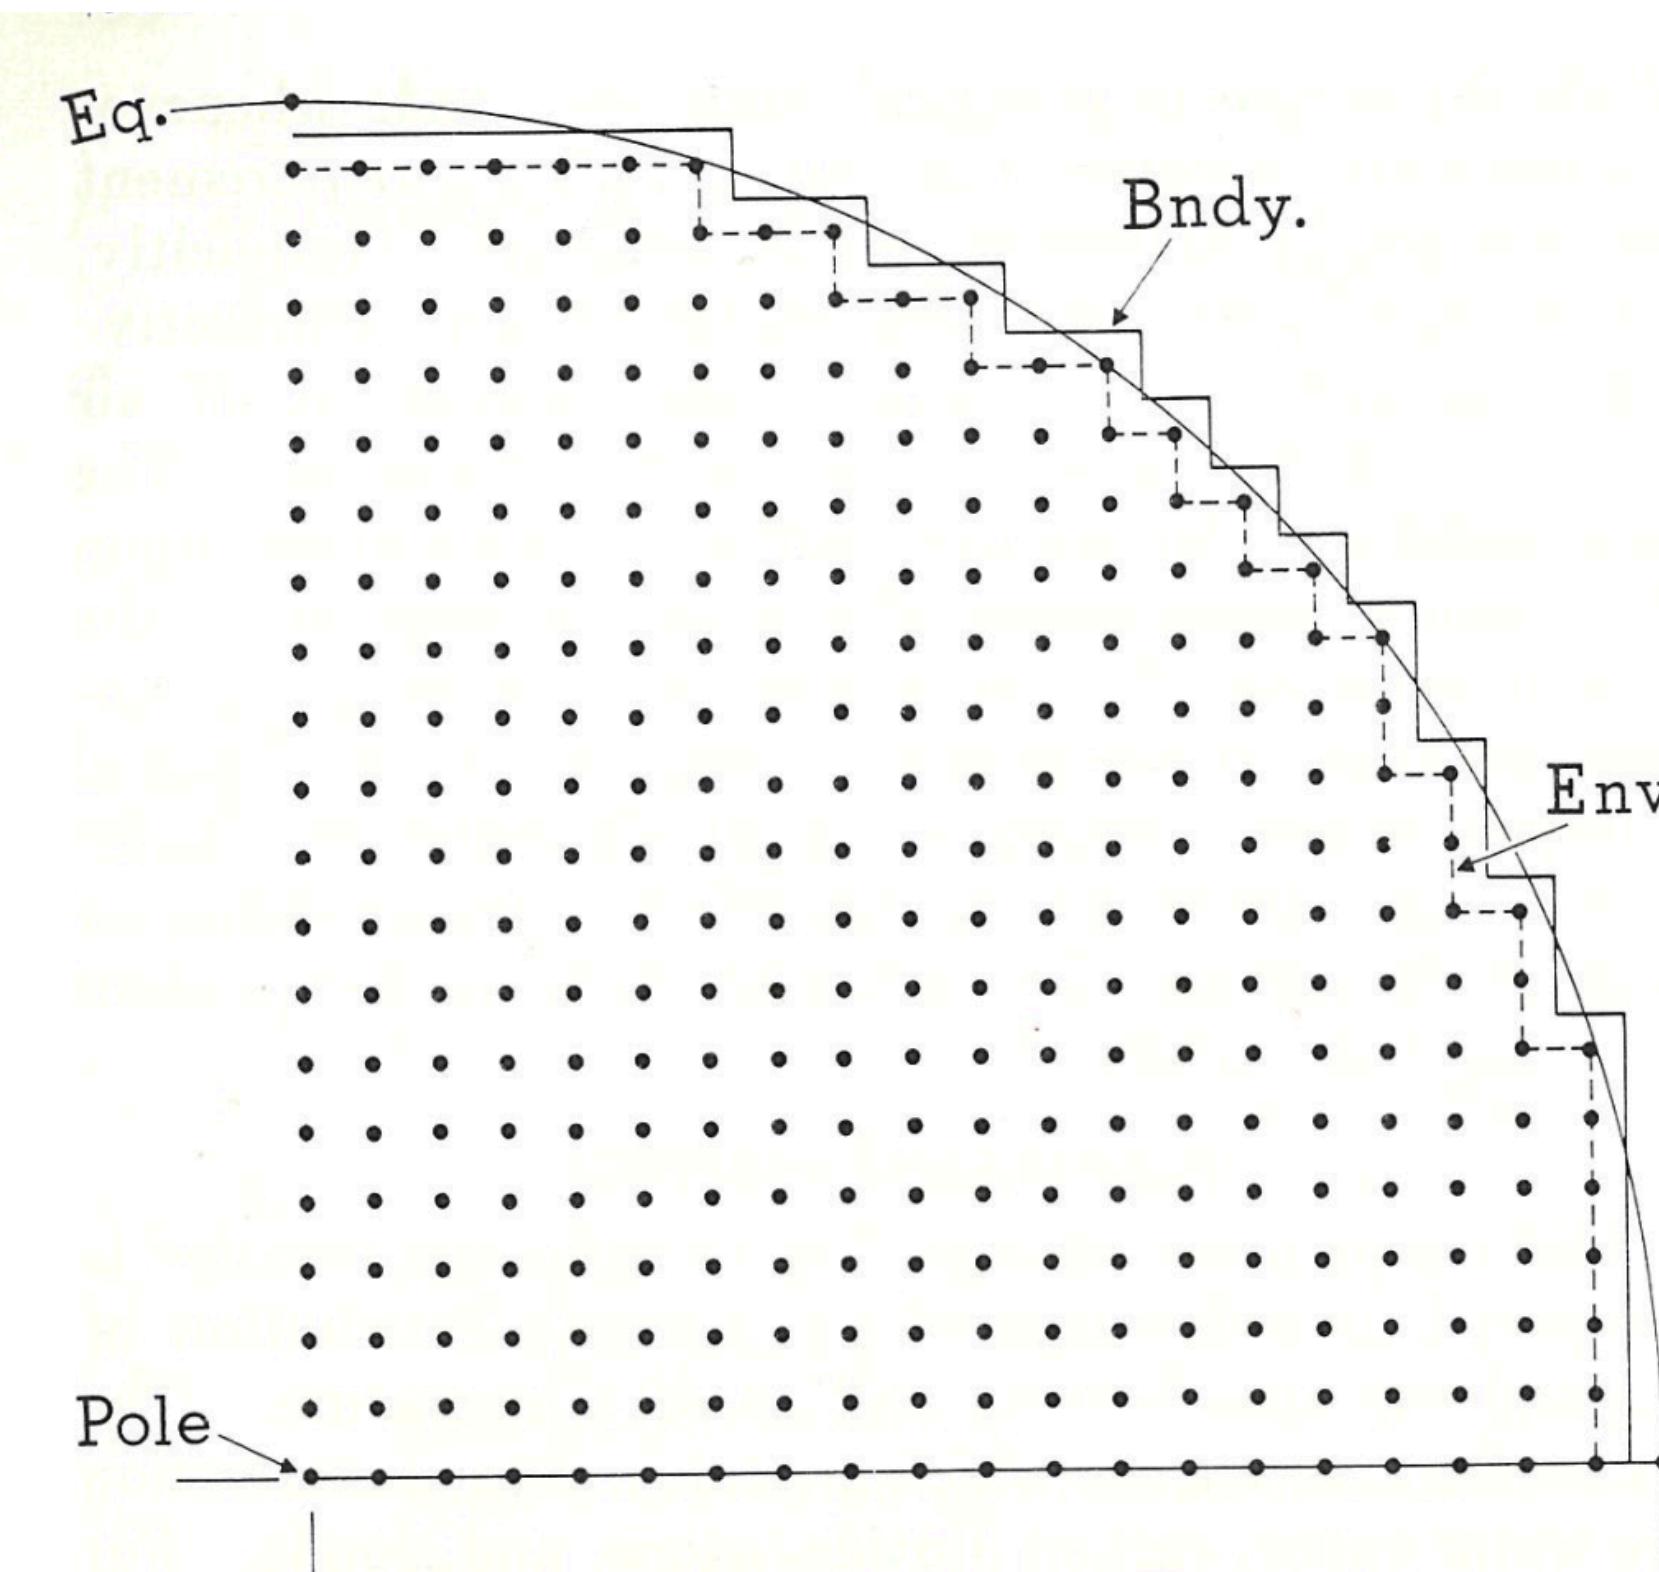
# a bit of history of climate modeling

first atmosphere  
model



Joseph Smagorinsky, Suki Manabe, & J. Leith Holloway  
Numerical results from a nine-level general circulation model of the atmosphere.

Mon. Wether Rev. 1965



# a bit of history of climate modeling

first atmosphere  
model



Joseph Smagorinsky, Suki Manabe, & J. Leith Holloway  
Numerical results from a nine-level general circulation model of the atmosphere.

Mon. Wether Rev. 1965

first ocean  
model



Kirk Bryan & Michael Cox  
A numerical investigation of the oceanic general circulation

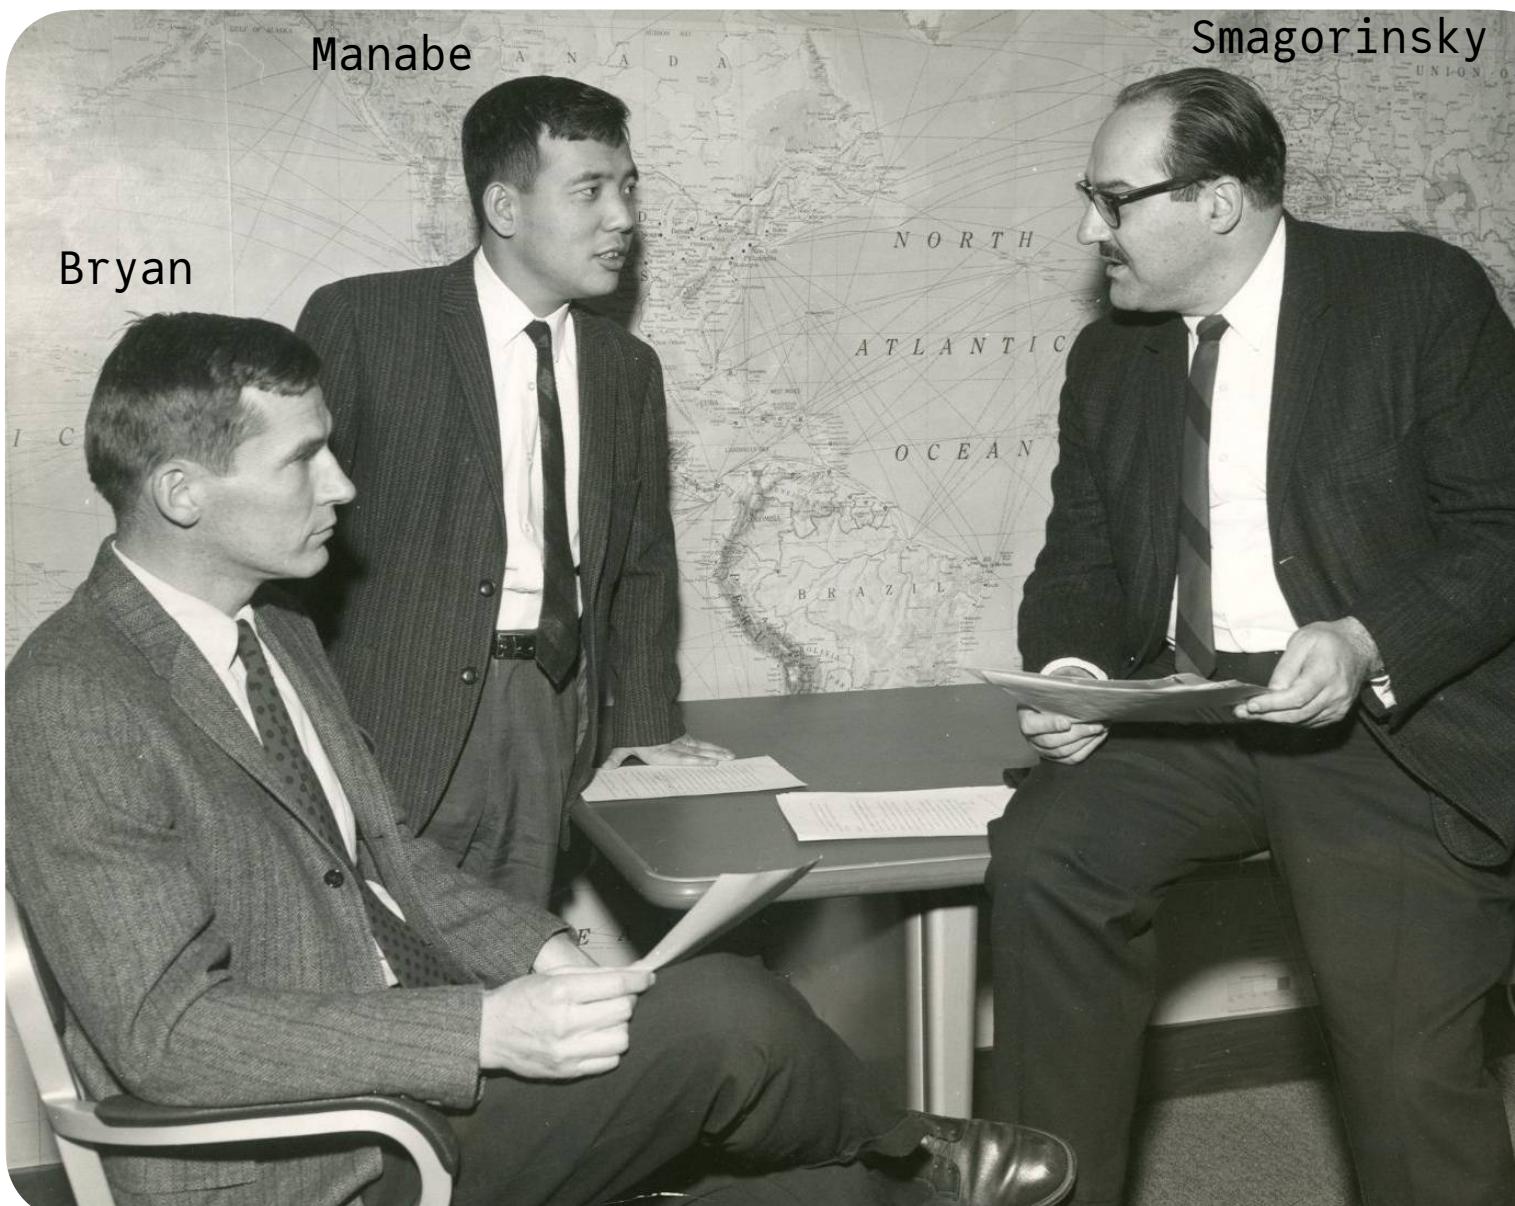
Tellus 1967  
(paper received 1965)

first coupled atmos-ocean  
model



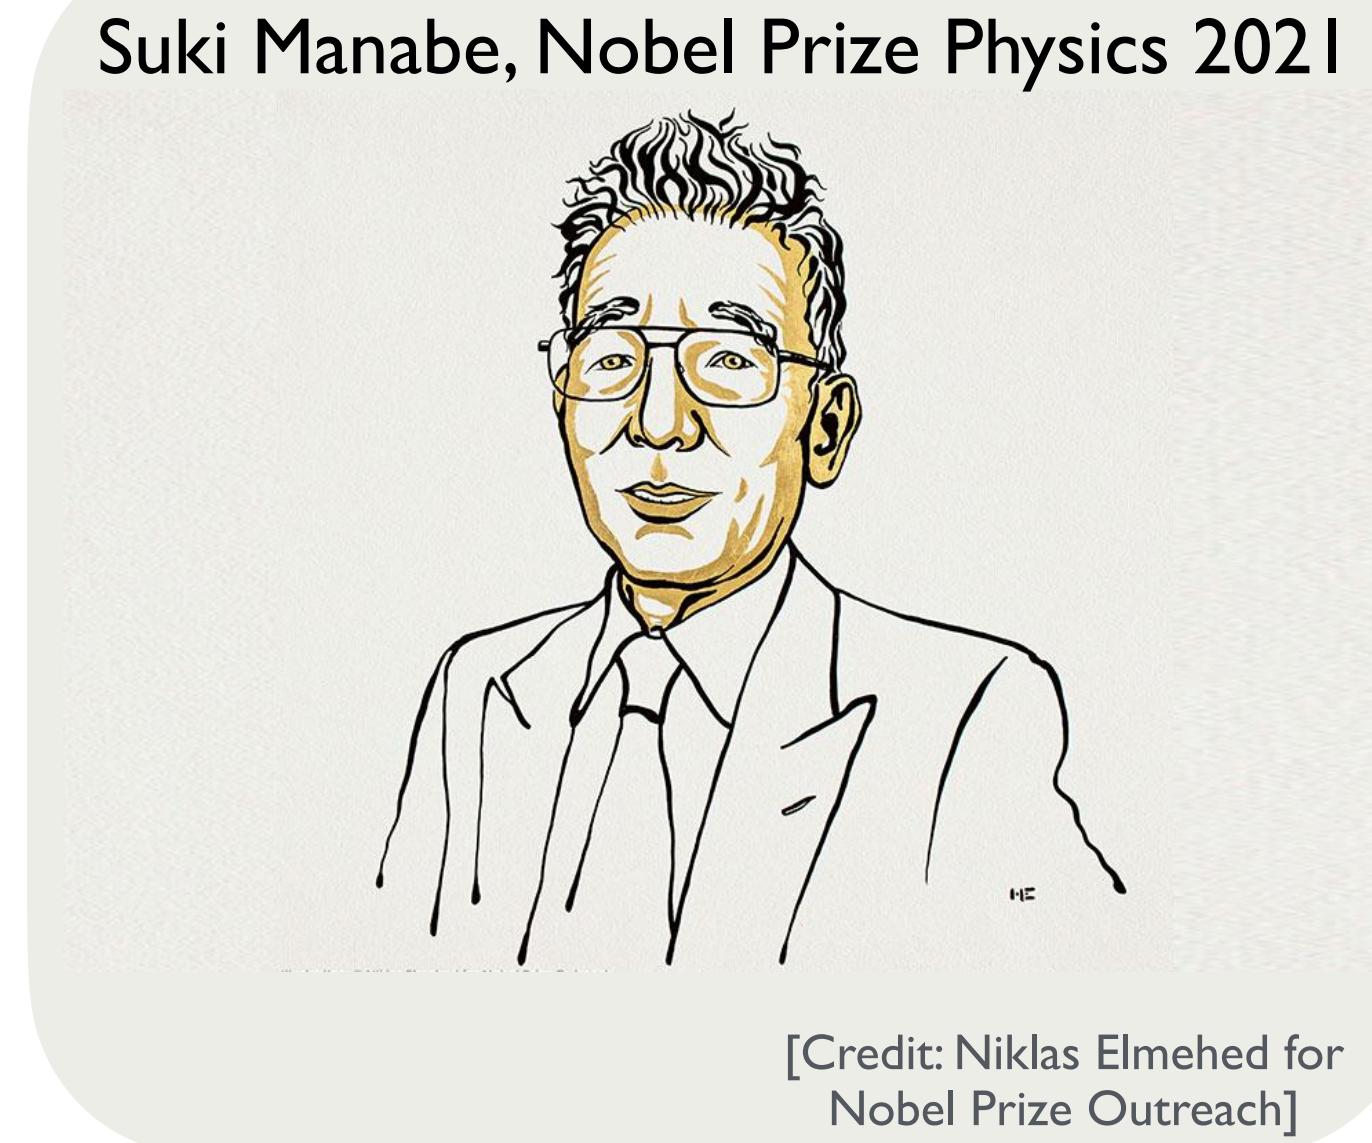
Suki Manabe & Kirk Bryan  
Climate calculations with a combined ocean-atmosphere model

J. Atmos. Sci. 1969



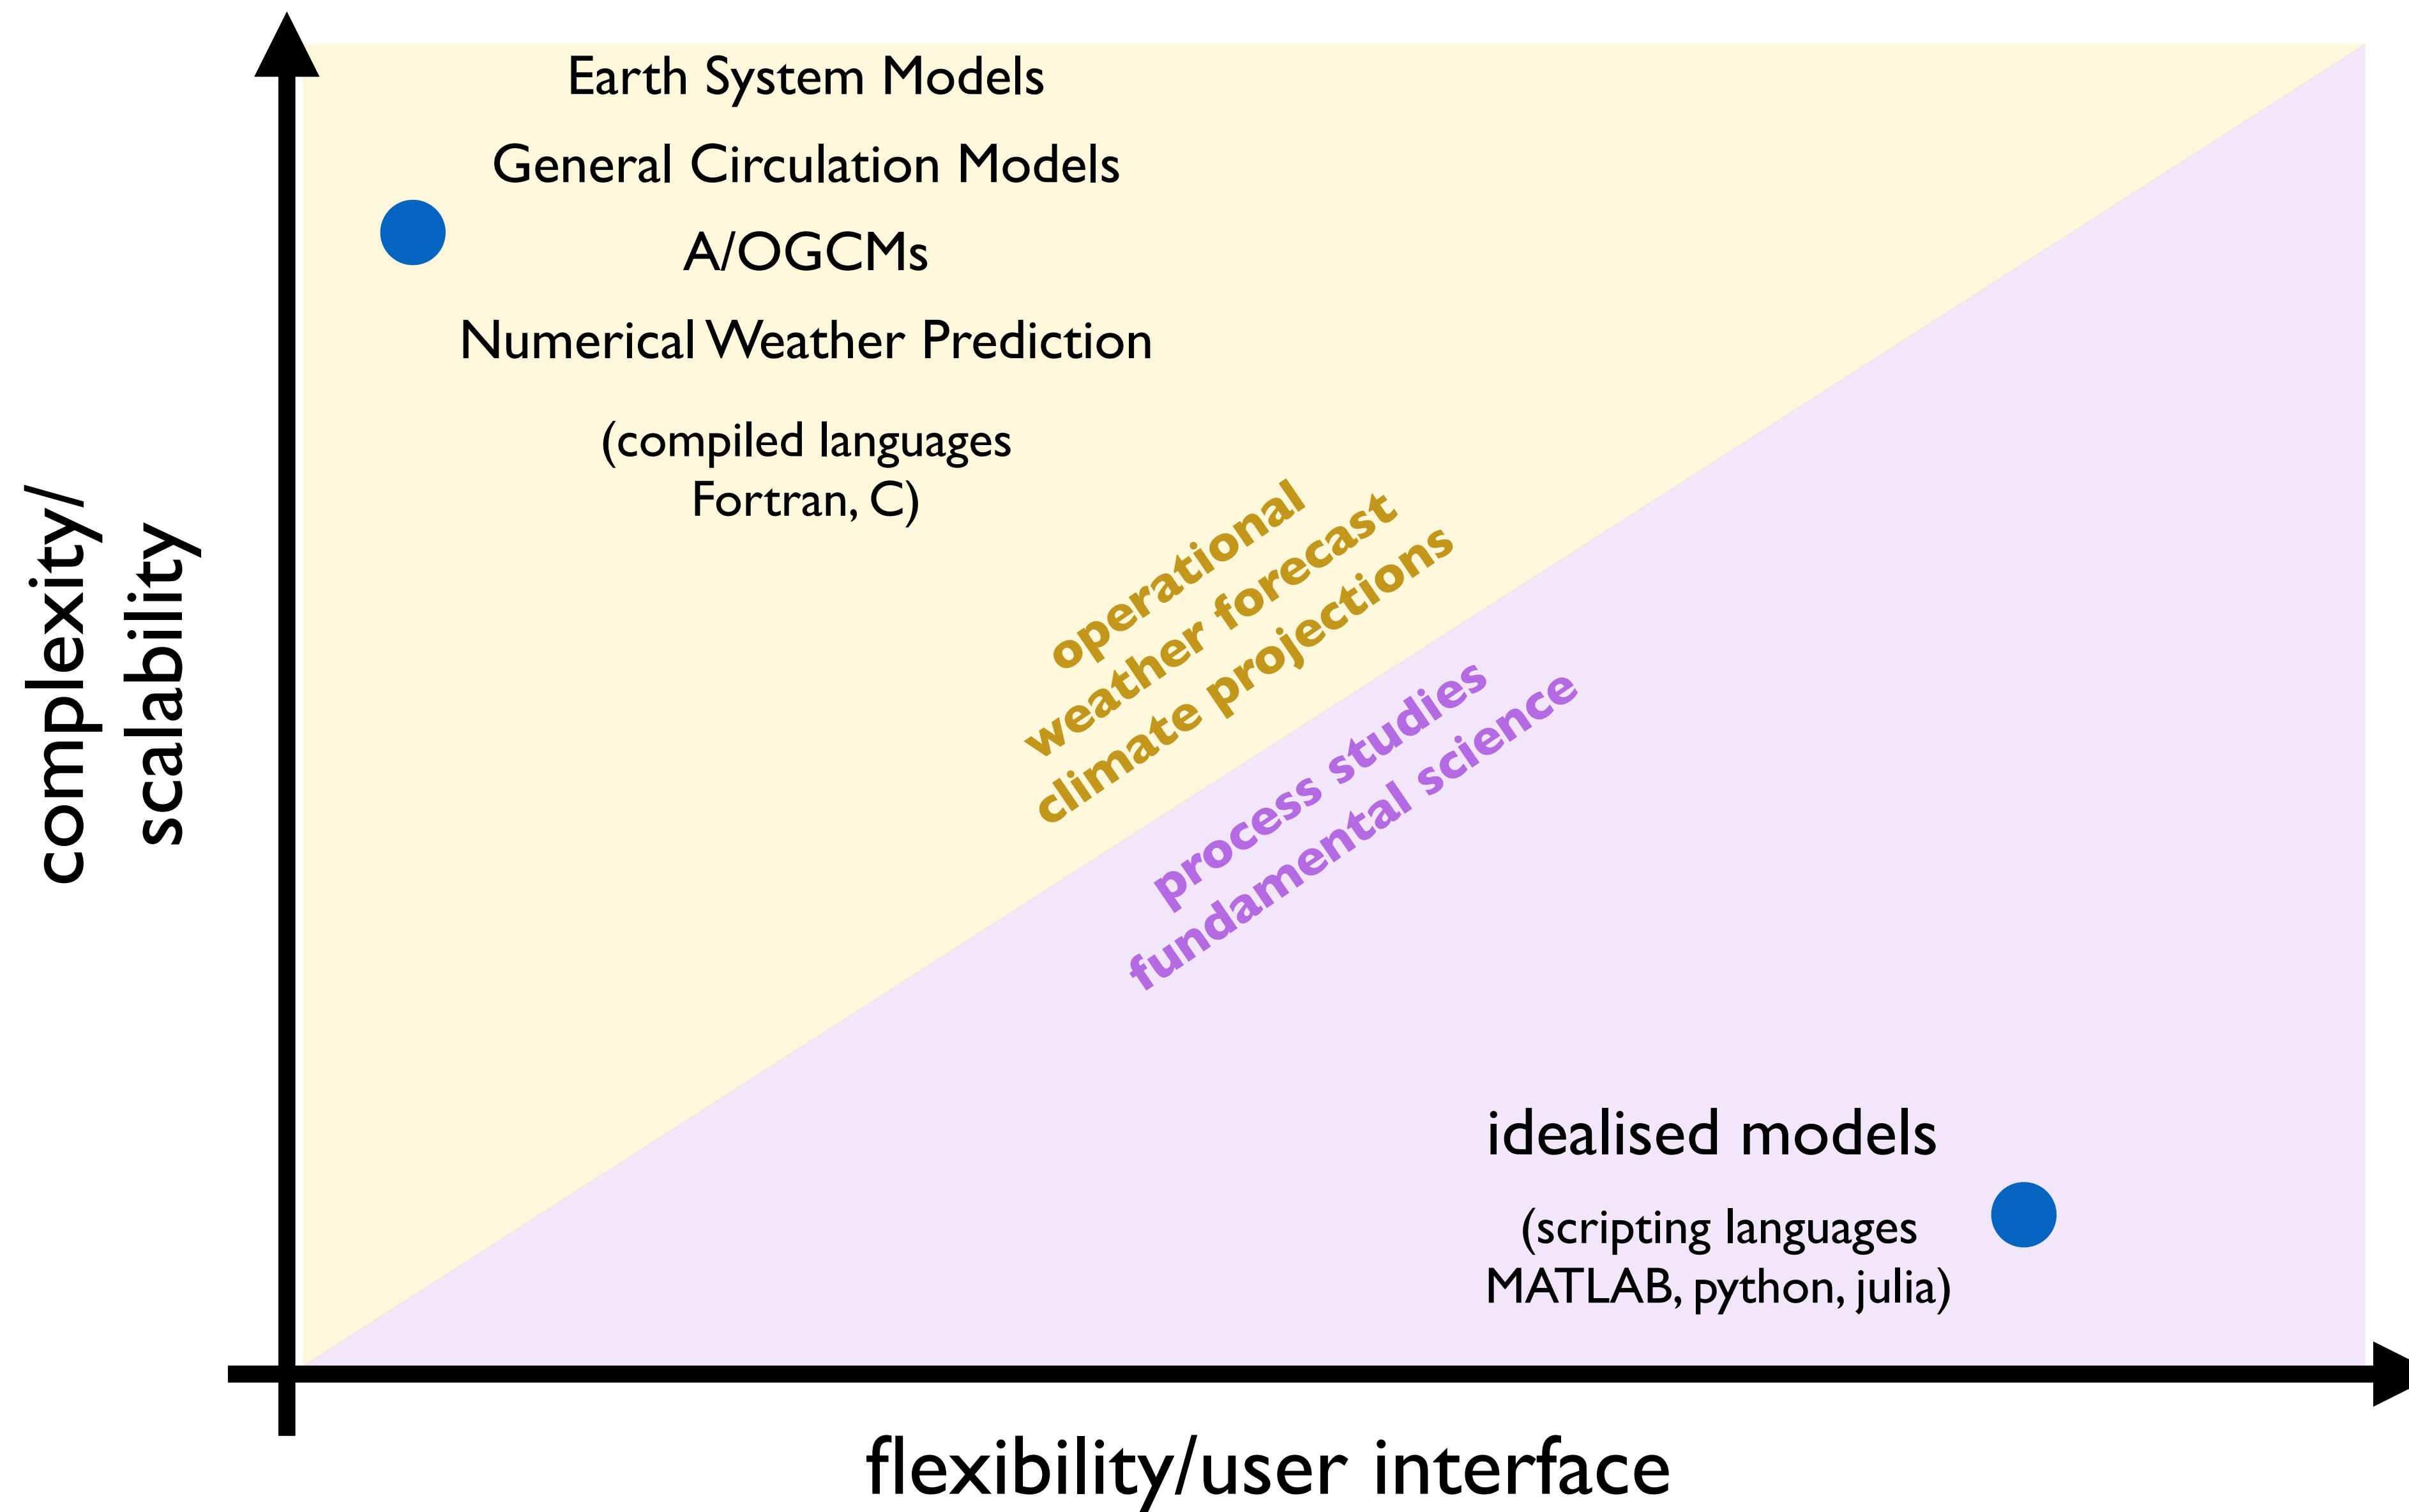
[photo courtesy of GFLD]

Suki Manabe, Nobel Prize Physics 2021

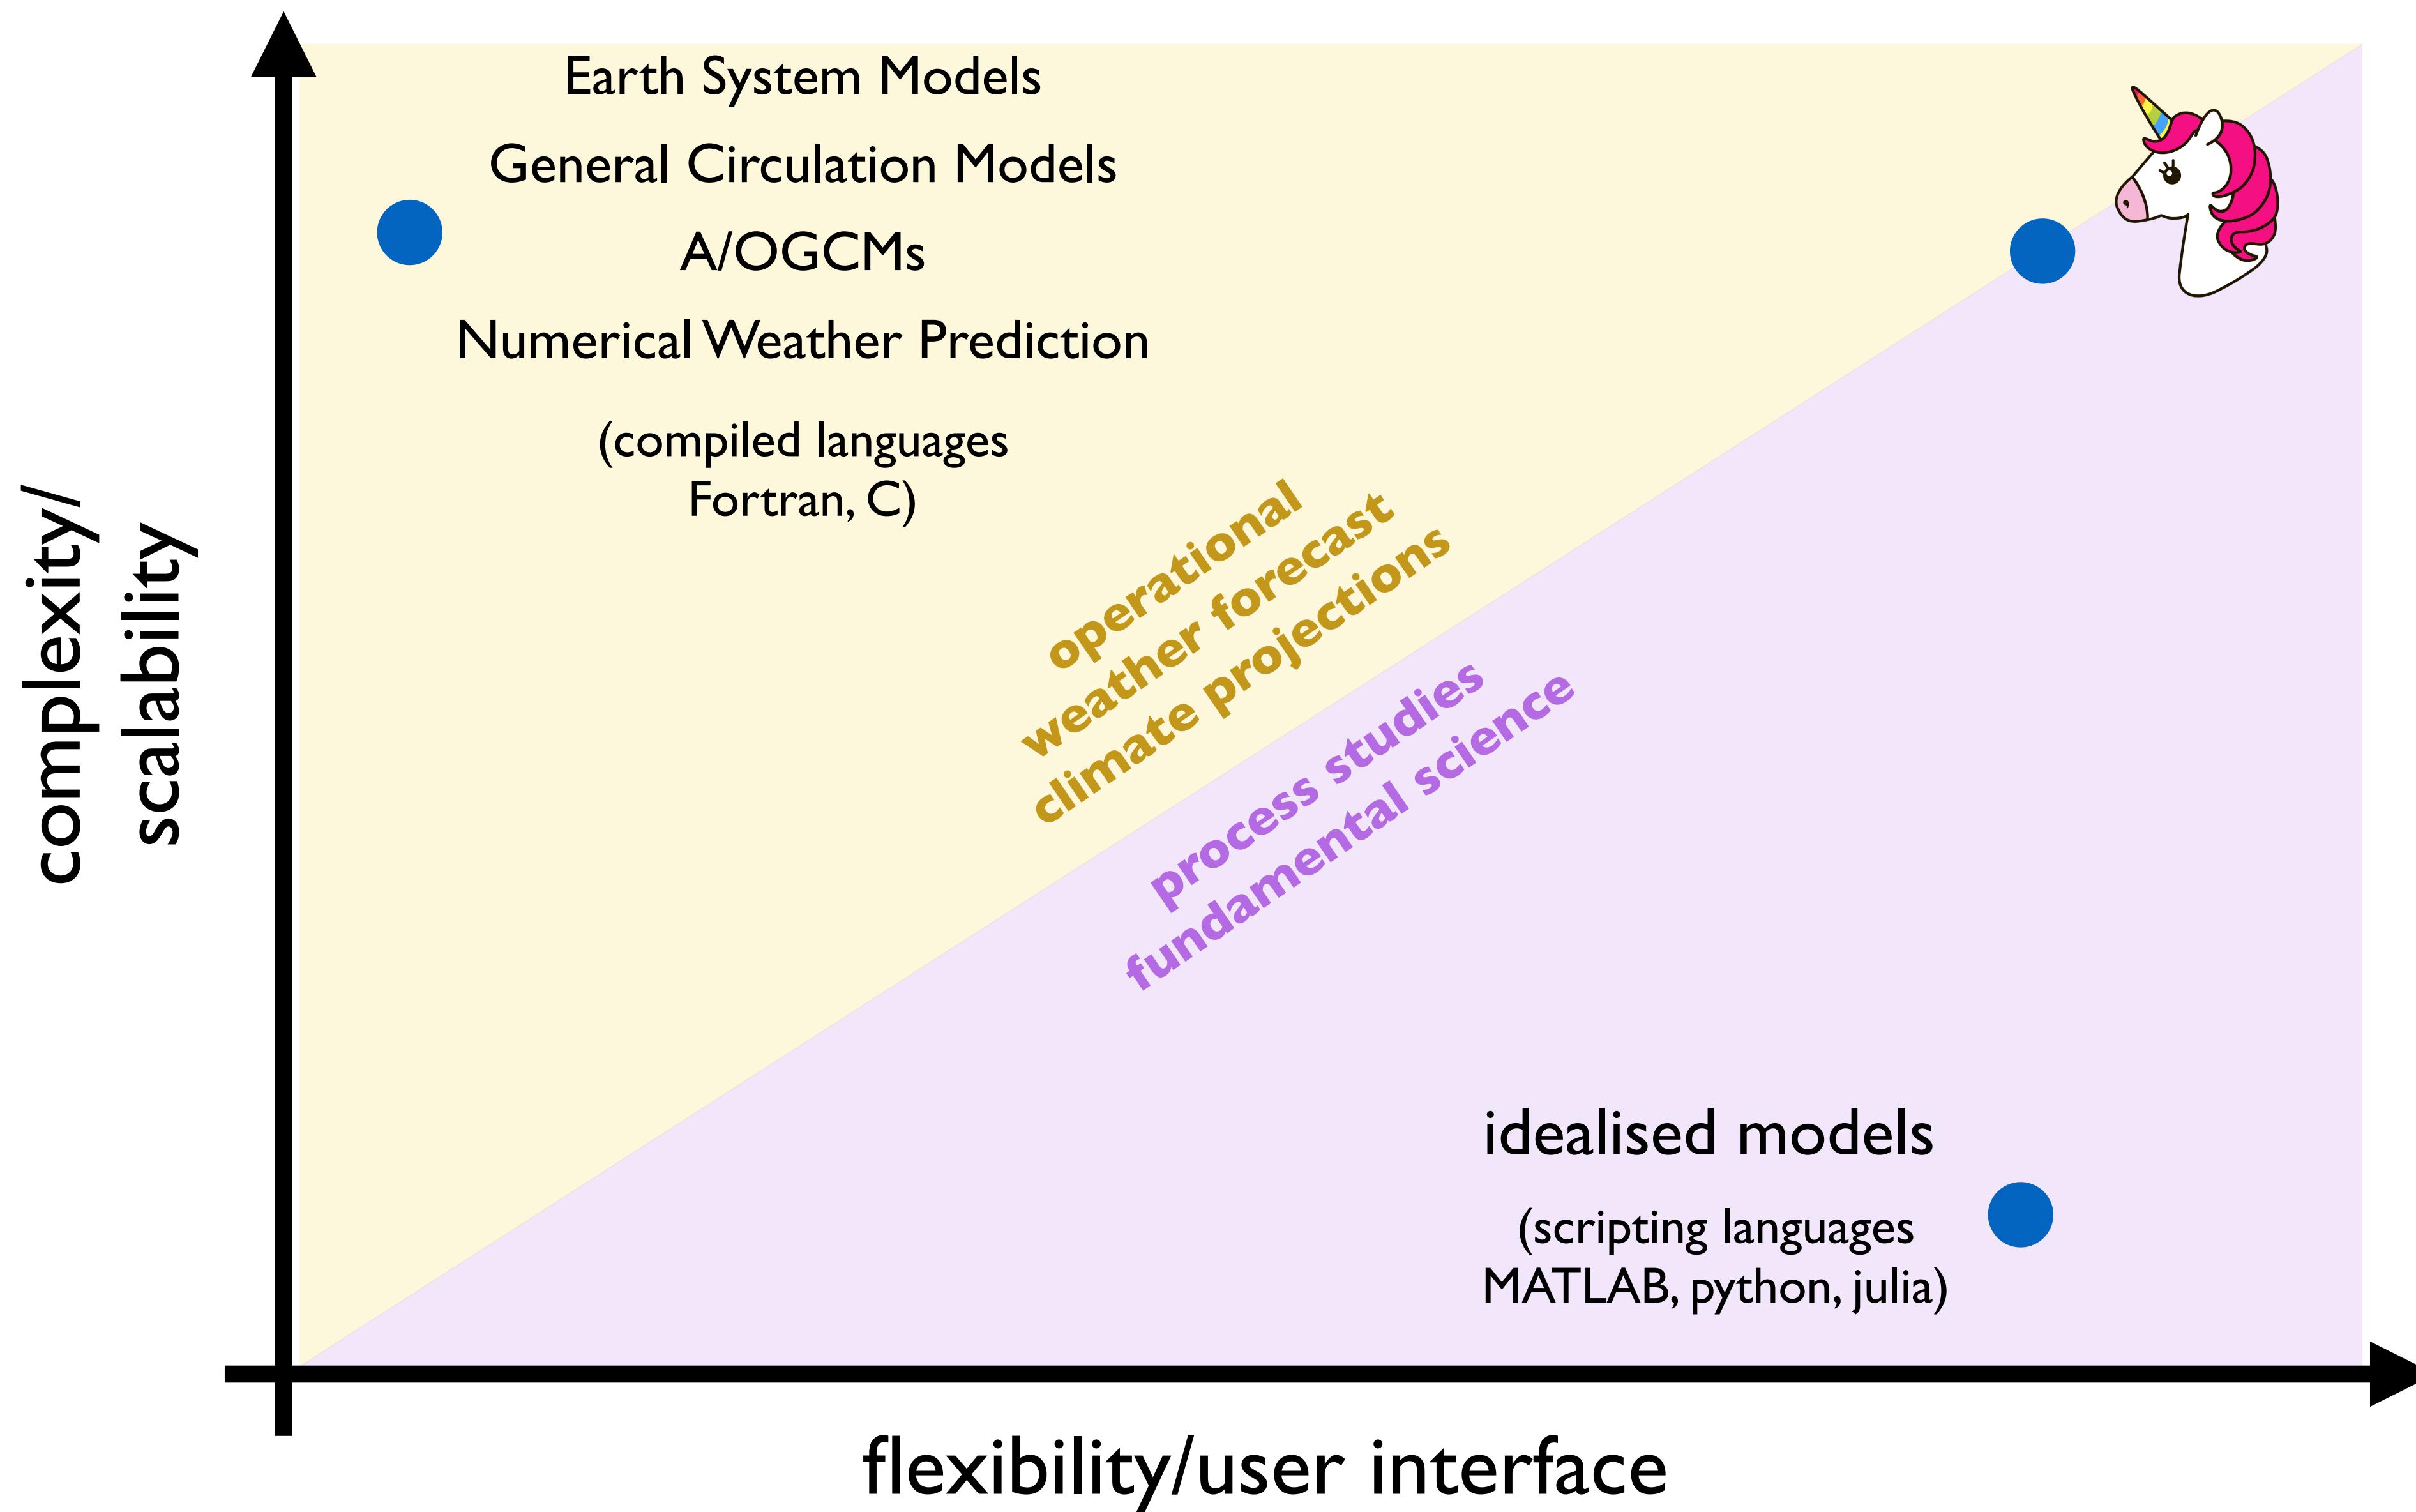


[Credit: Niklas Elmehed for  
Nobel Prize Outreach]

# different motive leads to different model



# different motive leads to different model



# is user interface important?

## internet on mobile phones



all devices have internet capabilities

but, when did people *actually* start using internet on their phones?

# namelist = list of parameter values

allow users to change parameters  
without modifying the source code

The NAMELIST input data file specifies the following:

```
NFIRST=1      start from scratch
NLAST=42      stop after timestep 42
NENERGY=20    print fields every 20 timesteps
NTSI=1        print single line of information every timestep
NMIX=10       do a mixing timestep every 10 timesteps
AMF=1.E9      horizontal mixing of momentum
AHF=2.E7      horizontal mixing of heat, salinity, tracers
FKPMF=1.0     vertical mixing of momentum
FKPHF=1.0     vertical mixing of heat, salinity, tracers
DTTSF=172800. length of timestep on temp., sal., tracers
DTUVF=7200.   length of timestep on internal mode velocities
DTSFF=7200.   length of timestep on stream function
MXSCAN=25     cut off relaxation at 25 scans if not converged
SORF=1.50    coefficient of over-relaxation
CRITF=4.E9    criterion for convergence of relaxation
ACORF=0.5    weight forward and backward timestep equally in
              the semi-implicit Coriolis term
TINITF=...    initial values of temperature, salinity, tracers
              (note.. for purposes of precision, salinity is
              carried in the model in units of parts per part,
              with .035 subtracted off, however TINITF is ex-
              pected in parts per part; .035 is subtracted
              after read-in)
ISIS=13,...   the island and its perimeter points are included
              in a box between I=13-16 and J=9-12
```



a namelist from the  
GFDL ocean model  
back in 1984

Model 1 is a duplicate of the sample model Semtner supplies in his report. It consists of a closed, rectangular basin with a shelf along the western wall and a small island in the center. The update file, NAMELIST input data file, and output file for this model are given on the following pages. The option record specifies that an island is included, that the model is to be run core contained, and that comments are desired in the code listing. The updates which follow simply alter the base code to yield a model like that of Semtner's sample. Note that many of the updates arise because an alternate equation of state is used. There are differences in the manner in which the computation is handled between the two models, so that the answers will not be exactly the same. In particular, Semtner's version uses a relaxation which does not converge as rapidly as that used here, and, in fact, is cut off on most timesteps at 26 scans, before convergence occurs. In the present model, the convergence is completed for all but a few timesteps. The result is that the answers are somewhat different but presumably more accurate in the present model.

# namelist = list of parameter values

allow users to change parameters  
without modifying the source code

```
NIGLOBAL = 1440
!
! The total number of thickness grid points in the x-direction in the physical
! domain. With STATIC_MEMORY_ this is set in MOM_memory.h at compile time.

NJGLOBAL = 450
!
! The total number of thickness grid points in the y-direction in the physical
! domain. With STATIC_MEMORY_ this is set in MOM_memory.h at compile time.

! LAYOUT = 24, 12
!
! The processor layout that was actually used.

! === module MOM ===

THICKNESSDIFFUSE_FIRST = True      ! [Boolean] default = False
                                    ! If true, do thickness diffusion before dynamics. This is only used if
                                    ! THICKNESSDIFFUSE is true.

DT = 600.0                          ! [s]
                                    ! The (baroclinic) dynamics time step. The time-step that is actually used will
                                    ! be an integer fraction of the forcing time-step (DT_FORCING in ocean-only mode
                                    ! or the coupling timestep in coupled mode.)

DT_THERM = 3600.0                   ! [s] default = 600.0
                                    ! The thermodynamic and tracer advection time step. Ideally DT_THERM should be
                                    ! an integer multiple of DT and less than the forcing or coupling time-step,
                                    ! unless THERMO_SPANS_COUPLING is true, in which case DT_THERM can be an integer
                                    ! multiple of the coupling timestep. By default DT_THERM is set to DT.
```

a namelist from an  
IPCC-class GCM  
ocean model



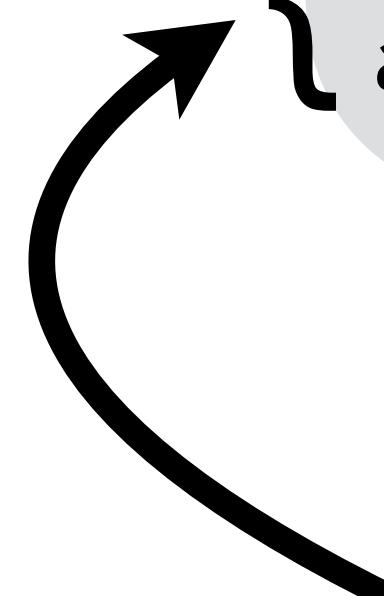
## parameters.namelist

```
Nx = 360          ! Number of points in longitude  
Ny = 180          ! Number of points in latitude  
  
force_with_dataset = True      ! JRA55 atmospheric reanalysis dataset
```

what if we want to add surface heating fluxes  $SHF_{\text{surf}} = Q_0 \cos(\text{lat})$  ?

## parameters.namelist

```
Nx = 360                      ! Number of points in longitude  
Ny = 180                      ! Number of points in latitude  
  
force_with_dataset = False      ! JRA55 atmospheric reanalysis dataset  
  
{ force_with_coslat = True       ! Surface heat flux ~ cos(lat)  
  amplitude_heating_forcing = 340 ! W/m^2, only if force_with_coslat = True
```



what if we want to add surface heating fluxes  $SHF_{surf} = Q_0 \cos(\text{lat})$  ?

given that such thing is  
implemented in main codebase

# setting up and running a simulation

## `ocean_simulation.script`

```
Nx = 360                      # Number of points in longitude
Ny = 180                      # Number of points in latitude

grid = GlobalGrid(Nx, Ny)

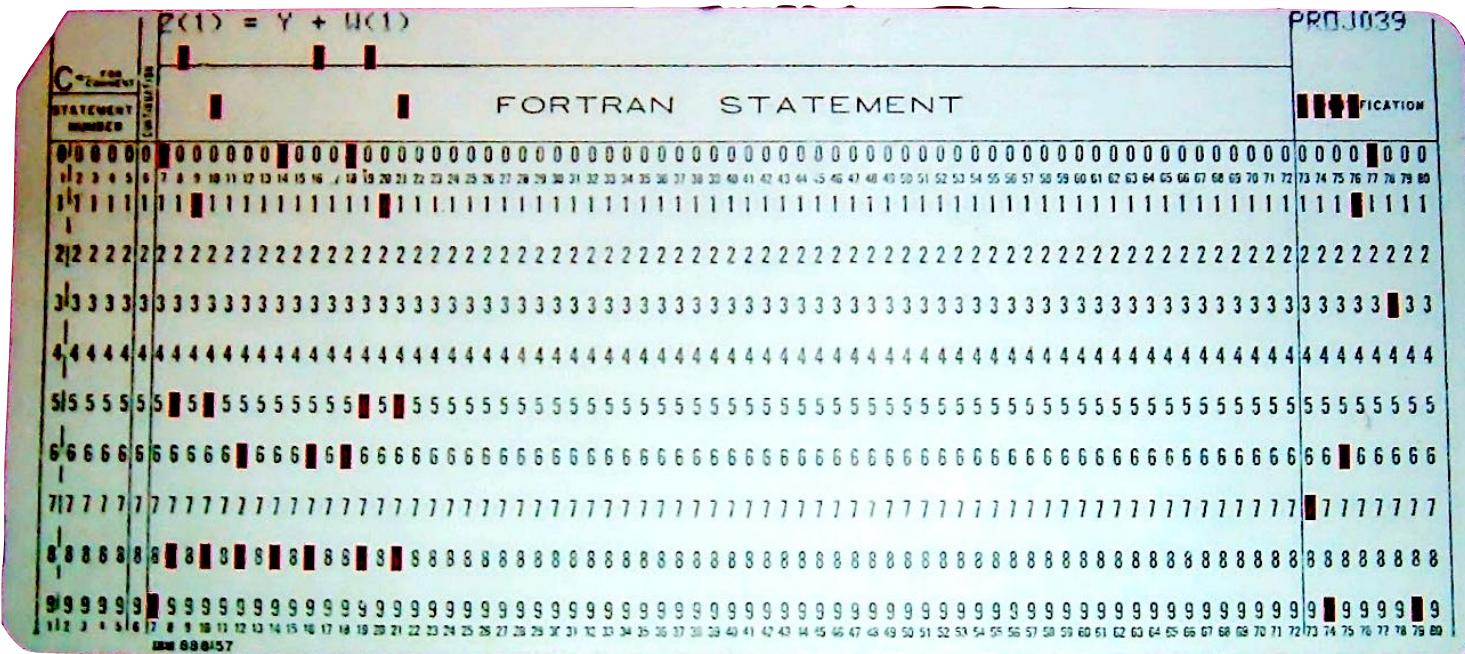
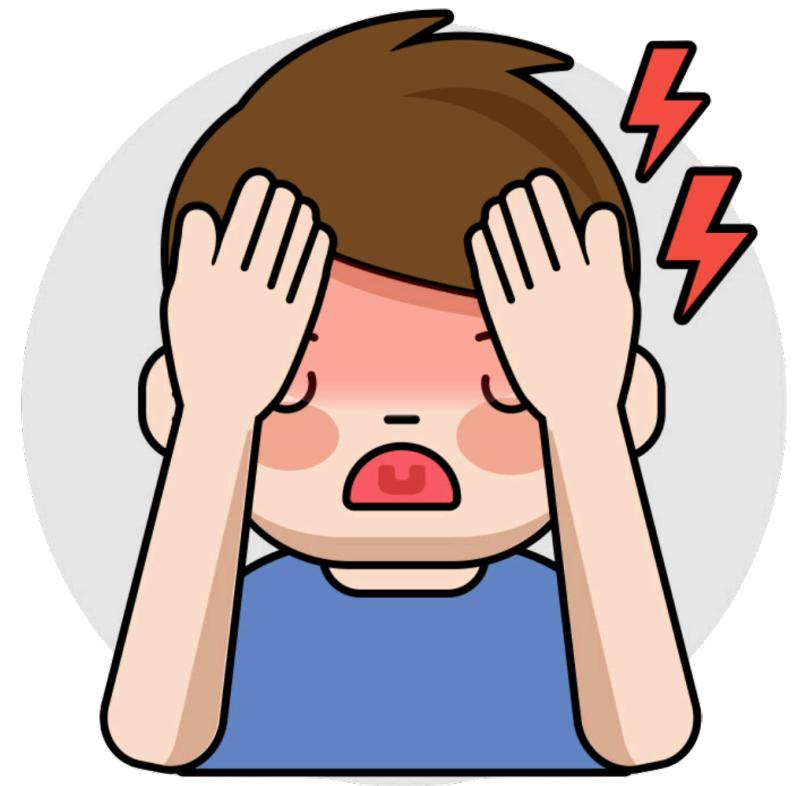
realistic_atmosphere = DataSet(path = "/data/atmospheric_reanalysis/")

idealised_atmosphere = IdealizedAtmosphere(grid)
set!(idealised_atmosphere, heat_flux = 340 * cos(deg2lat(grid.latitude)))

model = OceanModel(atmosphere = idealised_atmosphere, ...)

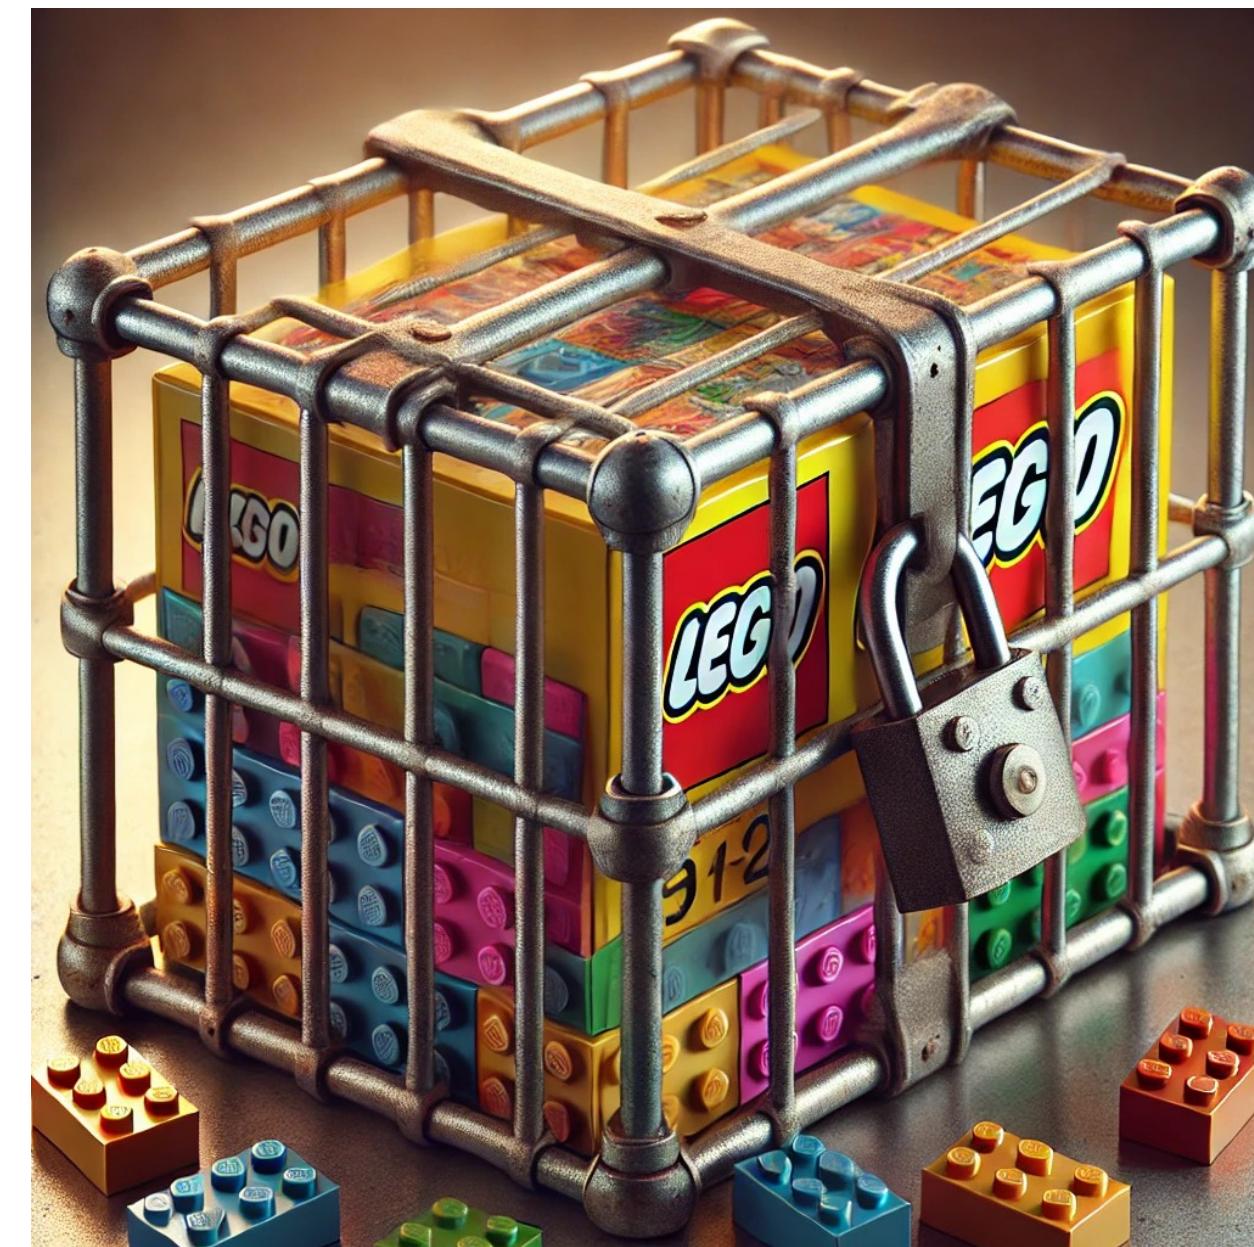
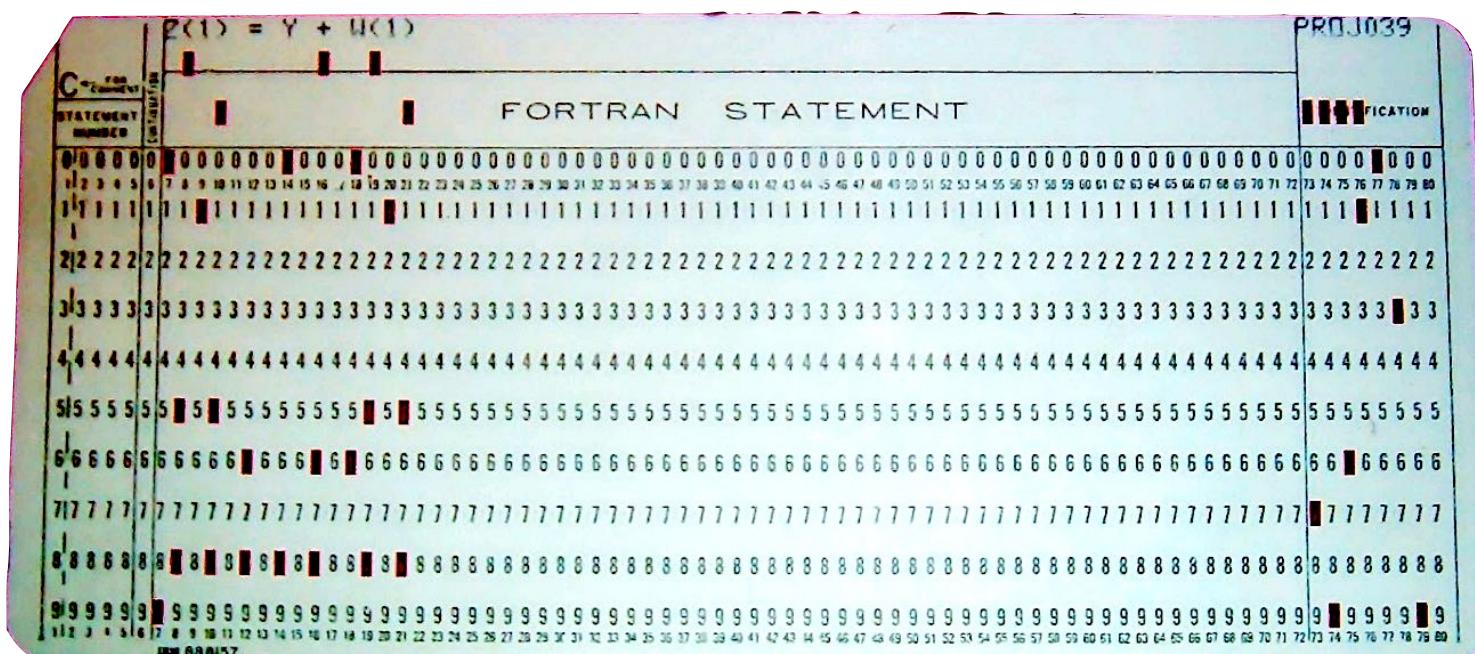
...
```

But we'd like more than namelists...  
models should be easy and fun!



# But we'd like more than namelists...

## models should be easy and fun!



[by Dall-E]



But we'd like more than namelists...  
models should be easy and fun!

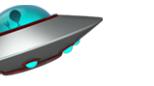
what if

want to switch model components in and out?

modify forcing, boundary conditions, bathymetry ,...?

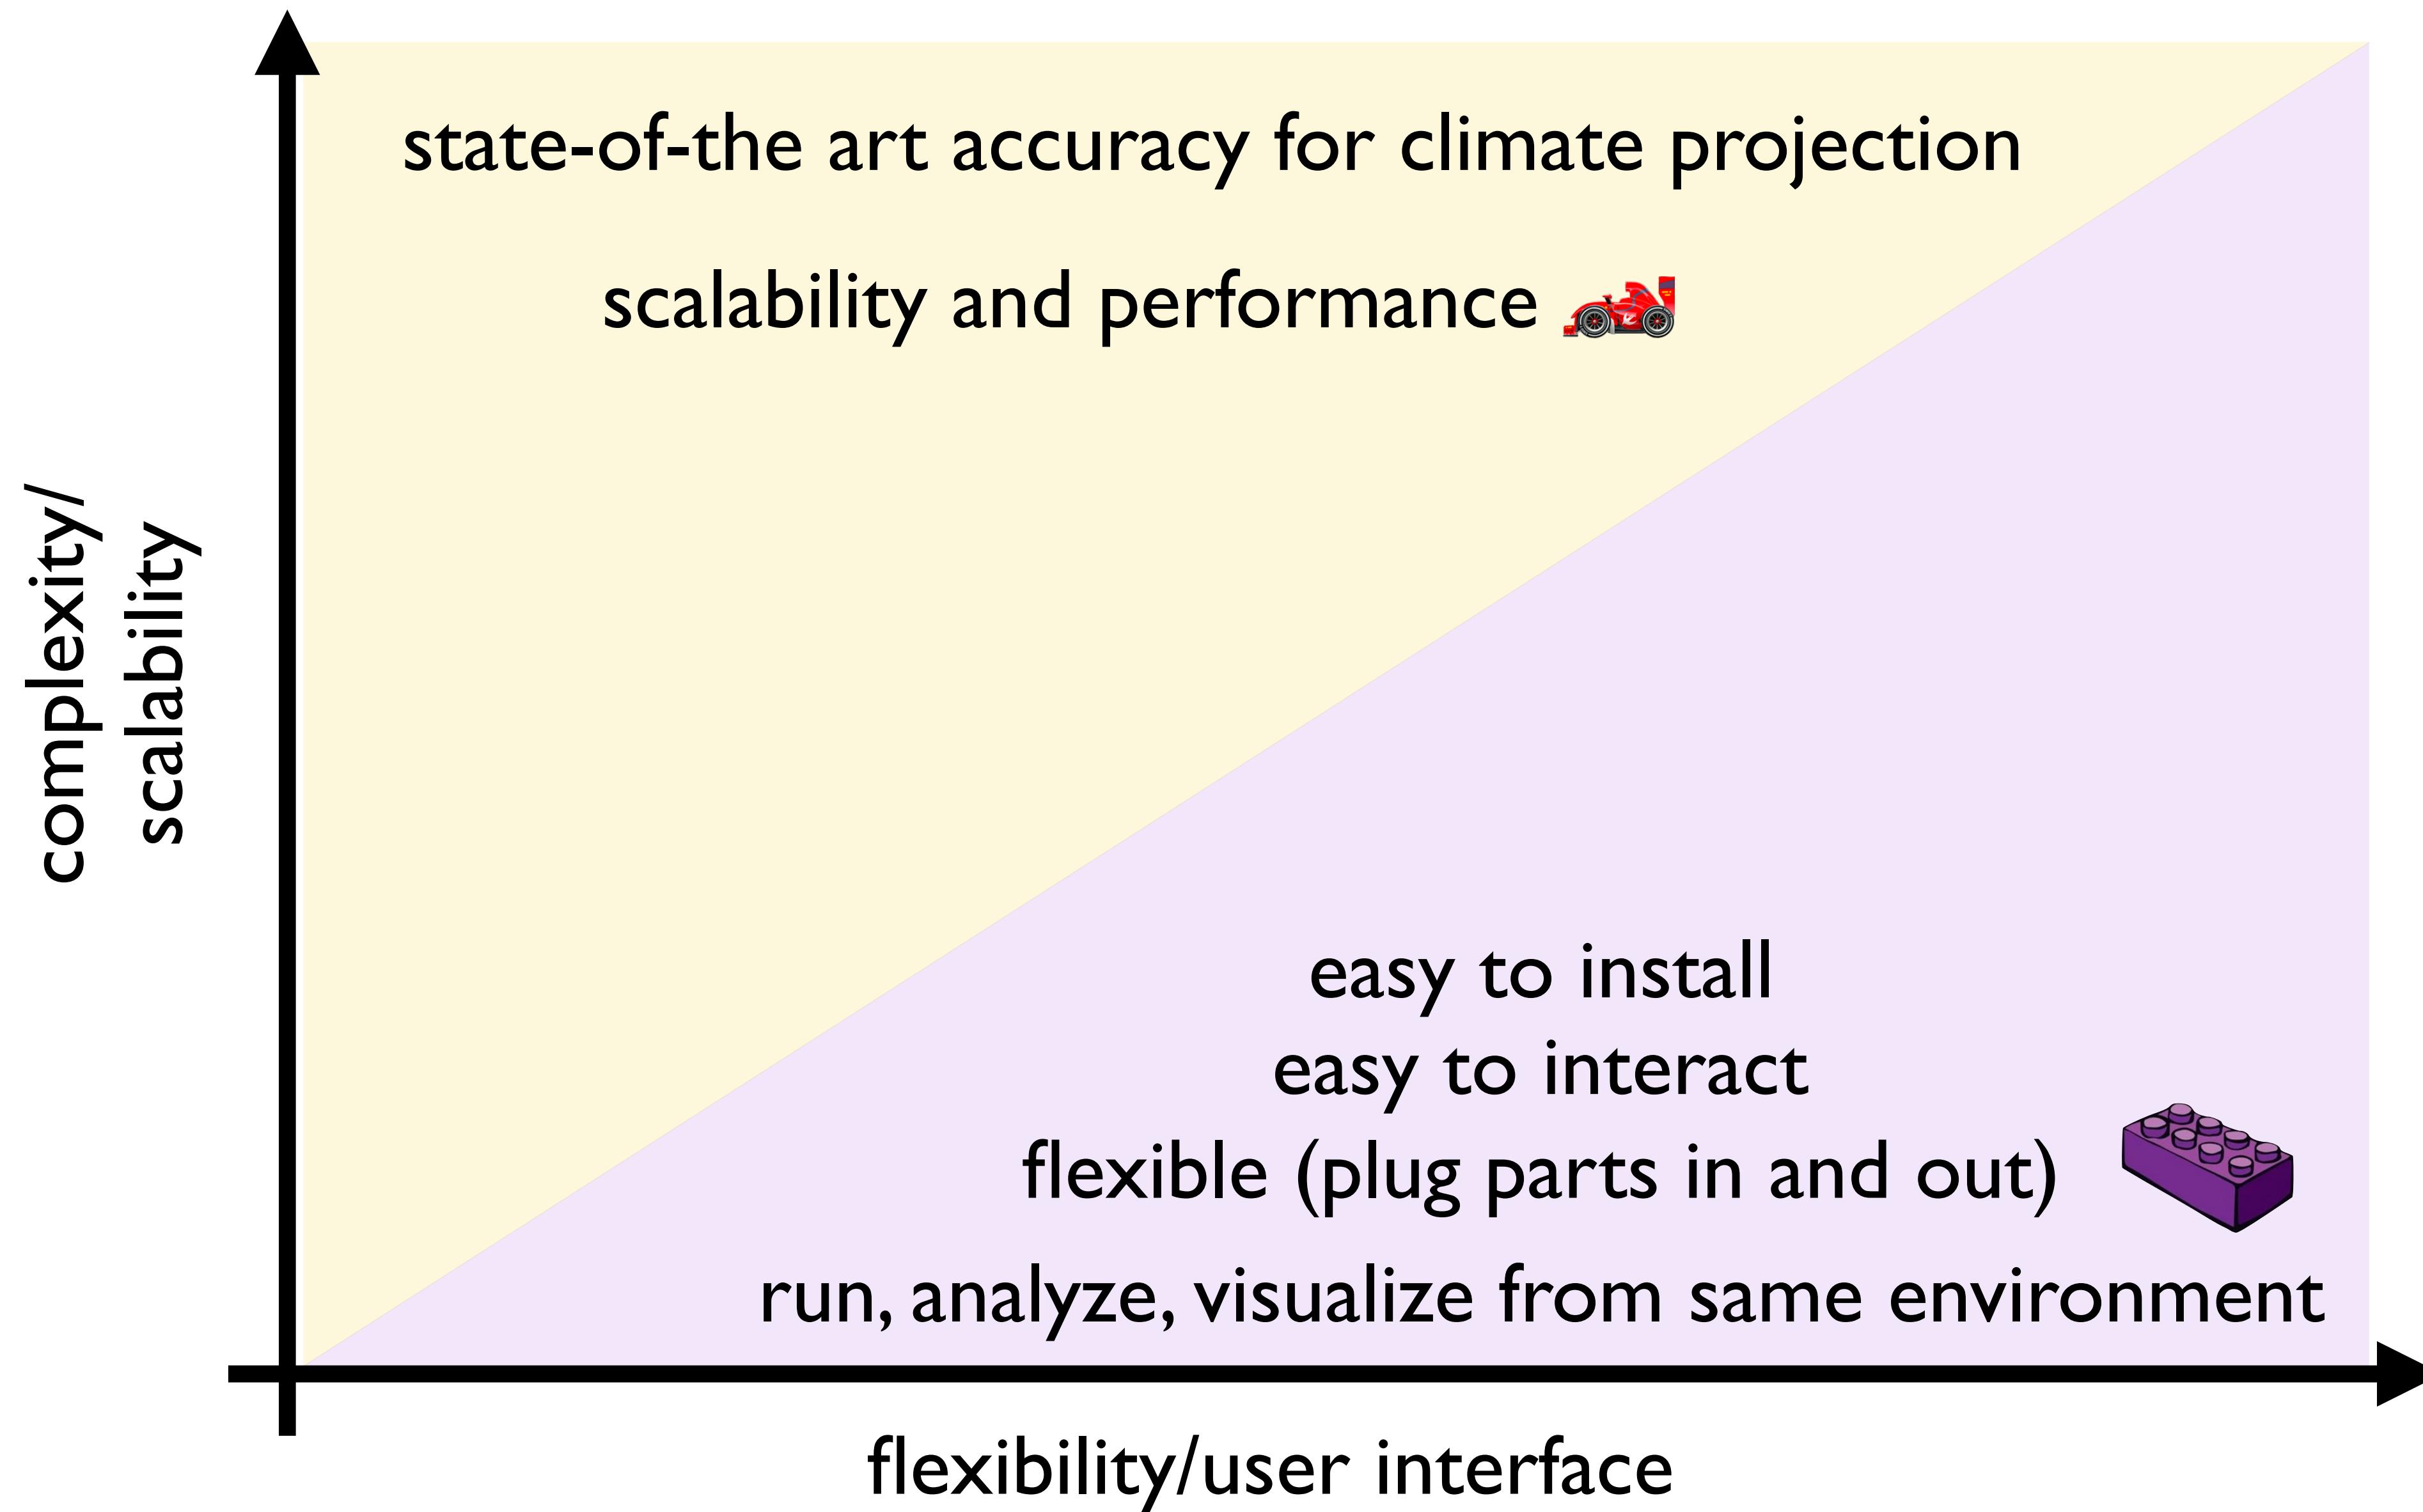
wanna share with other to reproduce your results?



modern software   
& its impact on our work flow in  
research  and teaching 

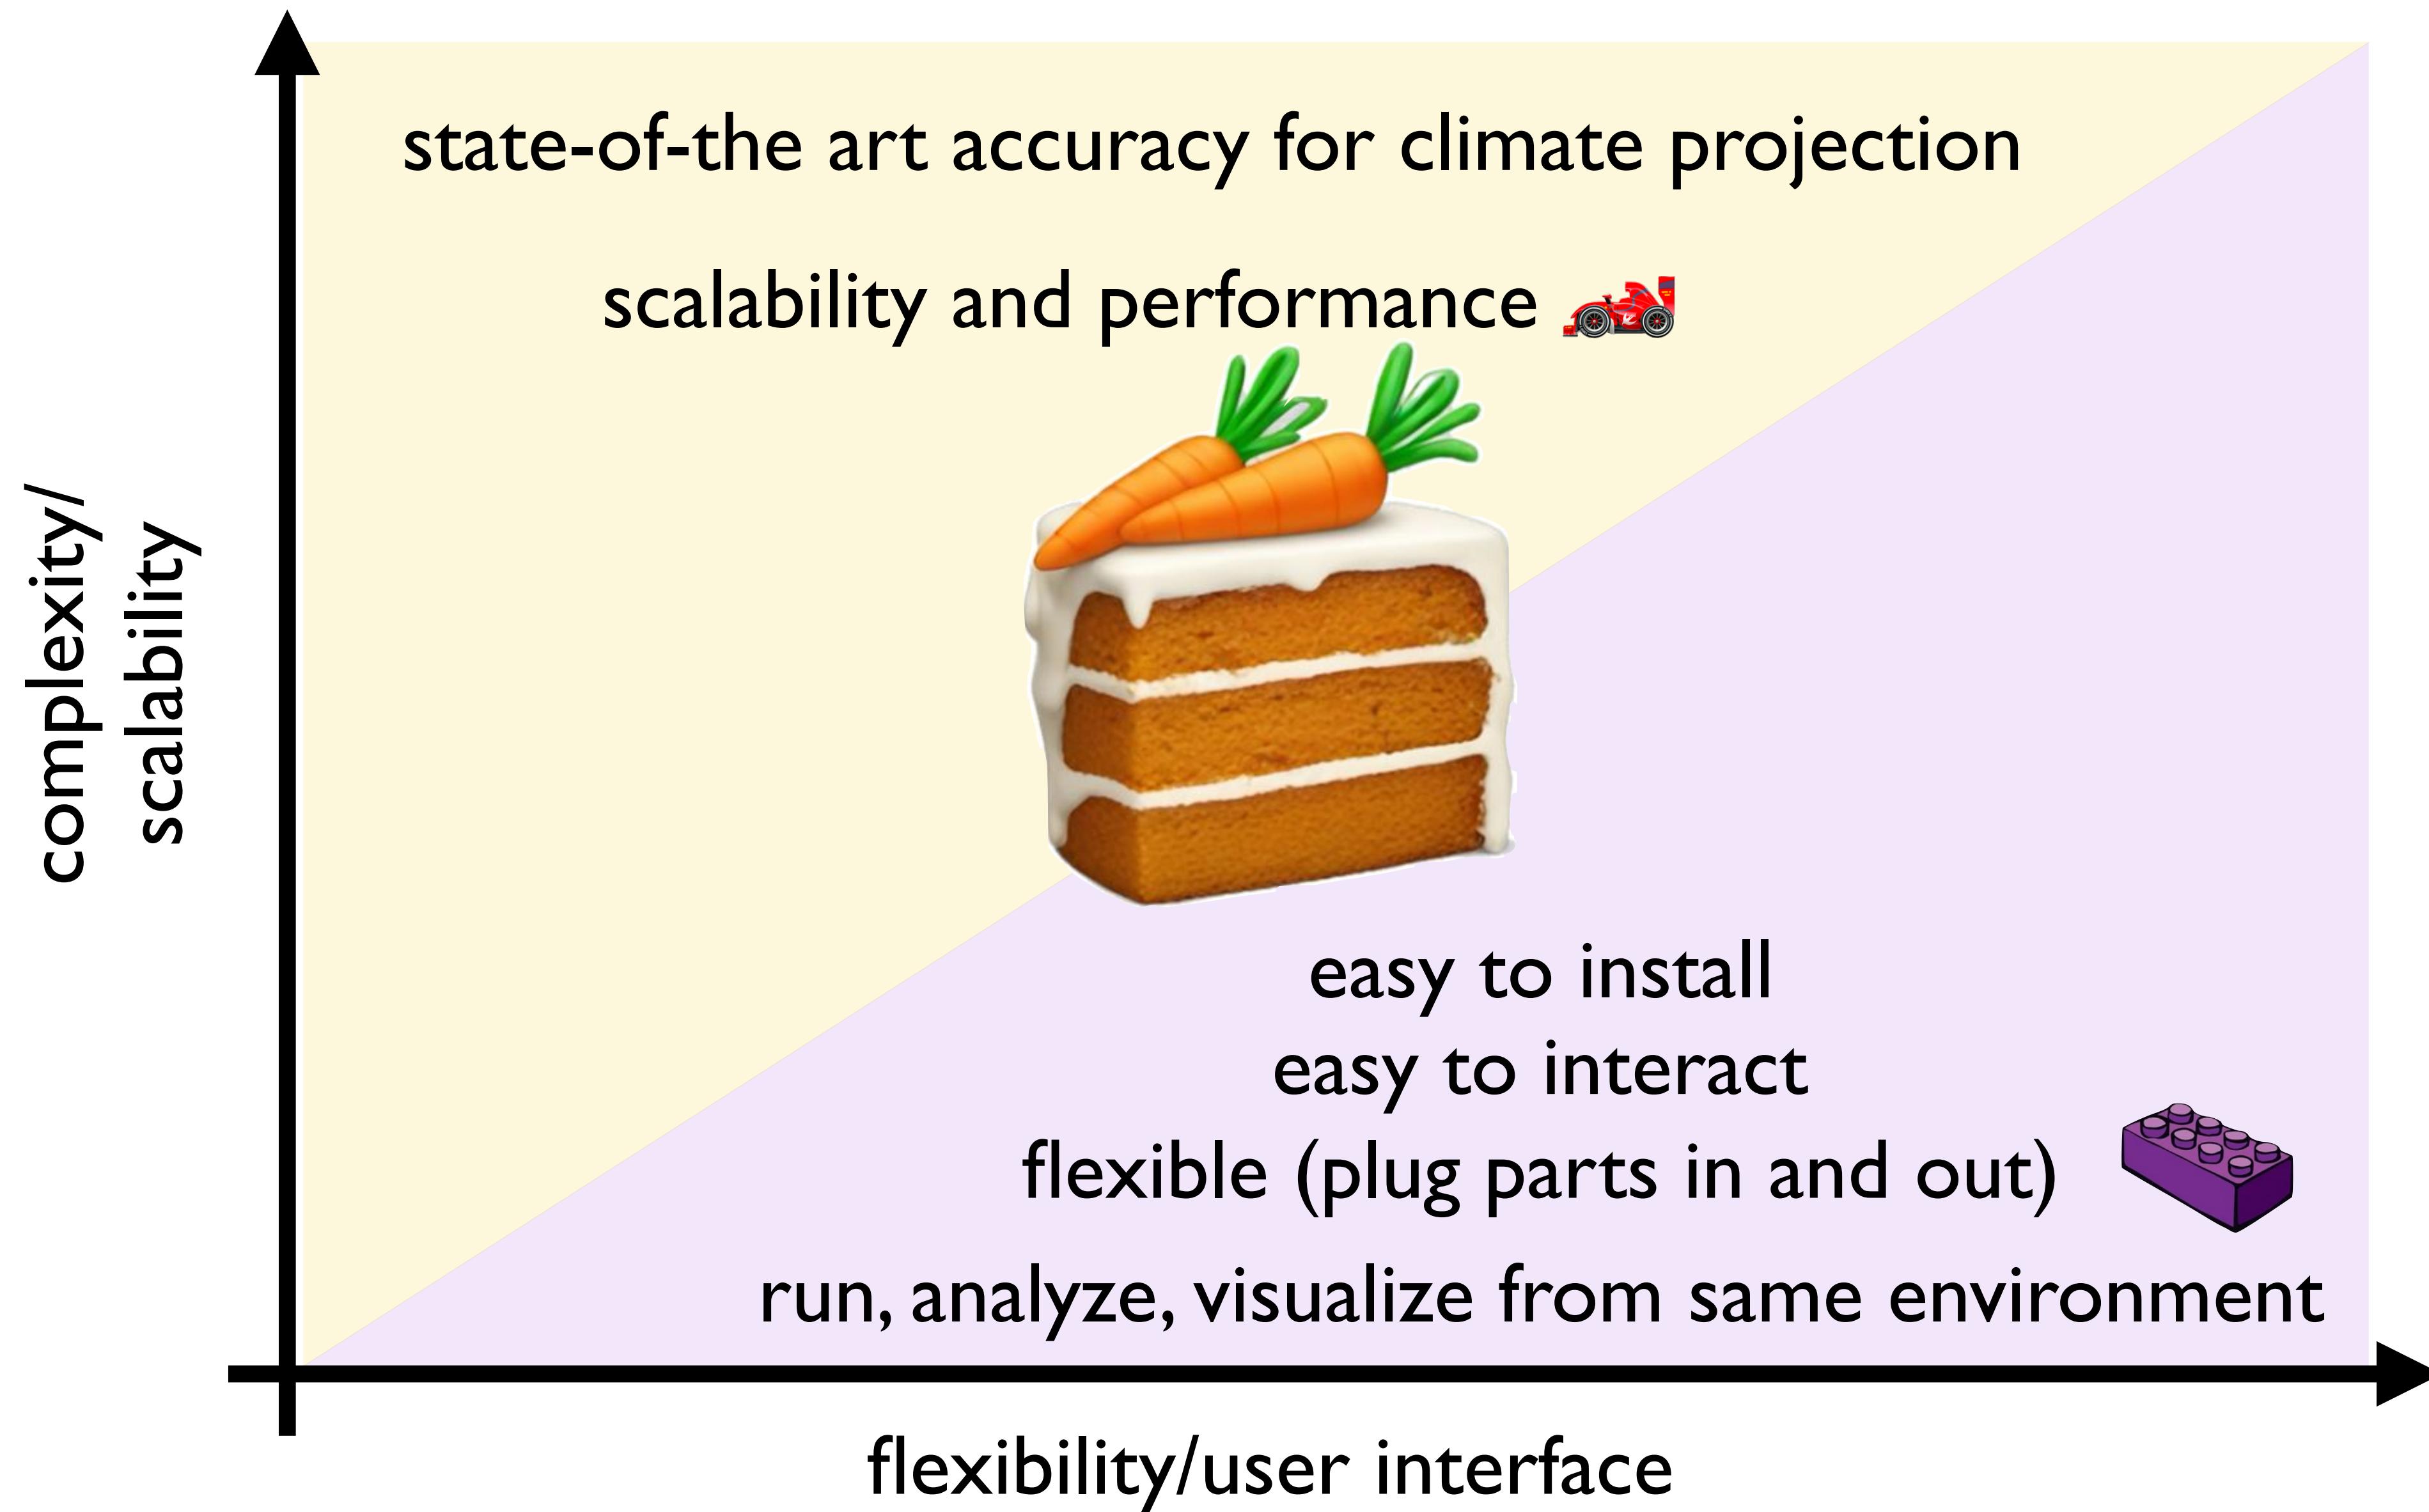


## ocean model



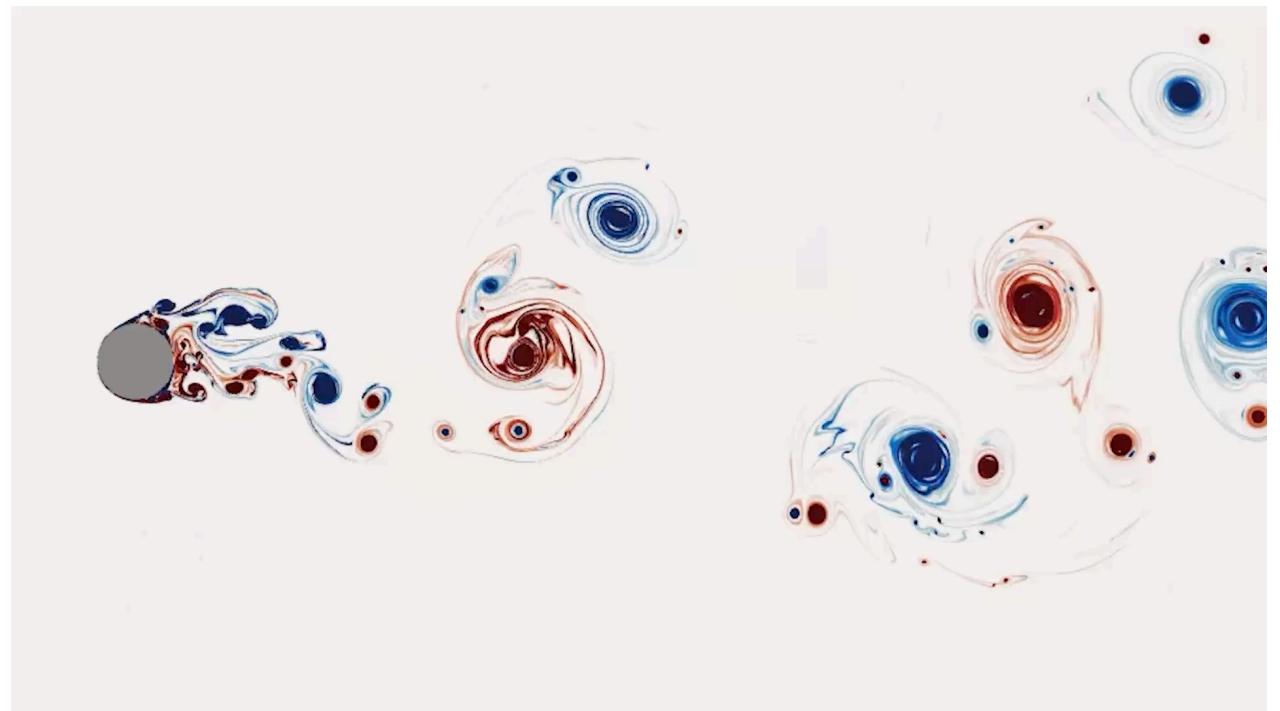


# ocean model



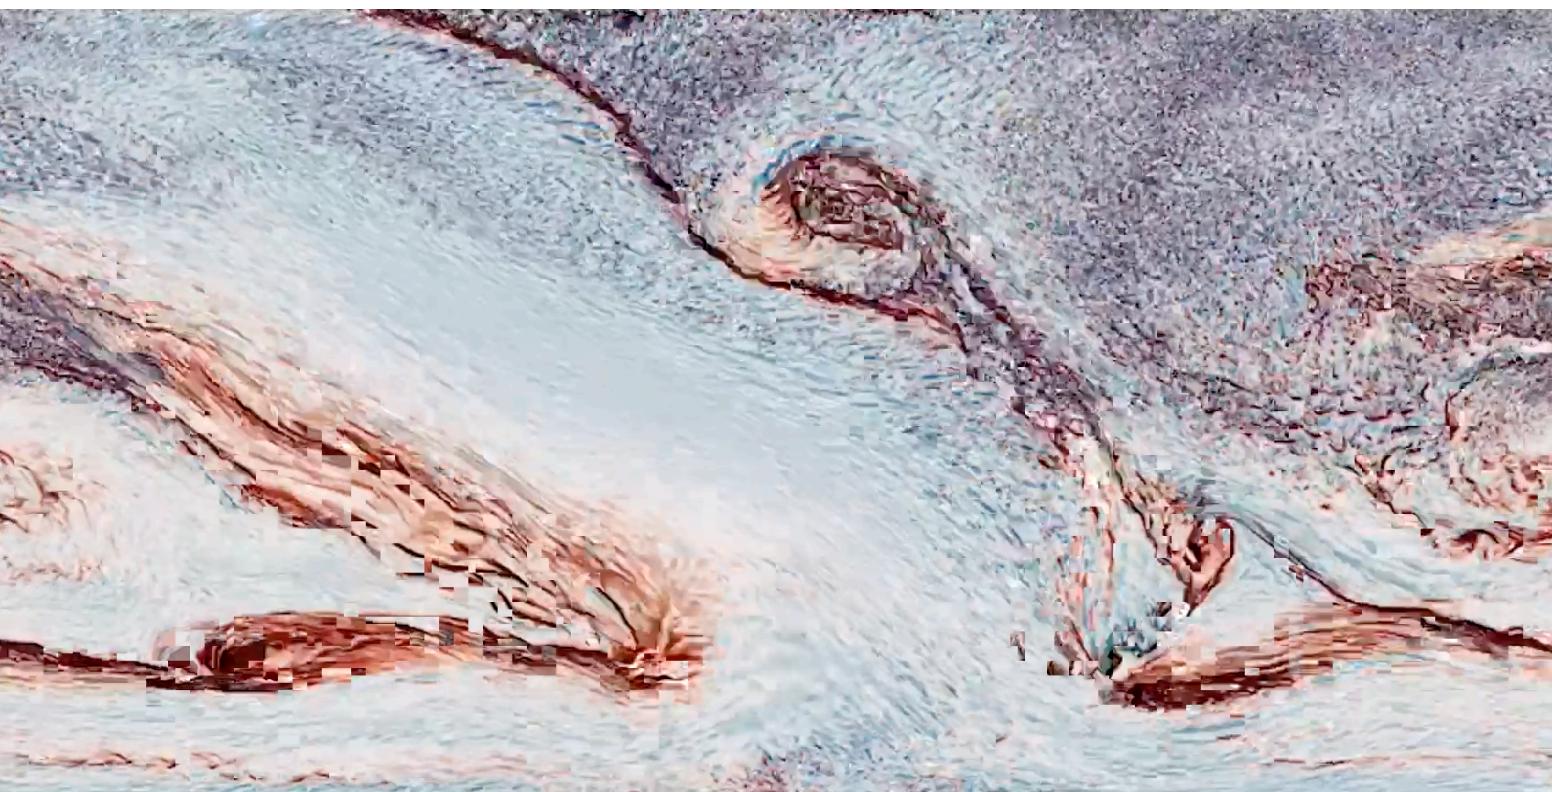


direct numerical simulation  
(DNS)



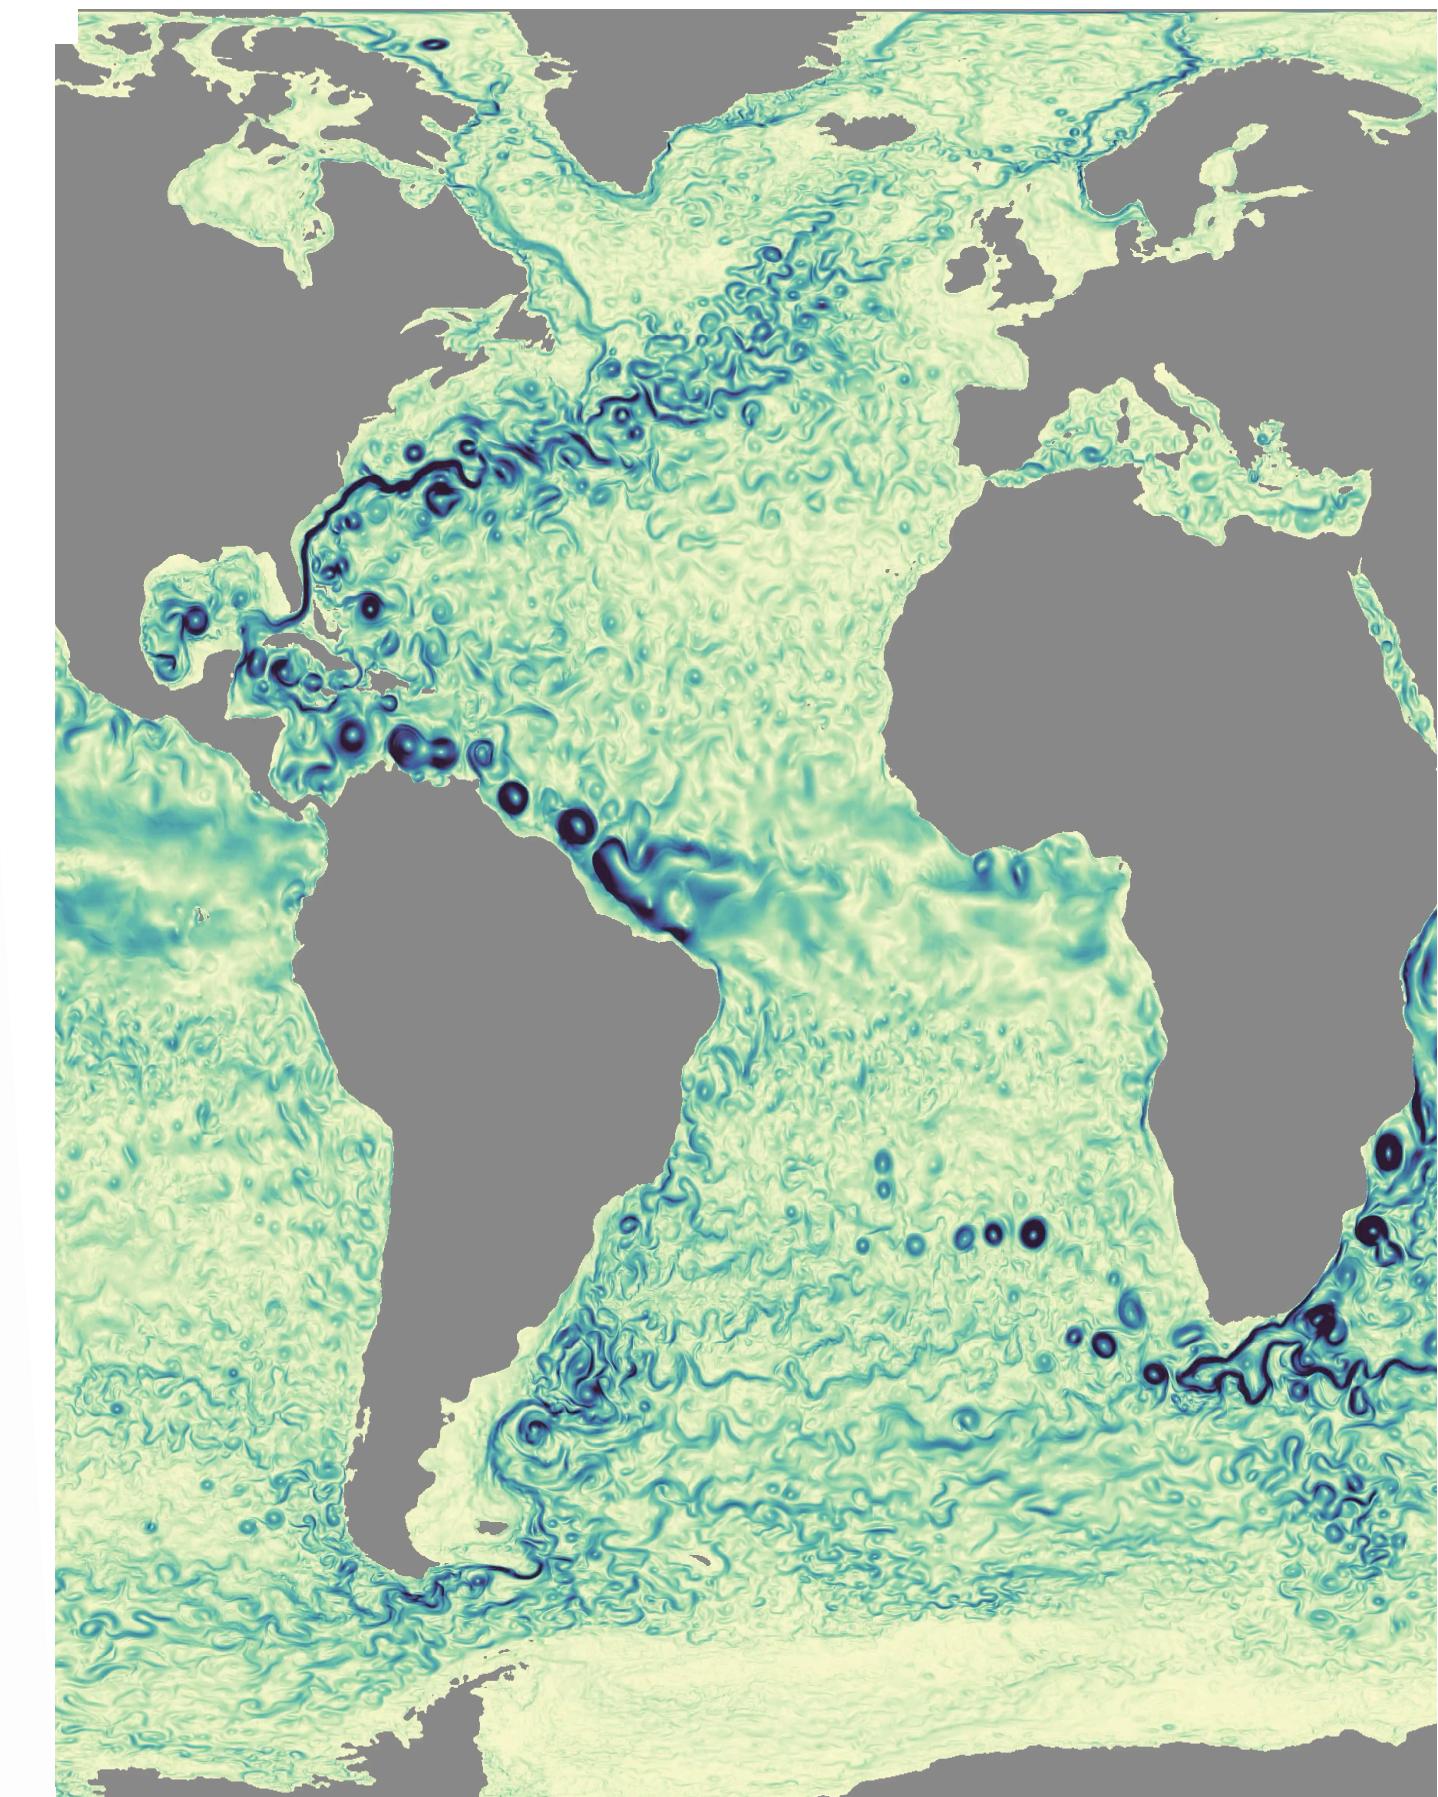
DNS of 2D flow around a cylinder

large eddy simulation  
(LES)



mixed-layer instability LES

general circulation modeling  
(GCM)

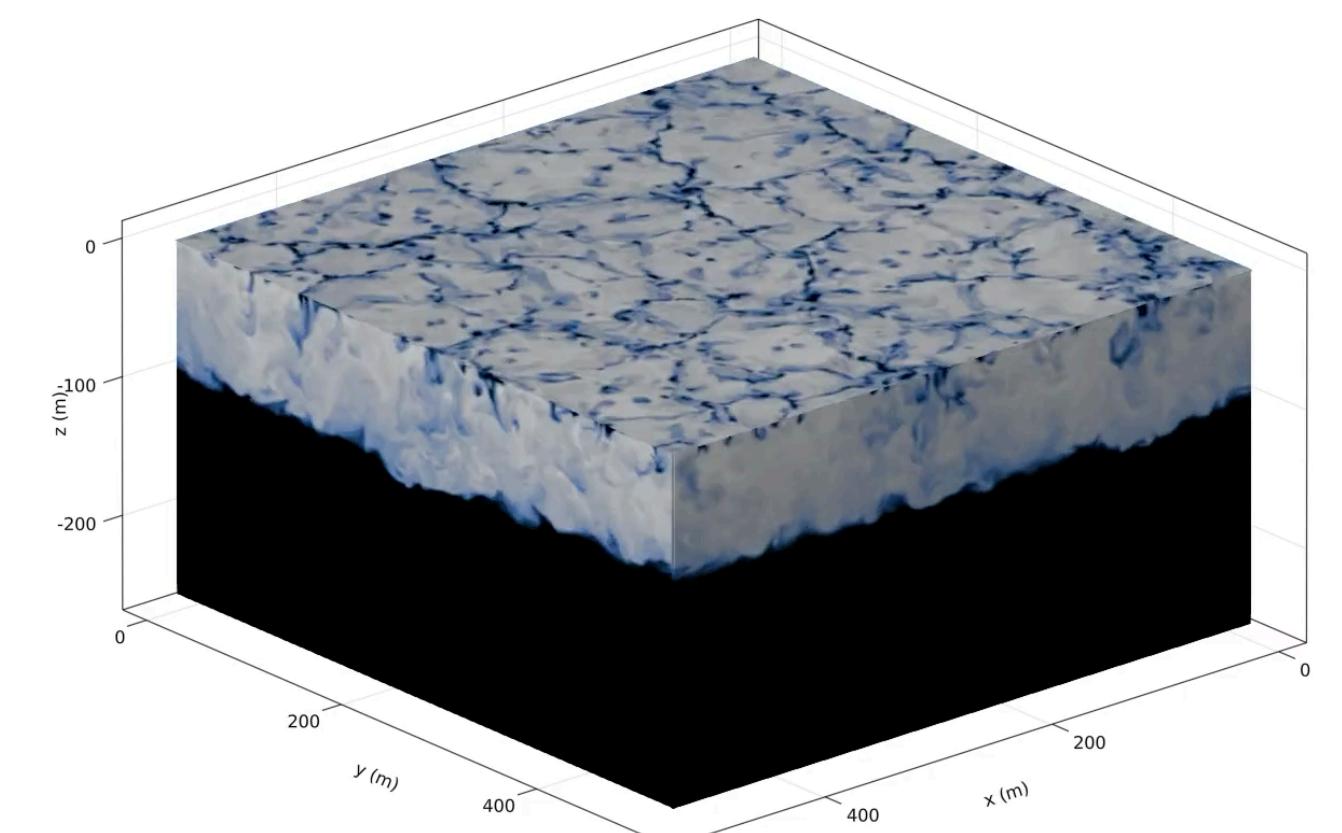


near-global simulation

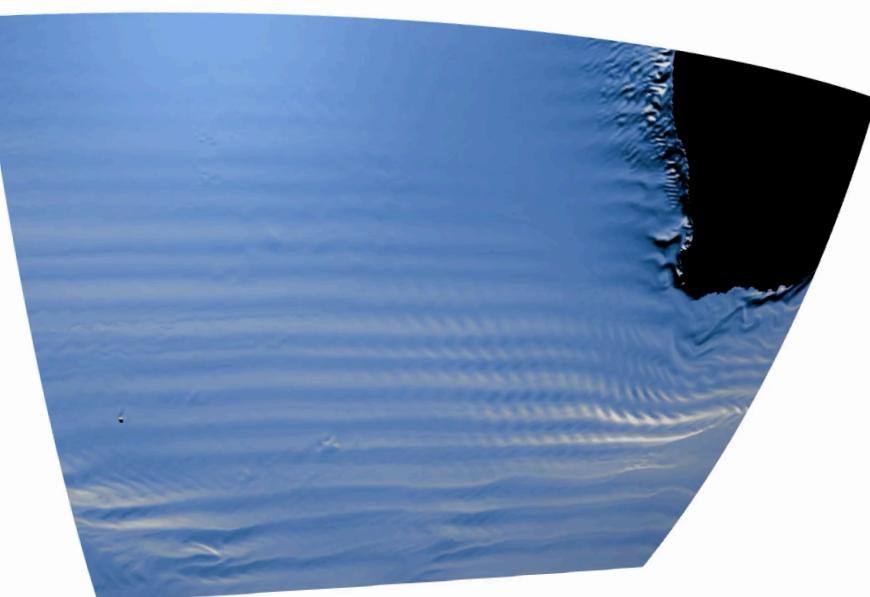
@ 1/12th degree horizontal resolution

DNS of cabbeling in freshwater

by ClimaOcean dev team  
(mostly Greg Wagner & Simone Silvestri)



Free convection LES



regional hydrostatic simulation



# Oceananigans in a nutshell



- Software for ocean-flavored fluid dynamics using finite-volume numerical methods

- **Fast:** written from scratch for GPUs



- **Friendly:** uses the Julia programming language

- ▶ *Simple simulations are easy*

- ▶ *Complex, creative simulations are possible*

- **Flexible**

- ▶ *Companion packages for biogeochemistry, sea ice, ocean-sea-ice coupled simulations*



Greg Wagner



Simone Silvestri



Many Contributors

I.

## decaying 2D turbulence

II.

## baroclinic instability on the sphere

III.

## near-global ocean simulation with continents, realistic atmosphere, and whatnot



## two\_dimensional\_turbulence.jl

```
grid = RectilinearGrid(size=(2048, 2048),
                       x=(0, 2π), y=(0, 2π),
                       topology = (Periodic, Periodic))

model = NonhydrostaticModel(; grid)

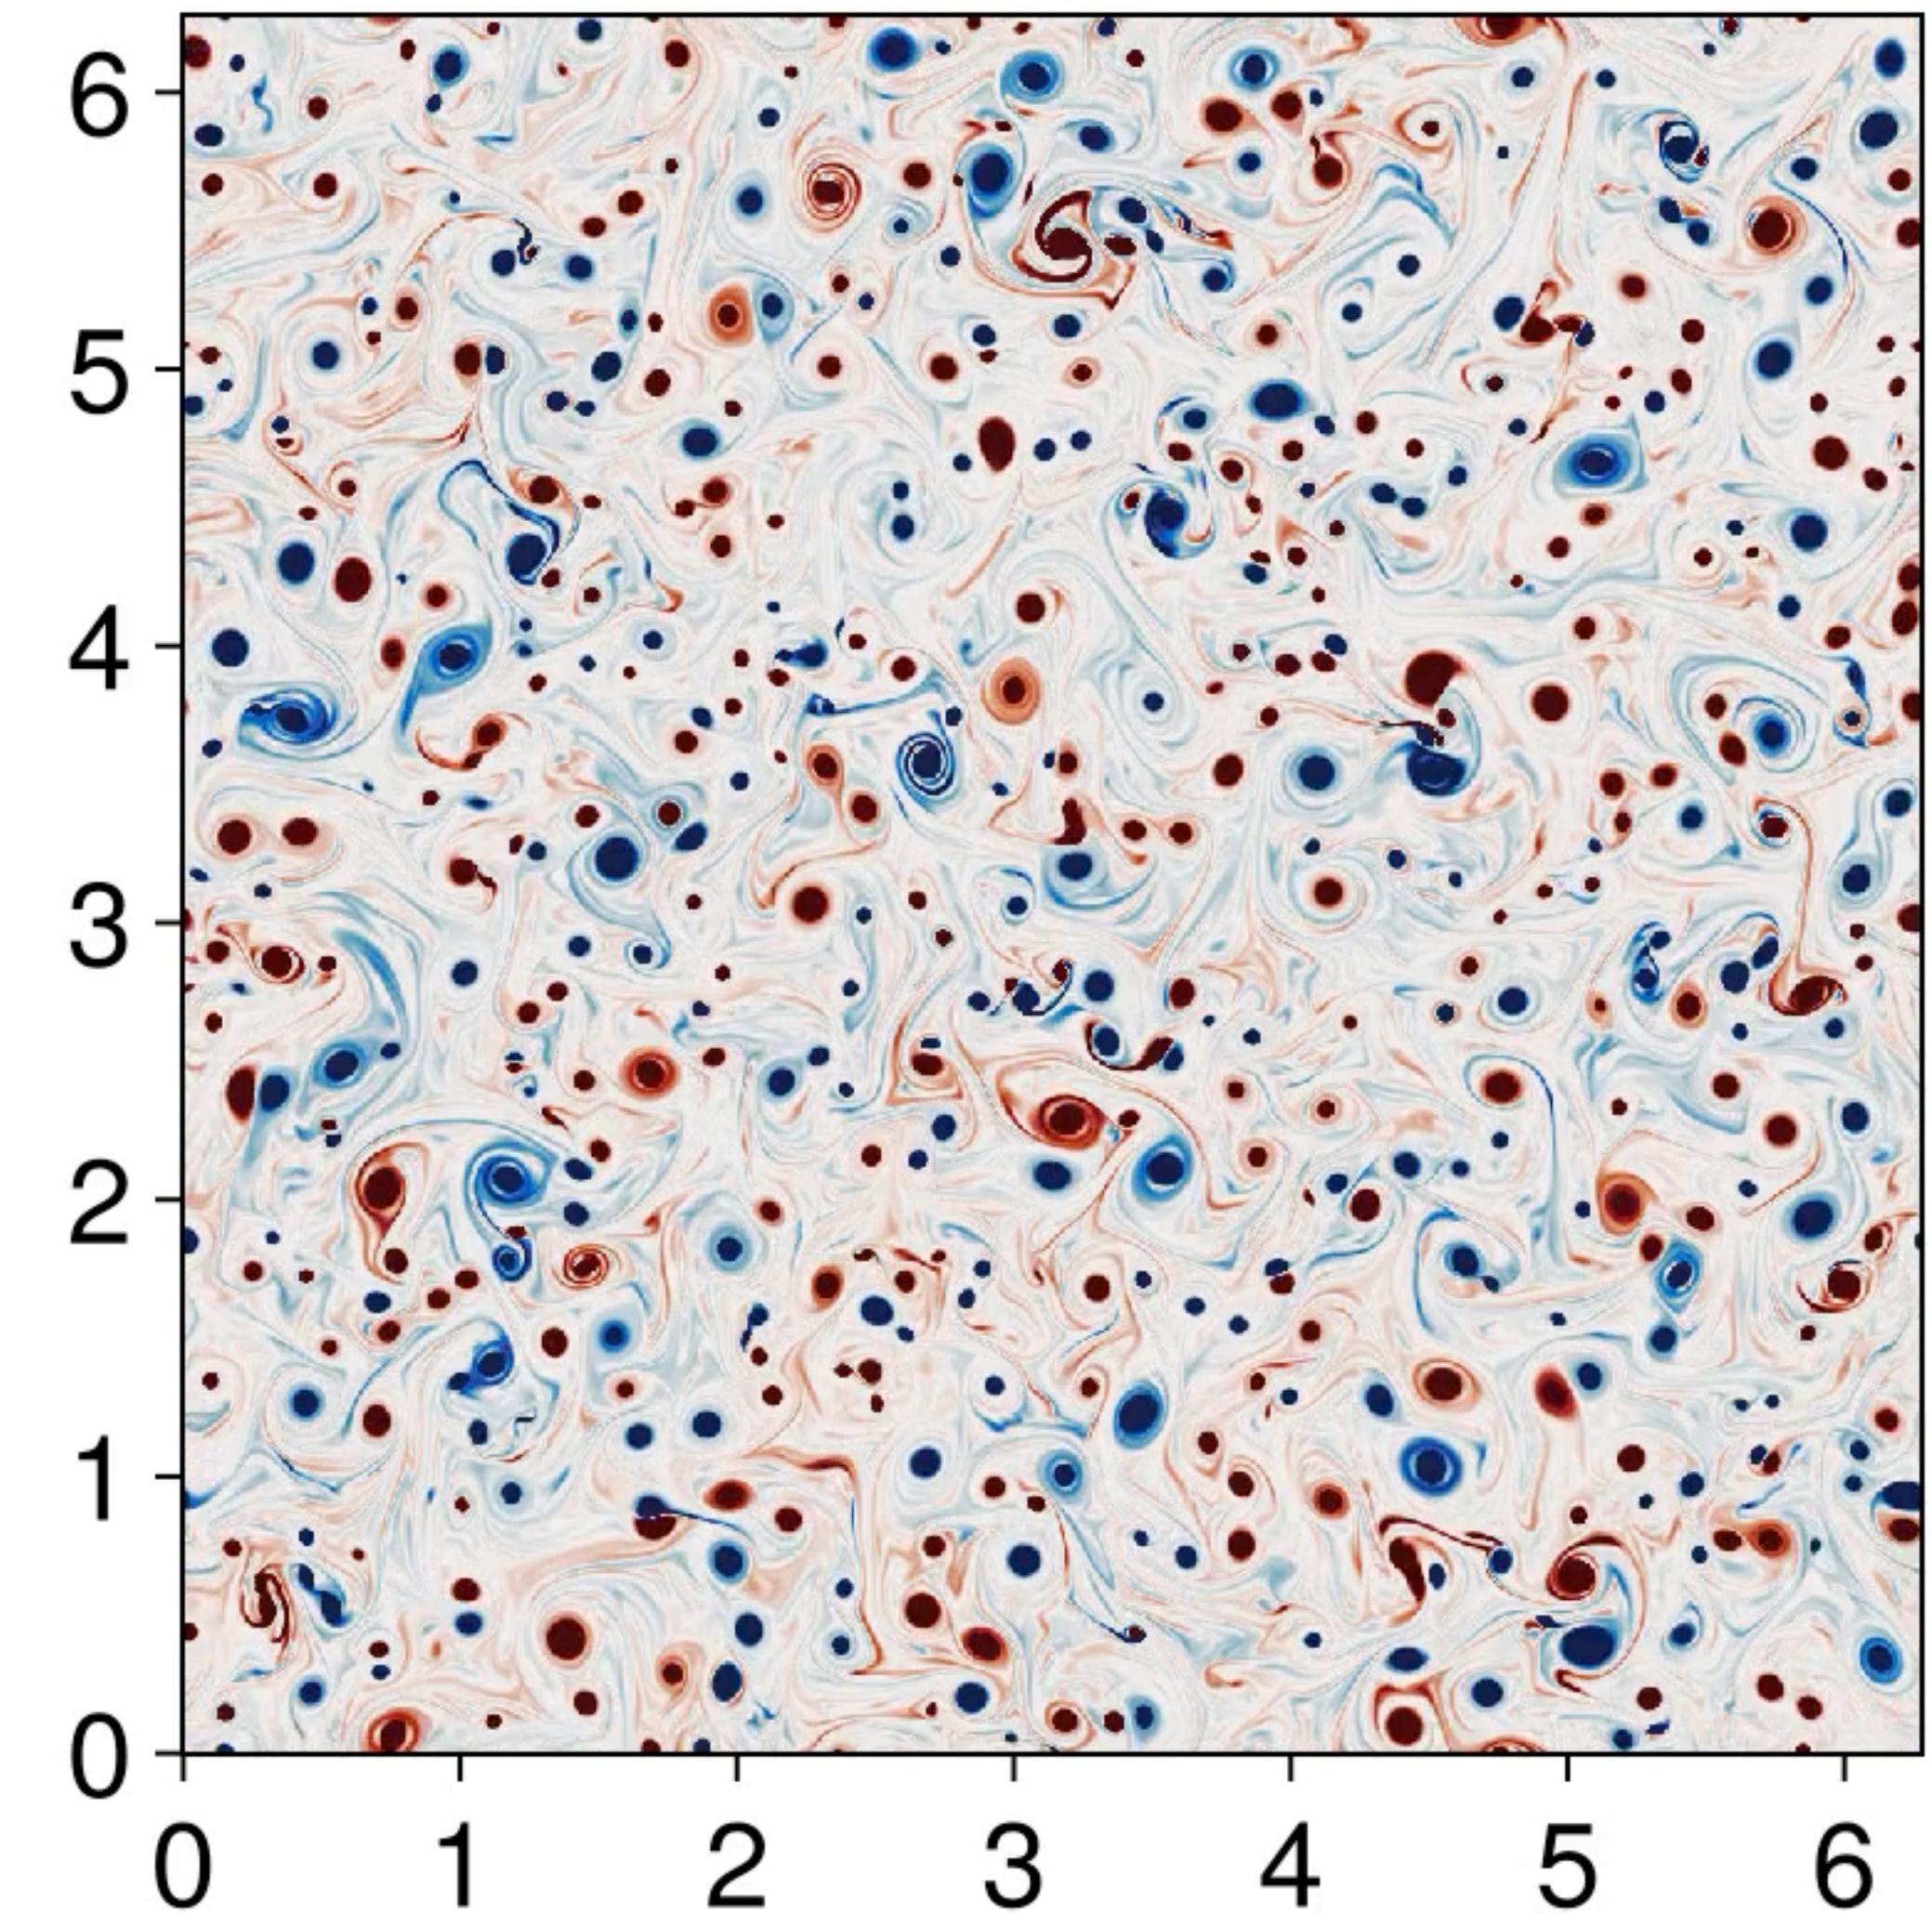
ε(x, y) = 2 * rand() - 1 # Uniformly-distributed random numbers ∈ [-1, 1]

set!(model, u=ε, v=ε)

u, v = model.velocities
ζ = ∂x(v) - ∂y(u)

heatmap(ζ)
```

**vorticity, t = 40.0**



## baroclinic\_turbulence\_sphere.jl

```
Nx, Ny, Nz = 720, 160, 10 # 1/2 degree horizontal resolution

grid = LatitudeLongitudeGrid(GPU()); size = (Nx, Ny, Nz),
      latitude = (-80, 80),
      longitude = (0, 360),
      z = (-3000, 0))

buoyancy = SeawaterBuoyancy(equation_of_state=TEOS10EquationOfState())

model = HydrostaticFreeSurfaceModel(; grid, buoyancy,
                                    coriolis = HydrostaticSphericalCoriolis(),
                                    tracers = (:T, :S),
                                    momentum_advection = WENOVectorInvariant(order=9),
                                    tracer_advection = WENO(order=7))

Ti(λ, φ, z) = 30 * (1 - tanh((abs(φ) - 45) / 8)) / 2 + rand() # initial temperature
Si(λ, φ, z) = 28 - 5e-3 * z + rand() # initial salinity

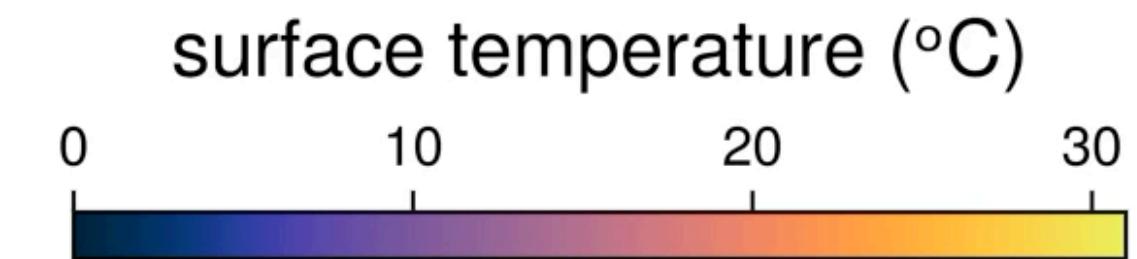
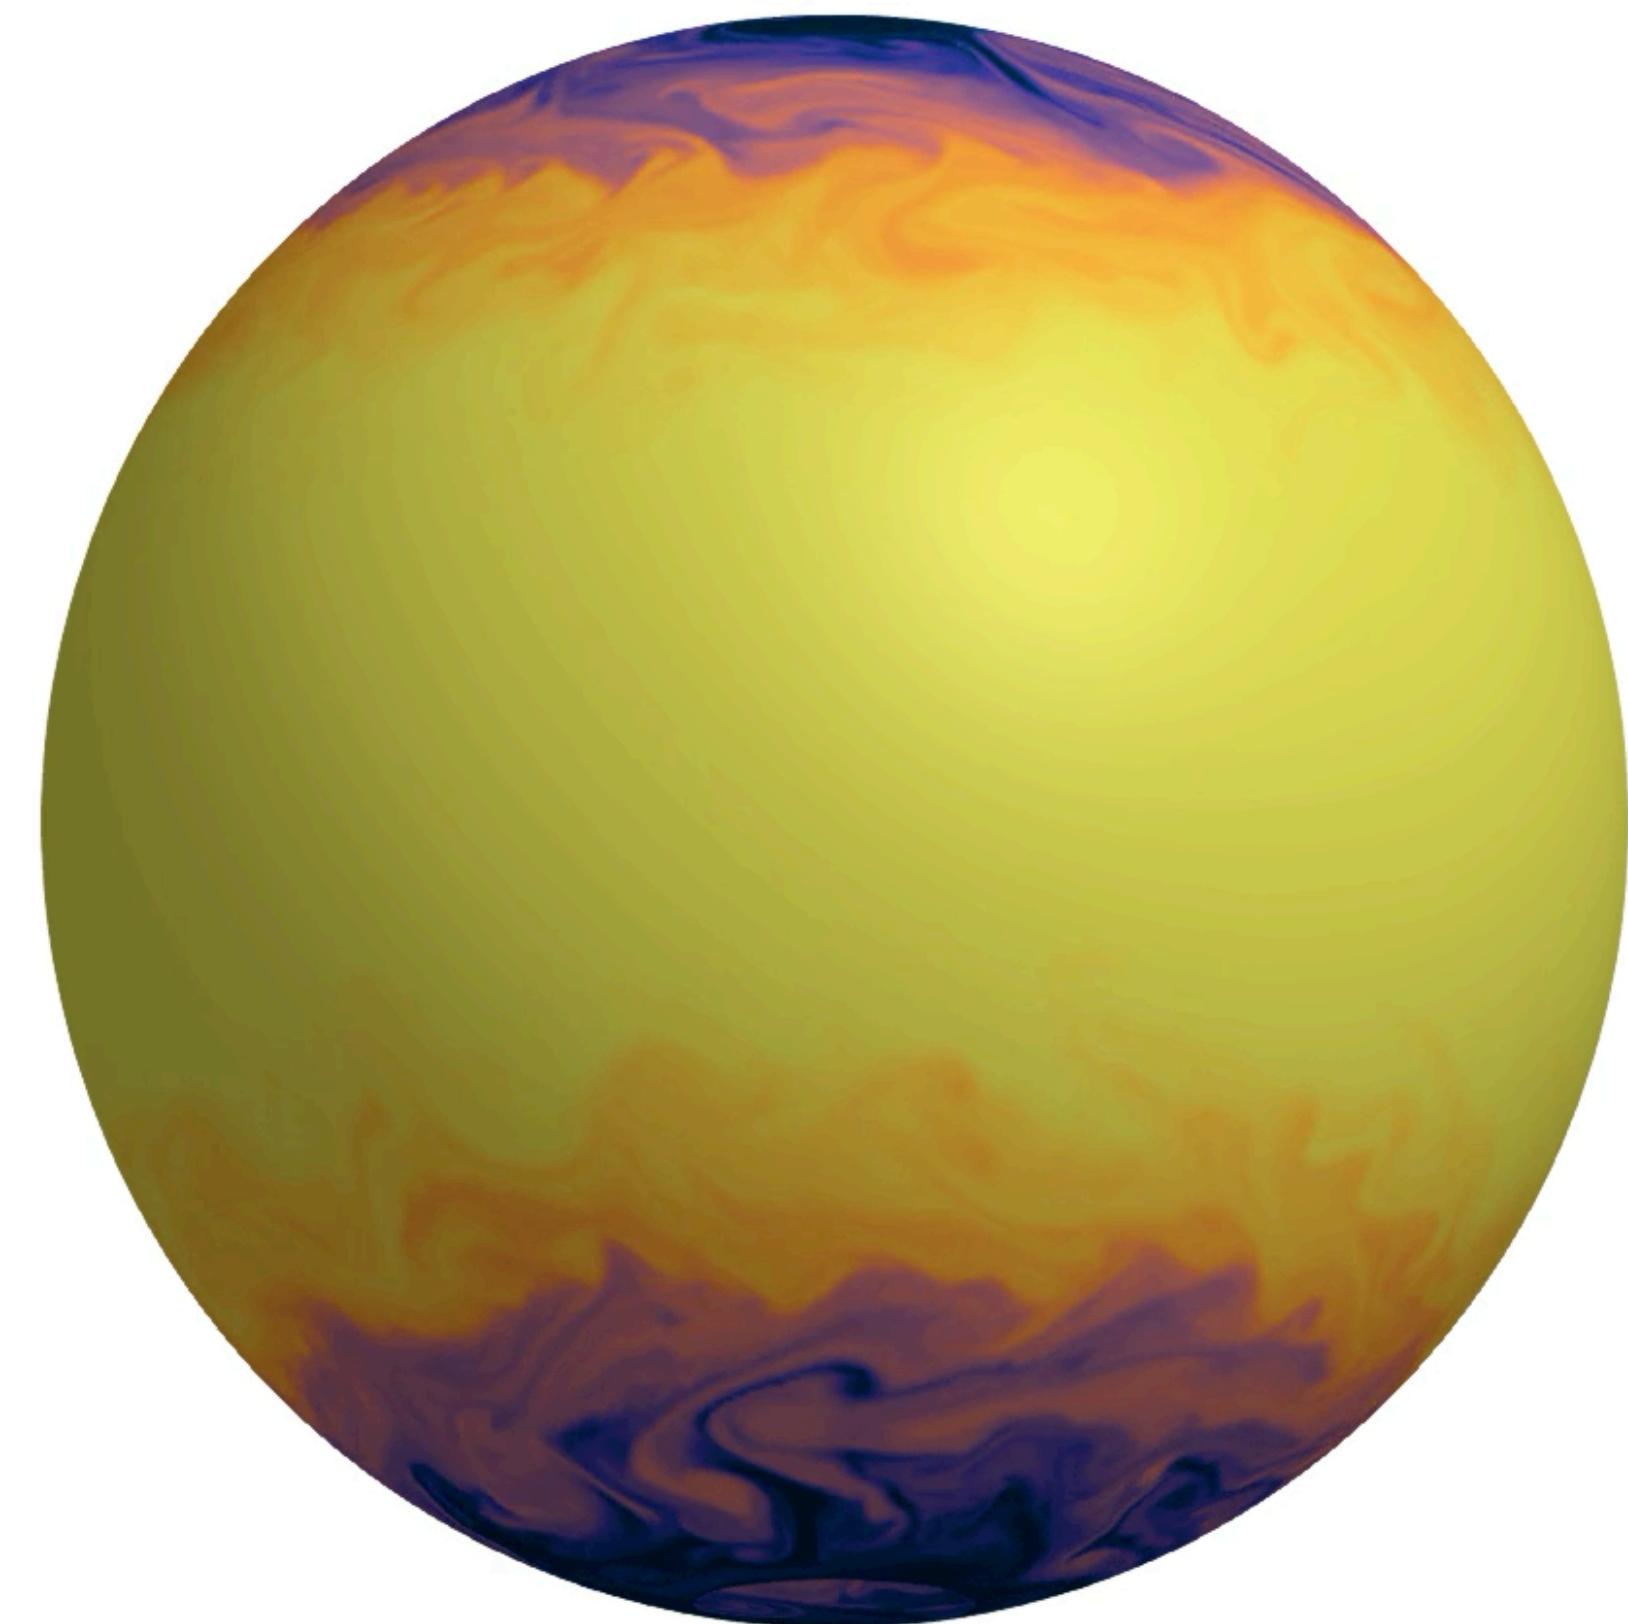
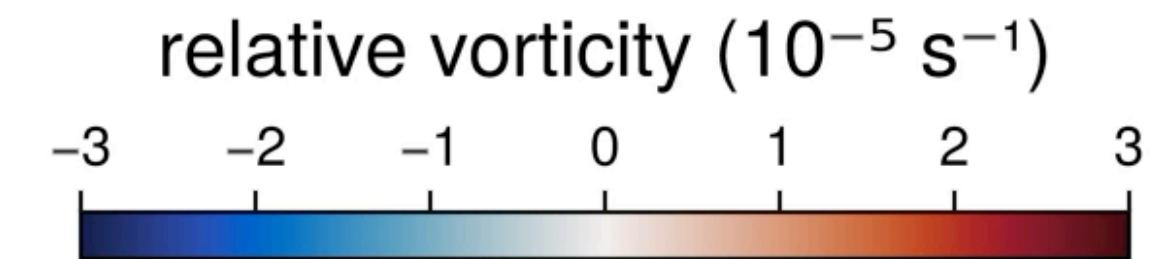
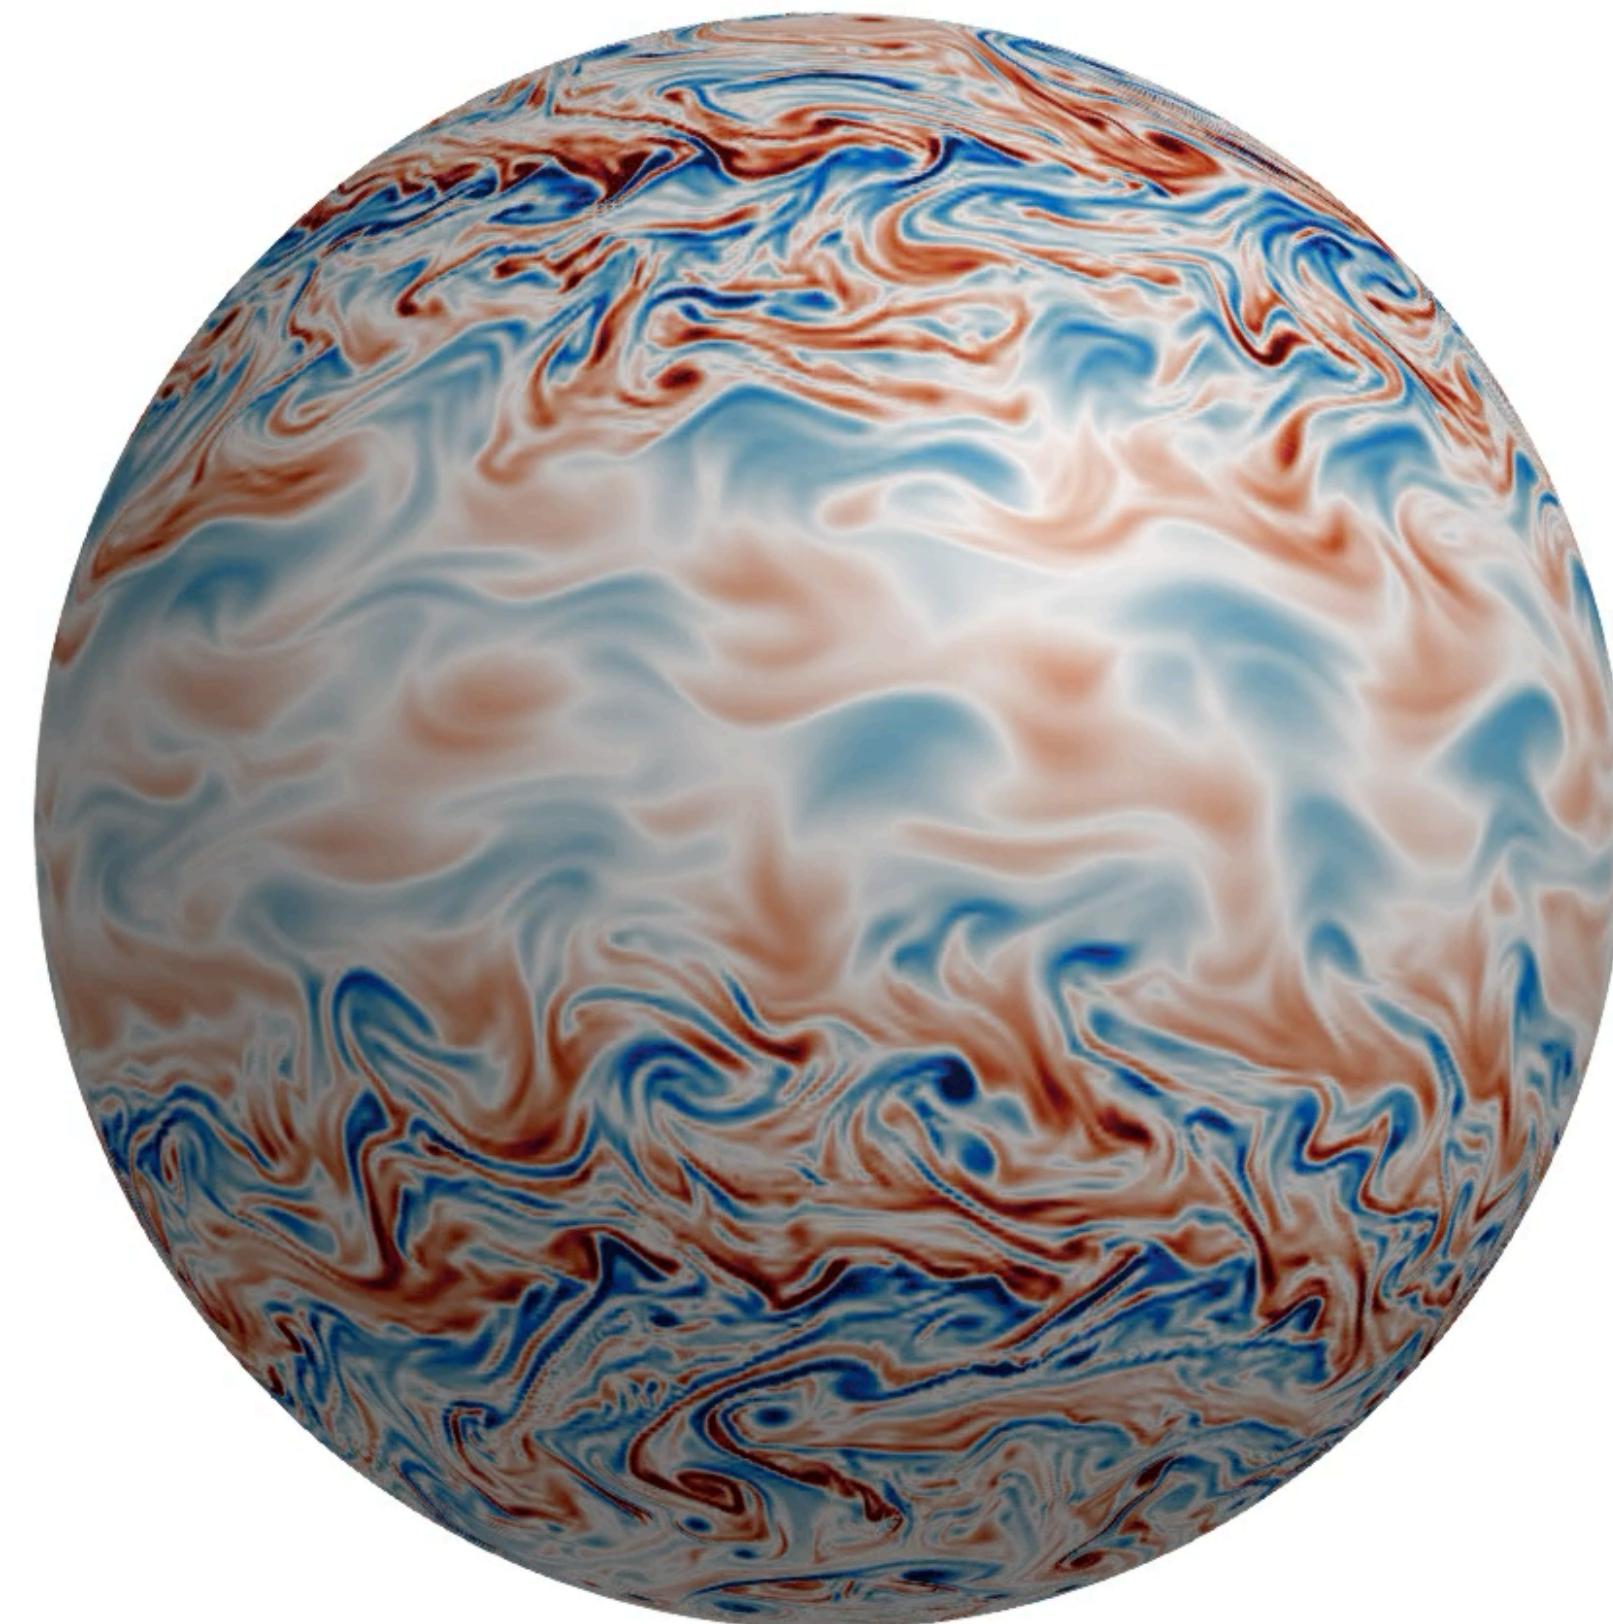
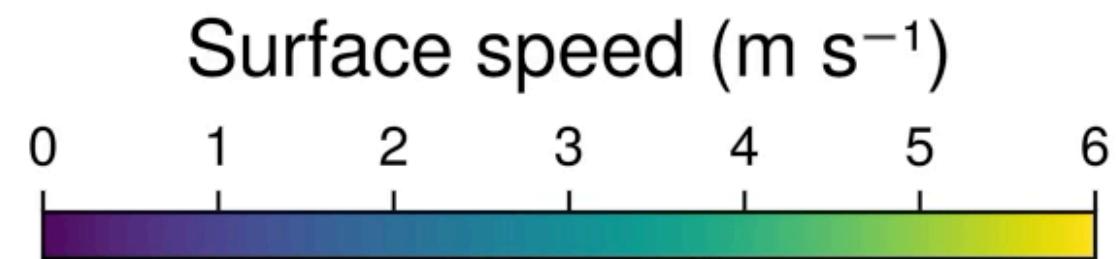
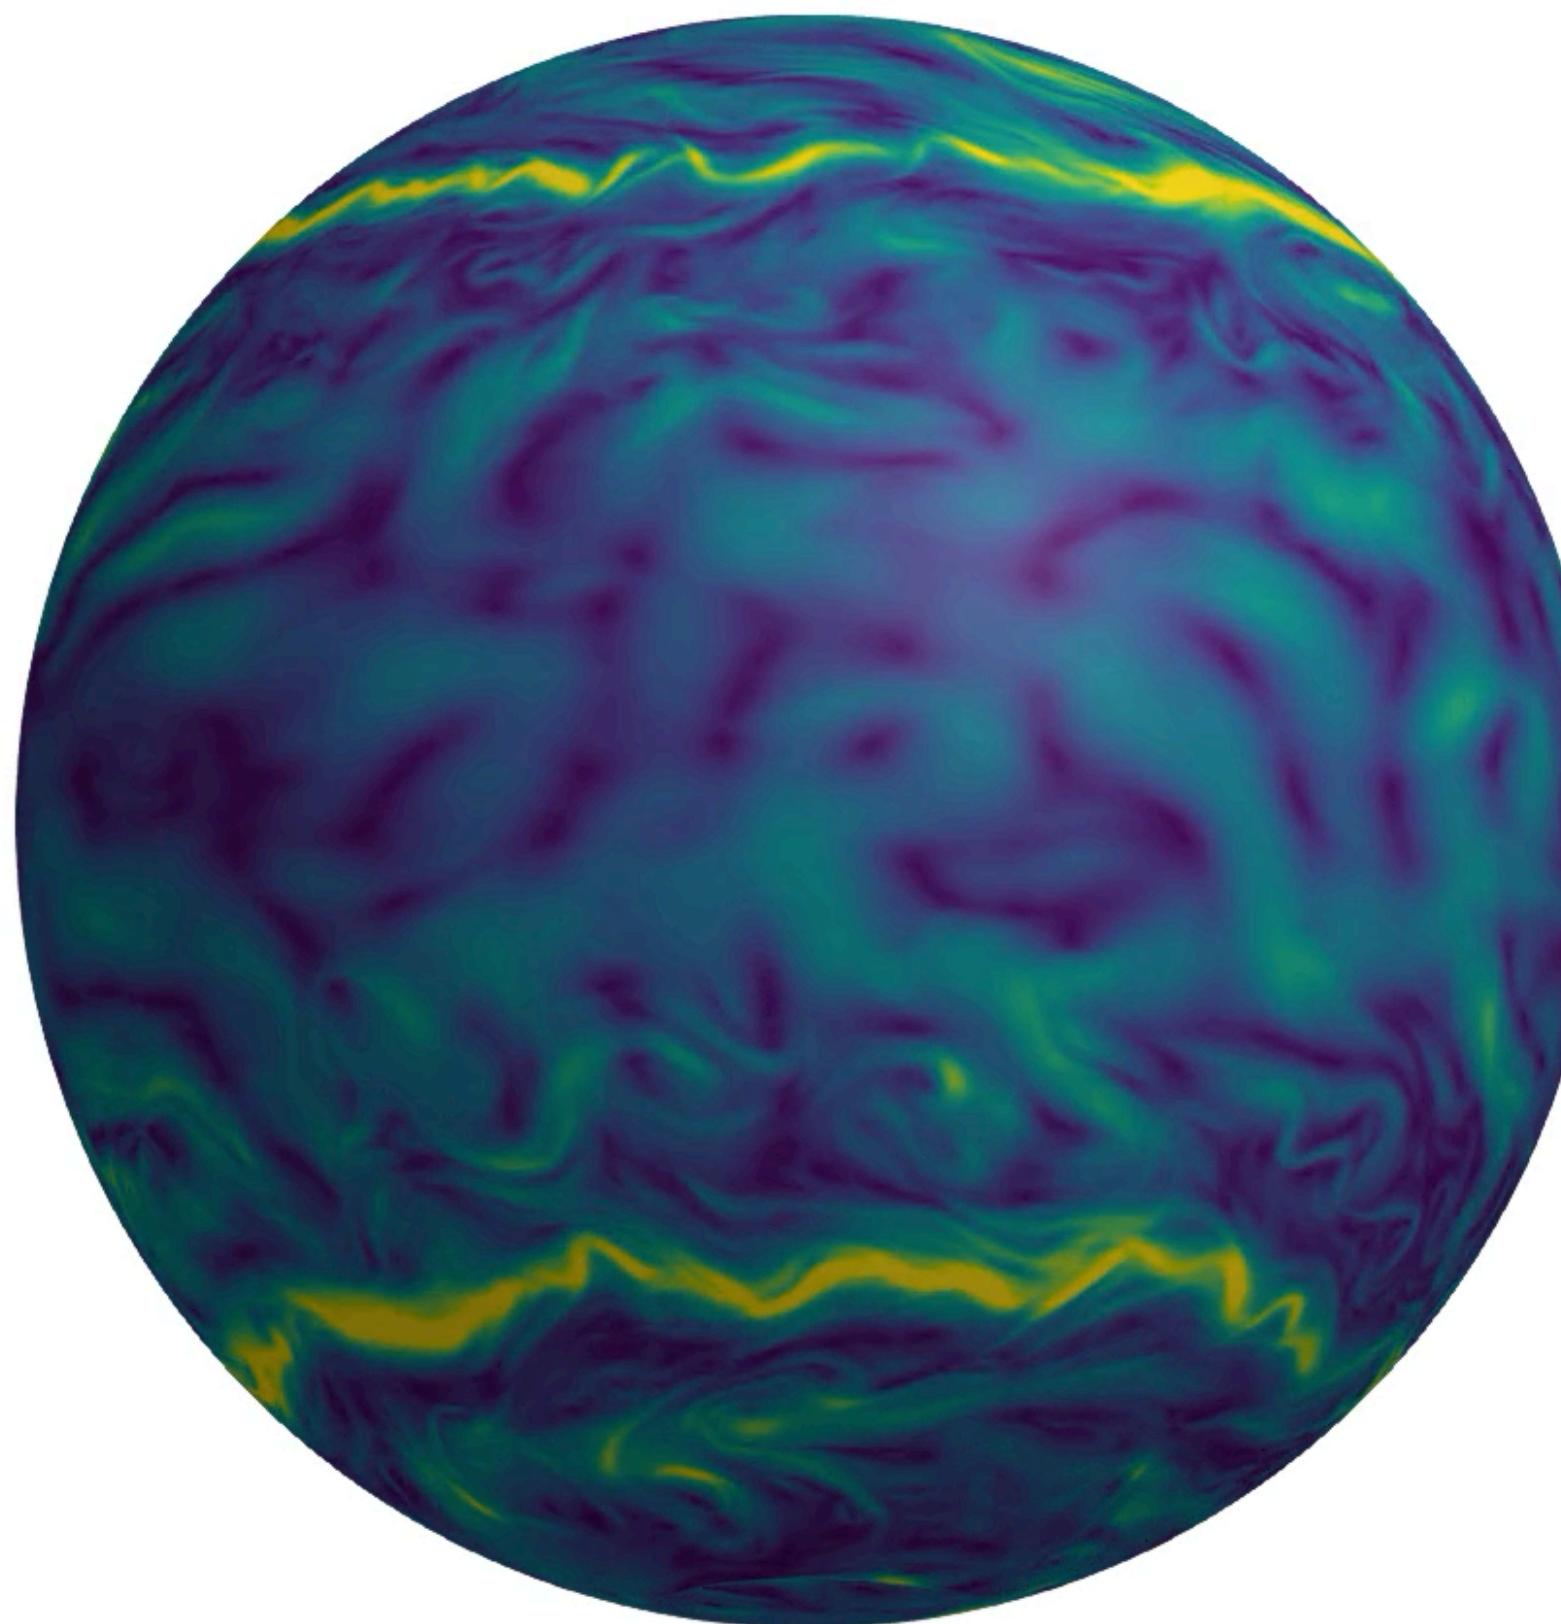
set!(model, T=Ti, S=Si)

simulation = Simulation(model, Δt=5minutes, stop_time=200days)

run!(simulation)
```



165.0 days



## near\_global\_ocean\_simulation.jl

```

Nx, Ny, Nz = 1440, 560, 10 # 1/4 degree horizontal resolution

grid = LatitudeLongitudeGrid(GPU()); size = (Nx, Ny, Nz),
                                         latitude = (-70, 70),
                                         longitude = (0, 360),
                                         z = (-3000, 0))

bathymetry = ClimaOcean.regrid_bathymetry(grid) # build gridded bathymetry based on ETOP01

grid = ImmersedBoundaryGrid(grid, GridFittedBottom(bathymetry))

# build an ocean simulation initialized to the ECCO state estimate on Jan 1, 1993
ocean = ClimaOcean.ocean_simulation(grid)
dates = DateTimeProlepticGregorian(1993, 1, 1)

set!(ocean.model, T = ClimaOcean.ECCOMetadata(:temperature; dates),
      S = ClimaOcean.ECCOMetadata(:salinity; dates))

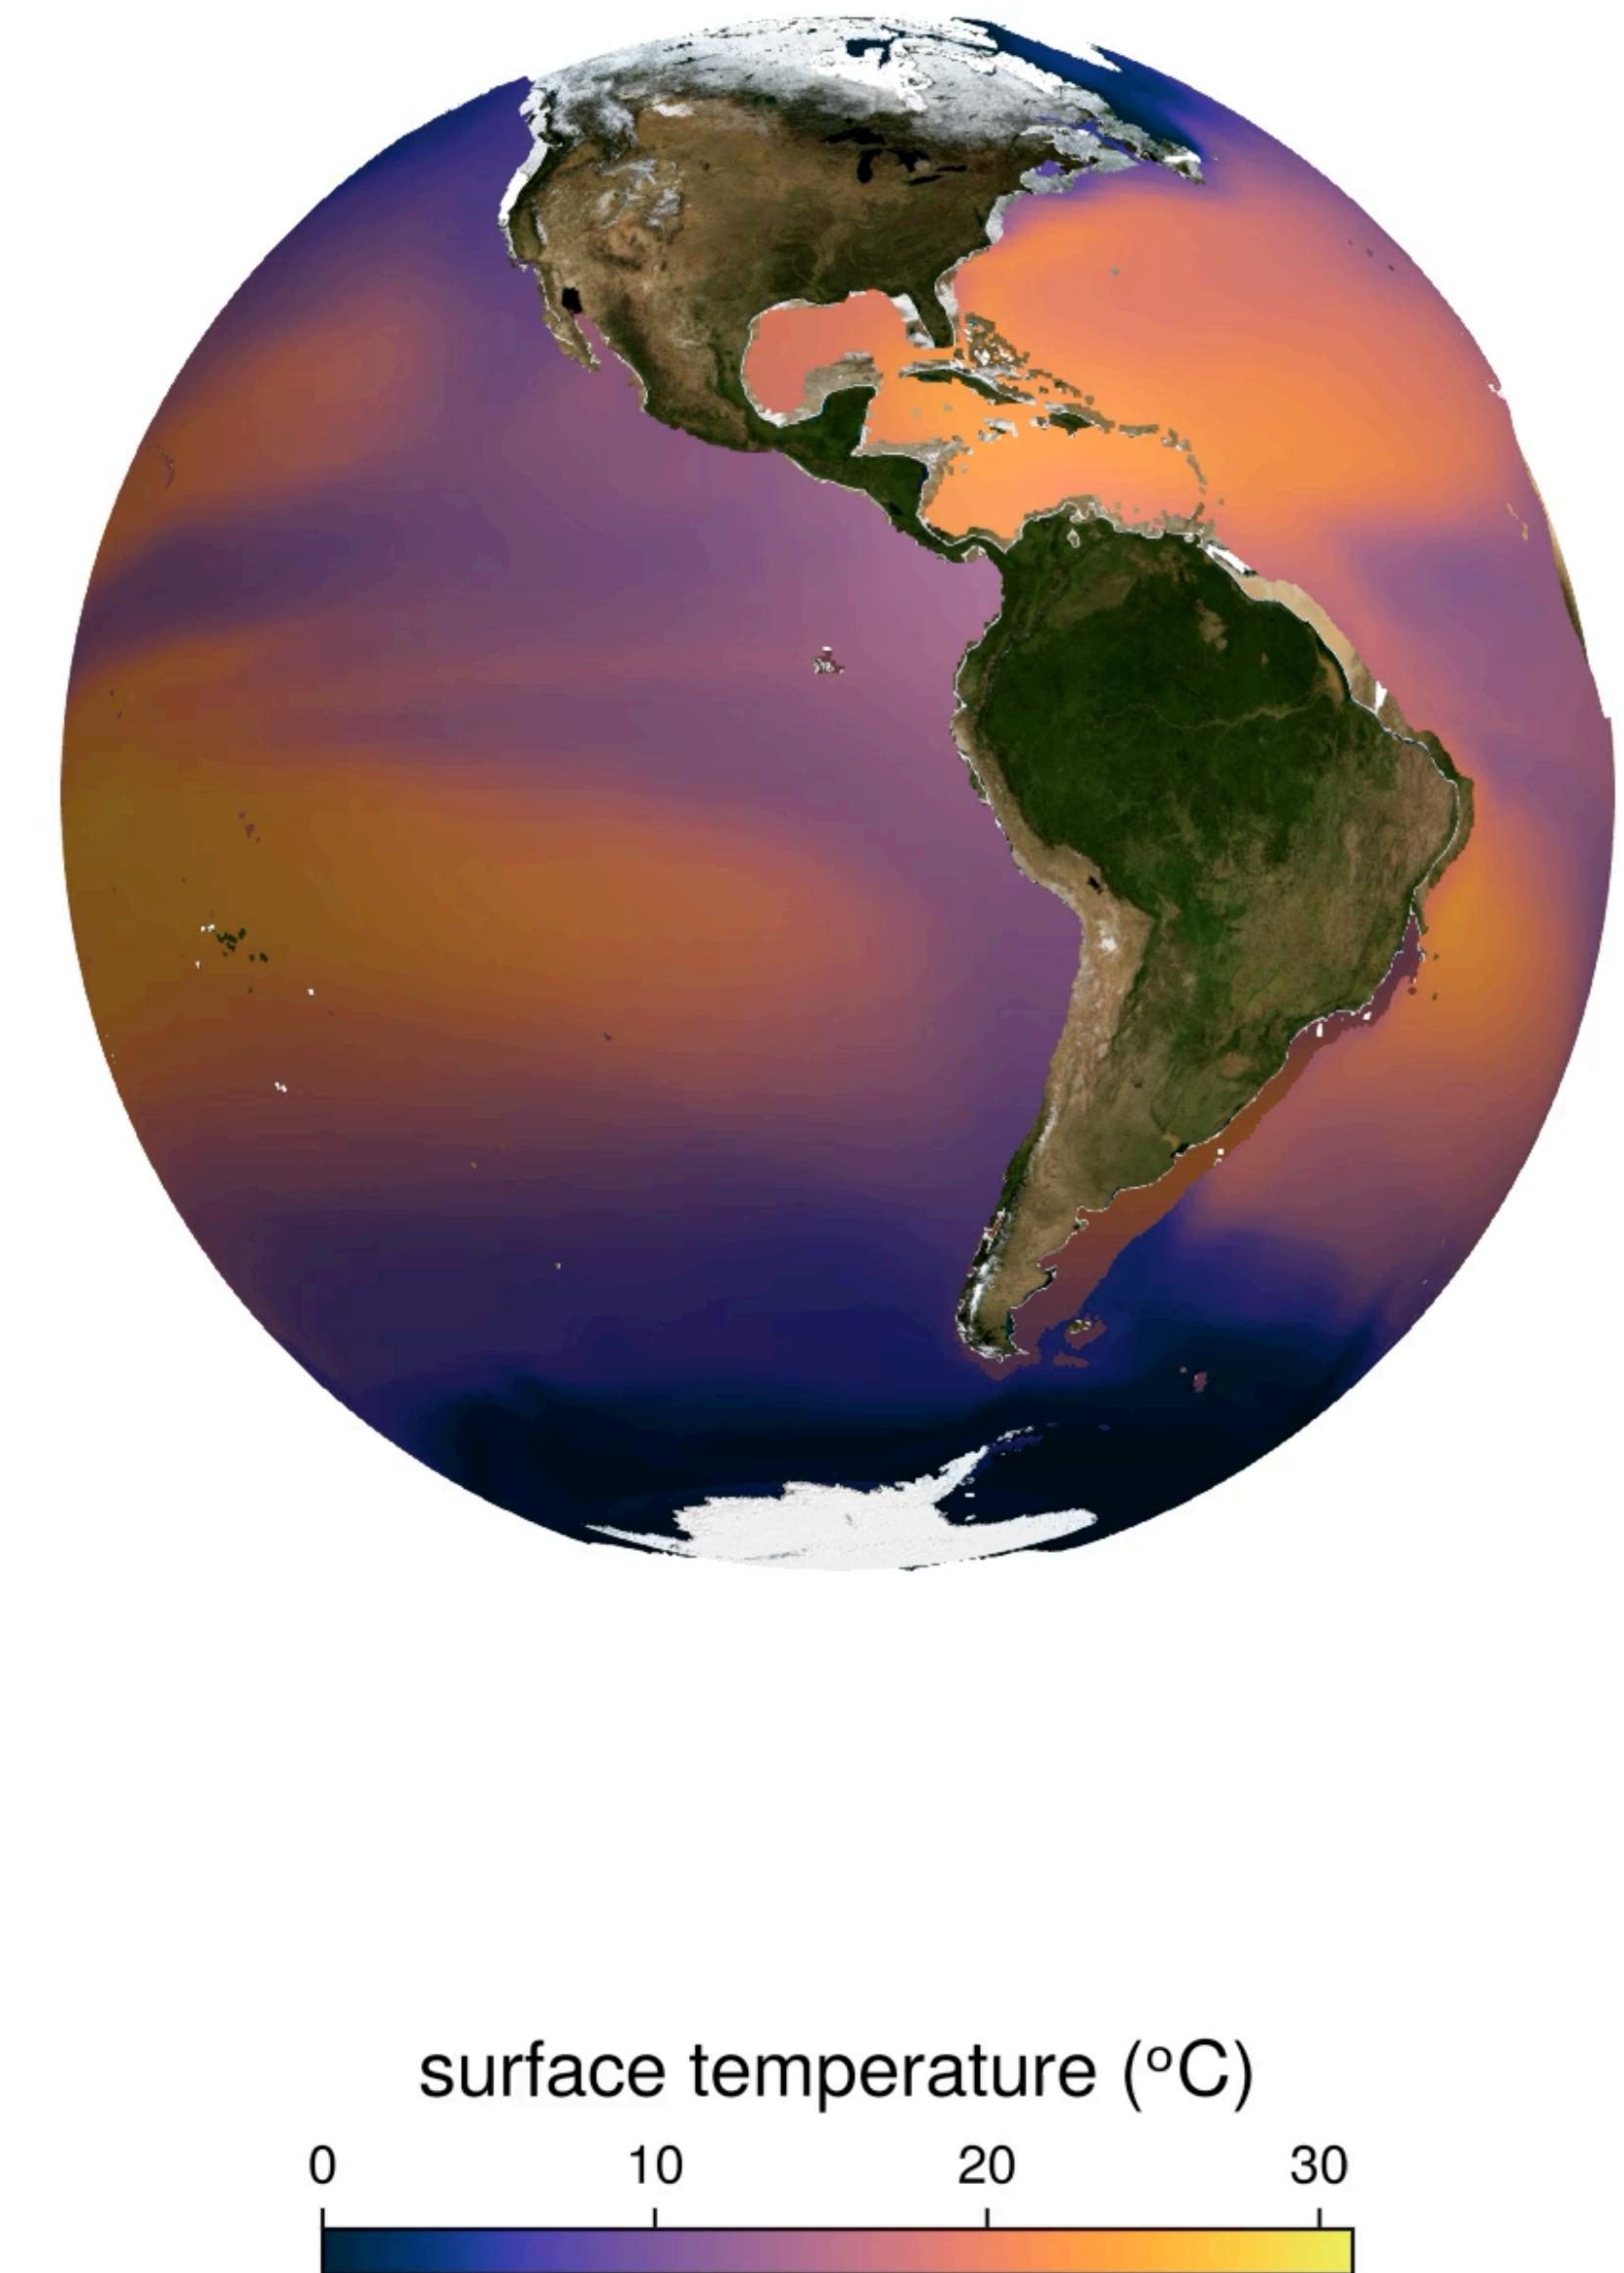
# prescribed atmosphere from JRA55 reanalysis
atmosphere = ClimaOcean.JRA55PrescribedAtmosphere(arch)

coupled_model = ClimaOcean.OceanSeaIceModel(ocean; atmosphere)

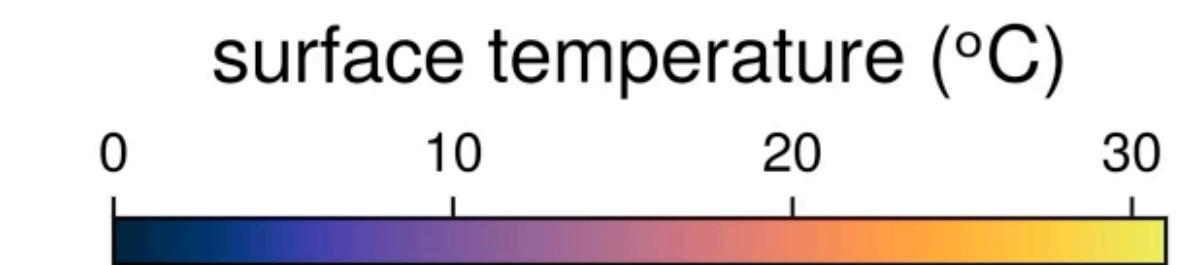
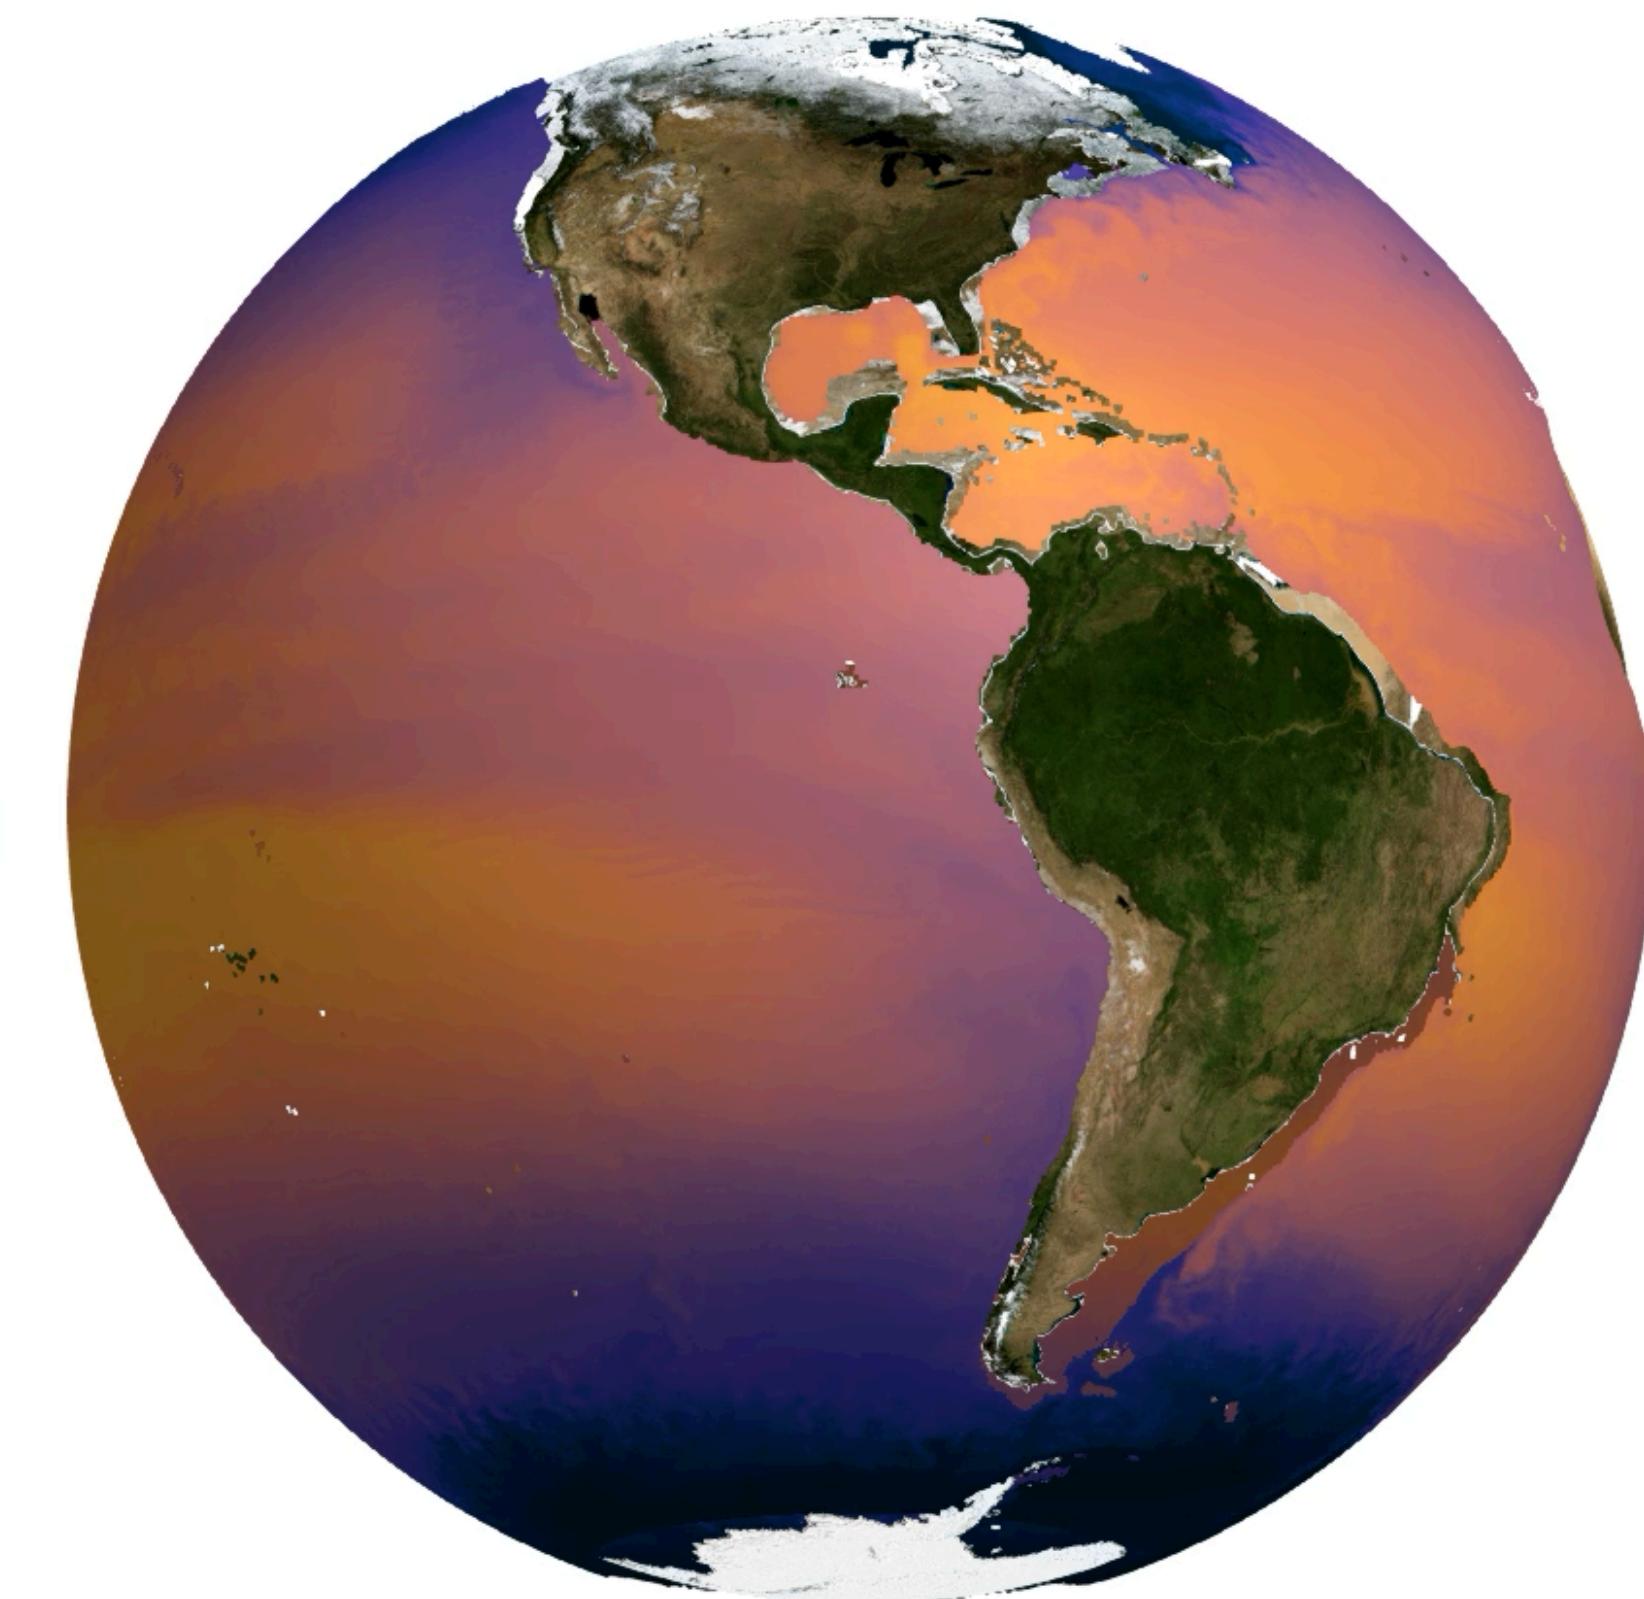
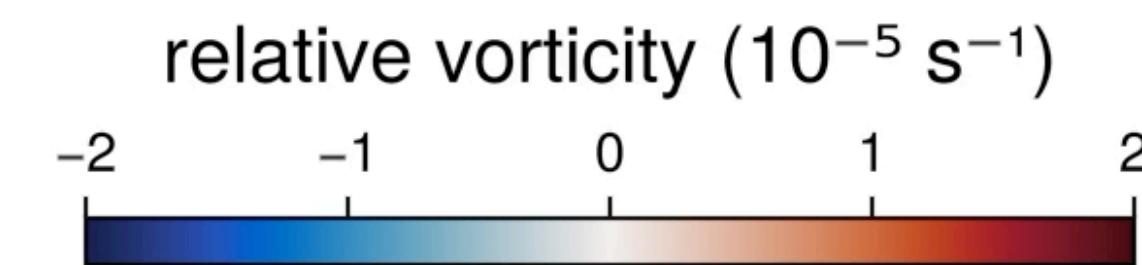
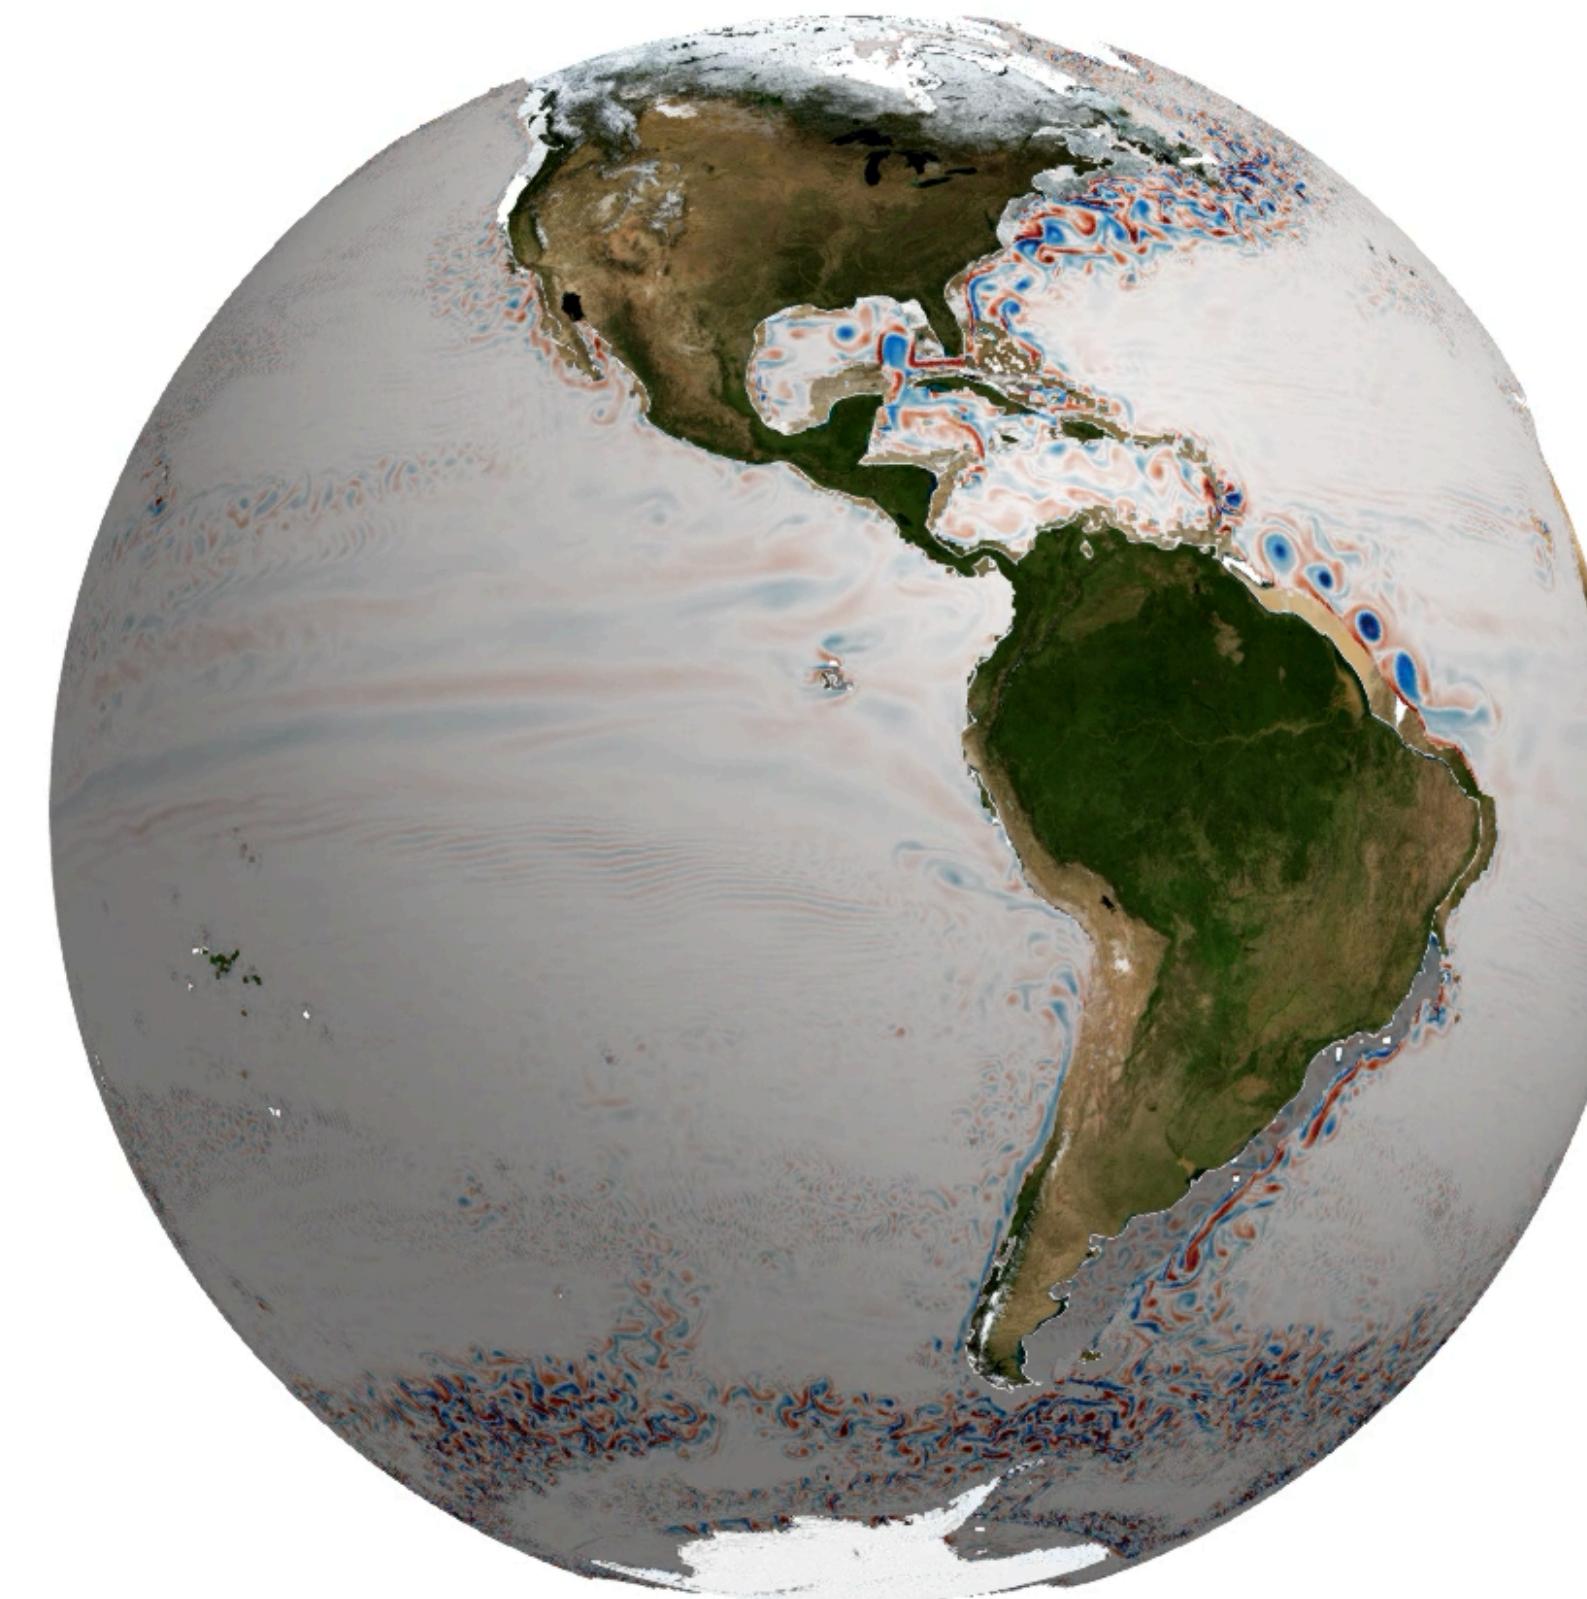
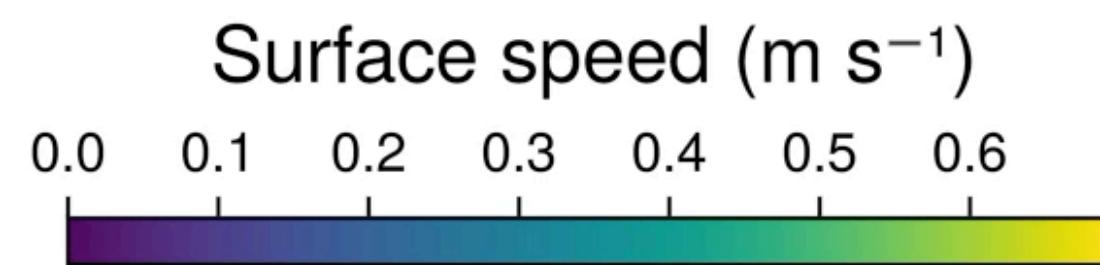
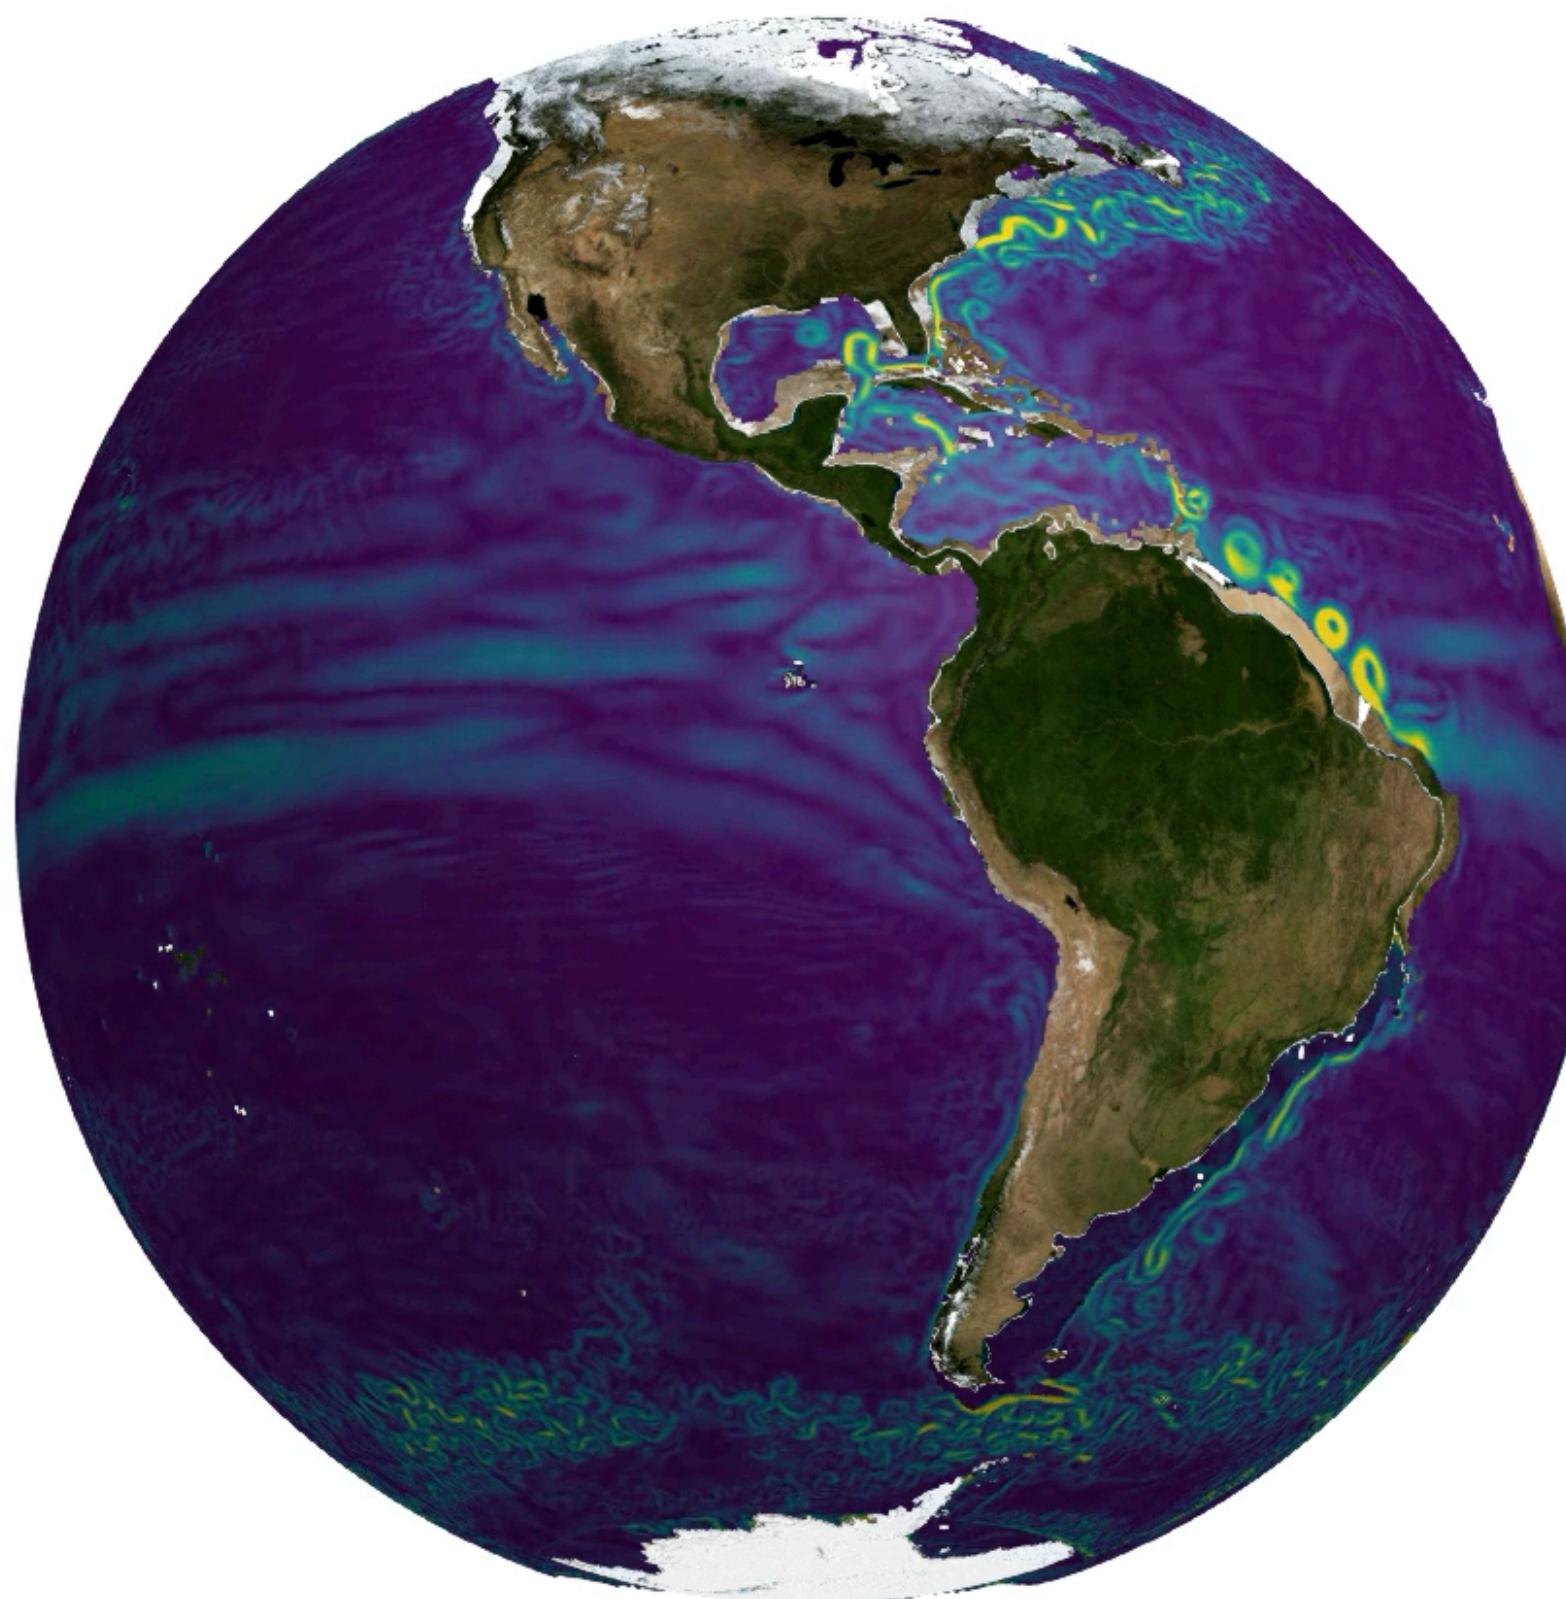
simulation = Simulation(coupled_model, Δt=5minutes, stop_time=30days)

run!(simulation)

```



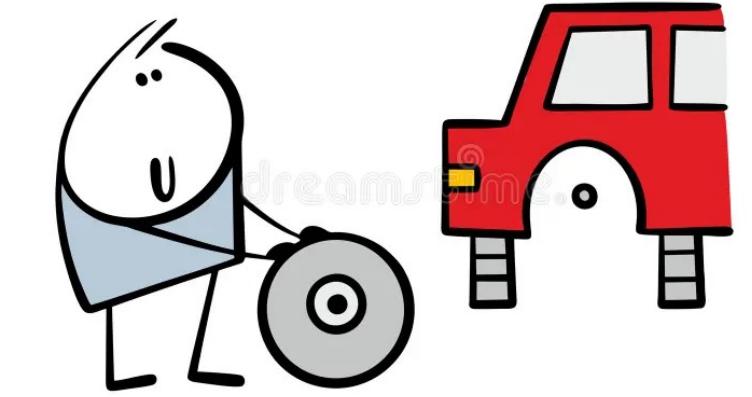
165.0 days



# Modeling software

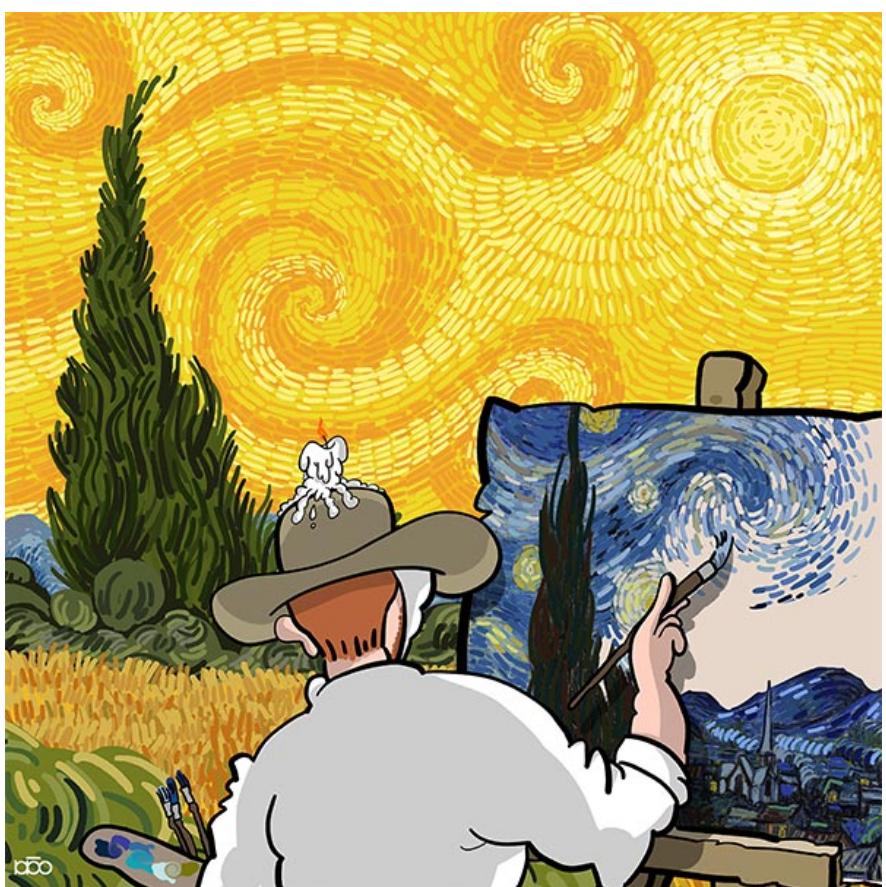


switch model components in and out easy



modify forcing, boundary conditions,

```
div = ∂x(u) + ∂y(v) + ∂z(w)  
set!(windstress = cos(latitude)^2)  
save(abs(∇T)^2, "temperature_dissipation.nc")
```



easy to share with other to reproduce results



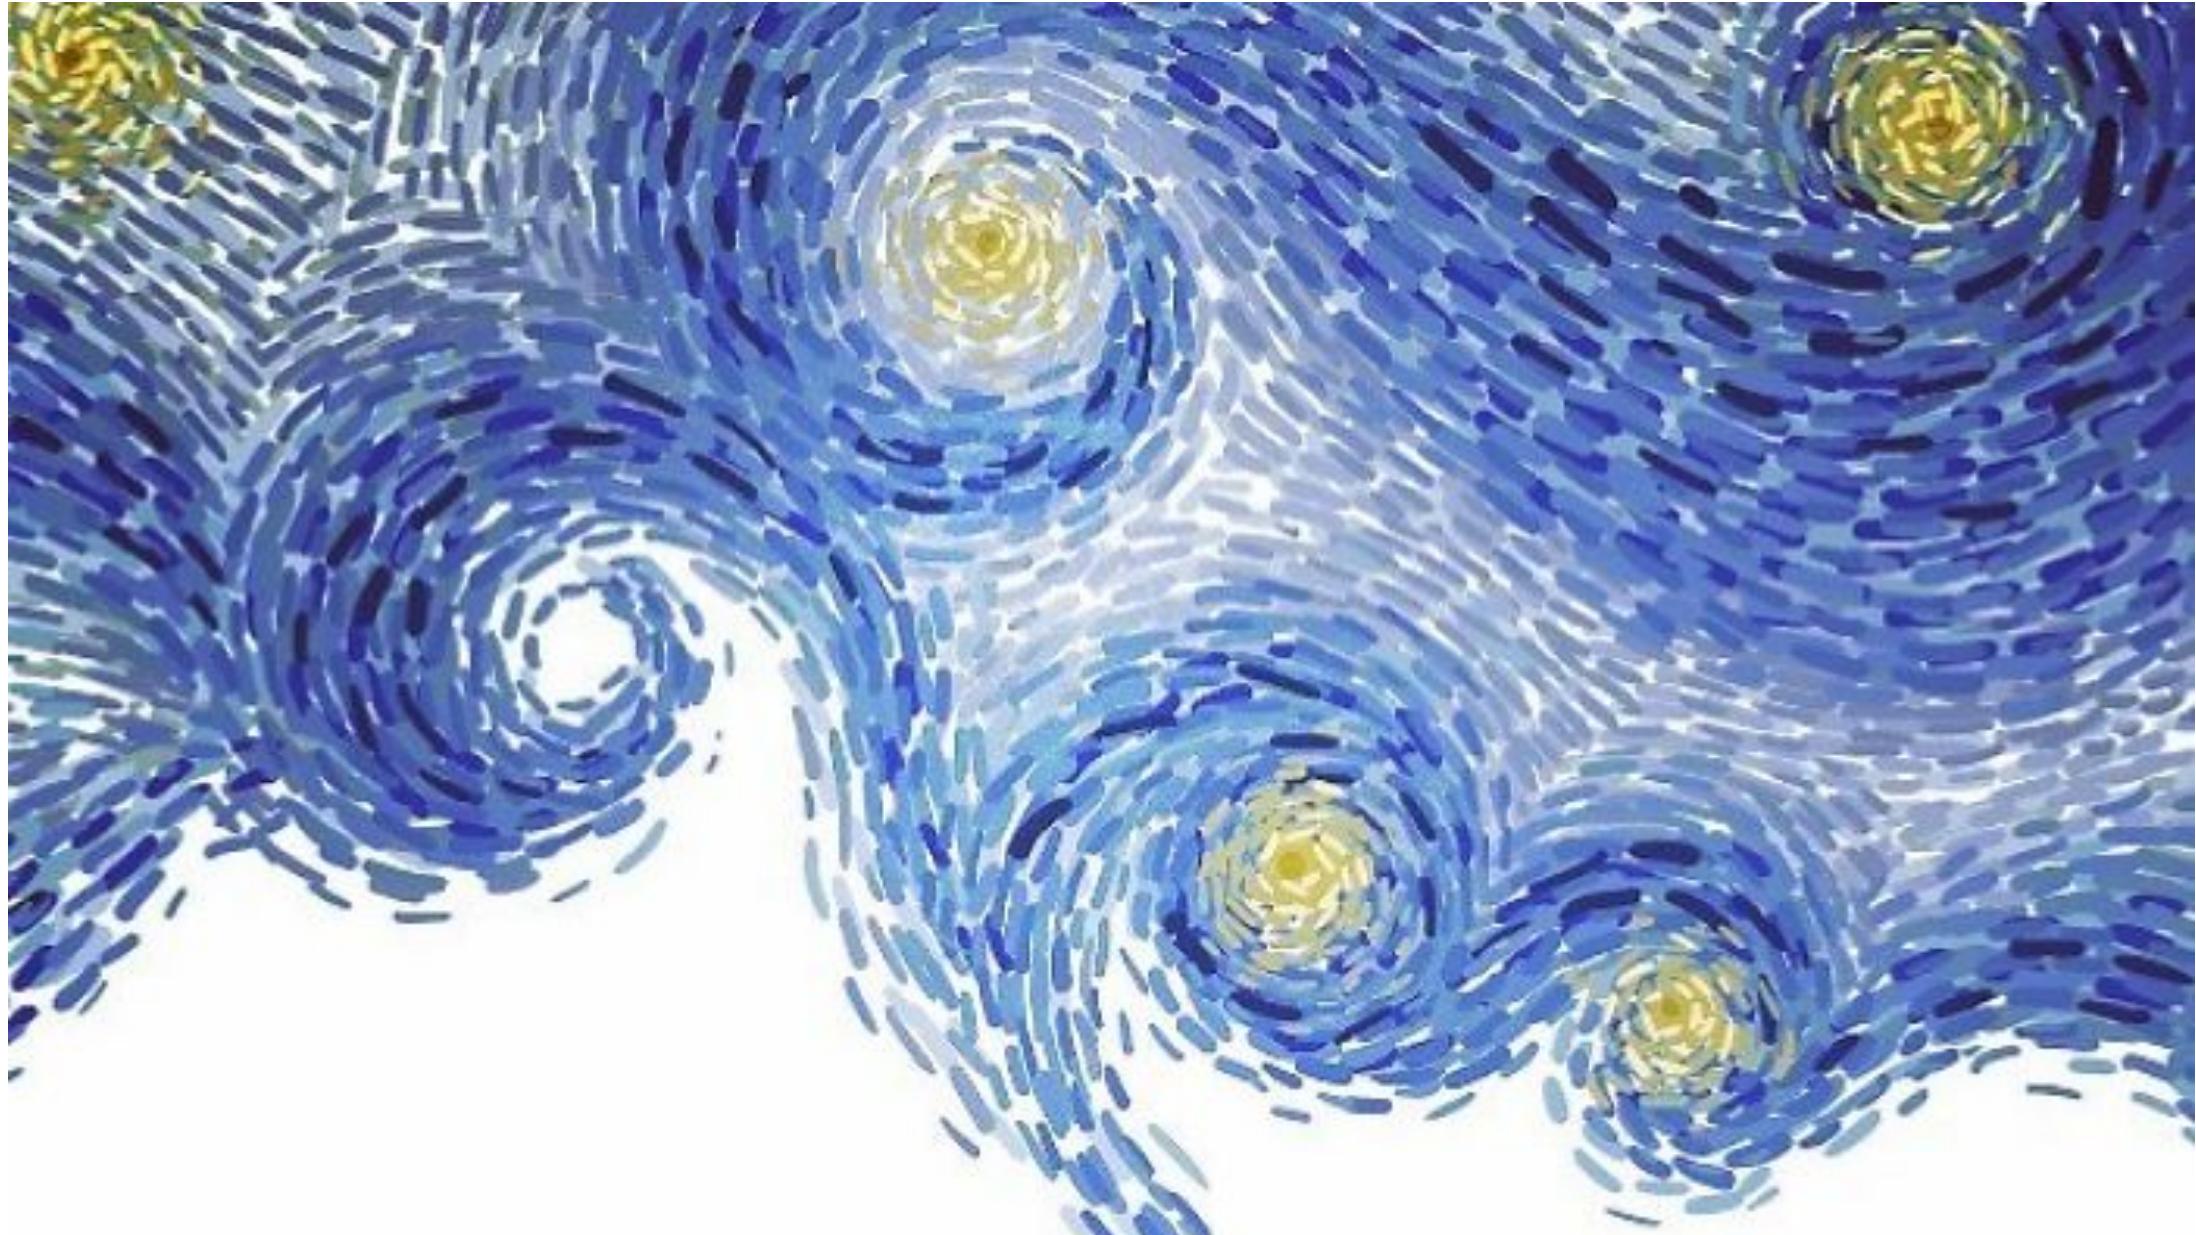
easy to visualise

# Modeling software



# anything to take home?

the ocean is beautiful with  
mesmerizing flow patterns!



models are our “eyes” to the future  
(climate projections)  
+ indispensable for research & teaching

software and science go hand in hand 

but the effort we put in software is very small  
( $\ll 1$ )



[drawing by Alireza Karimi Moghaddam; cartoonist, illustrator]



cute 🐧... *but does it seem tad angry?*

*because we don't invest much effort in software for ocean science* 🌊 ...